

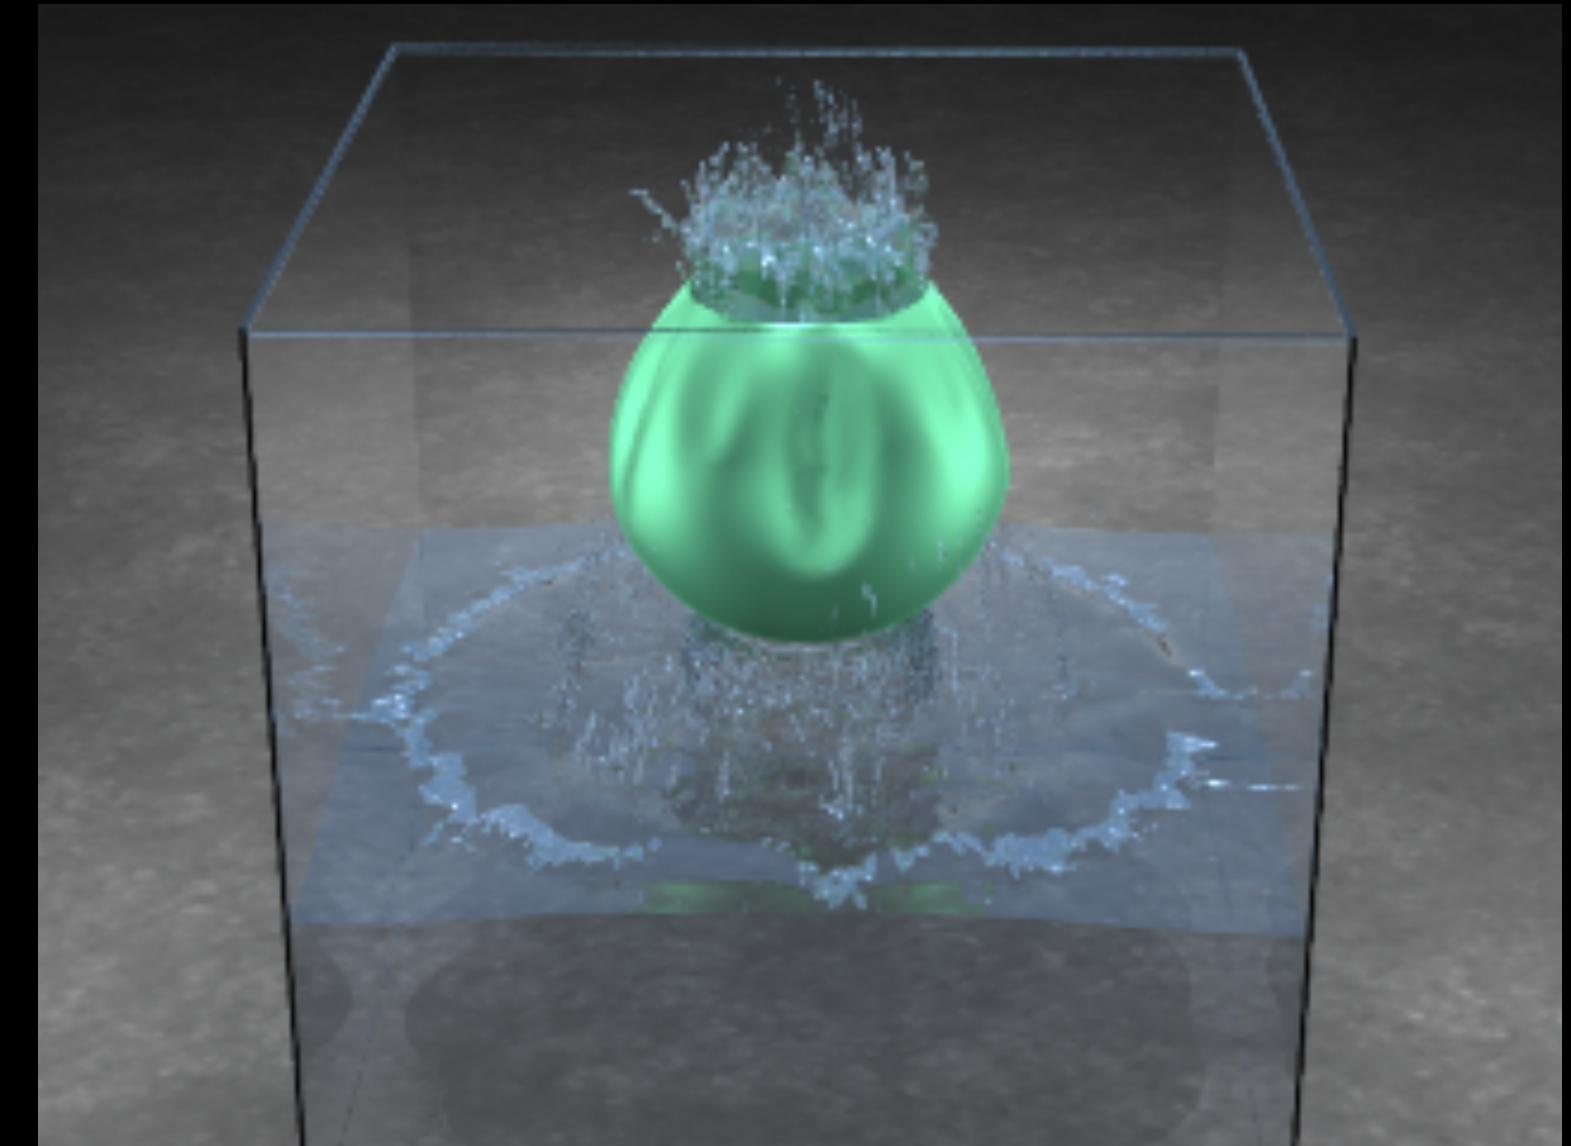
Unified Particle Physics for Real-Time Applications

Miles Macklin, Matthias Müller, Nuttapong Chentanez, Tae-Yong Kim

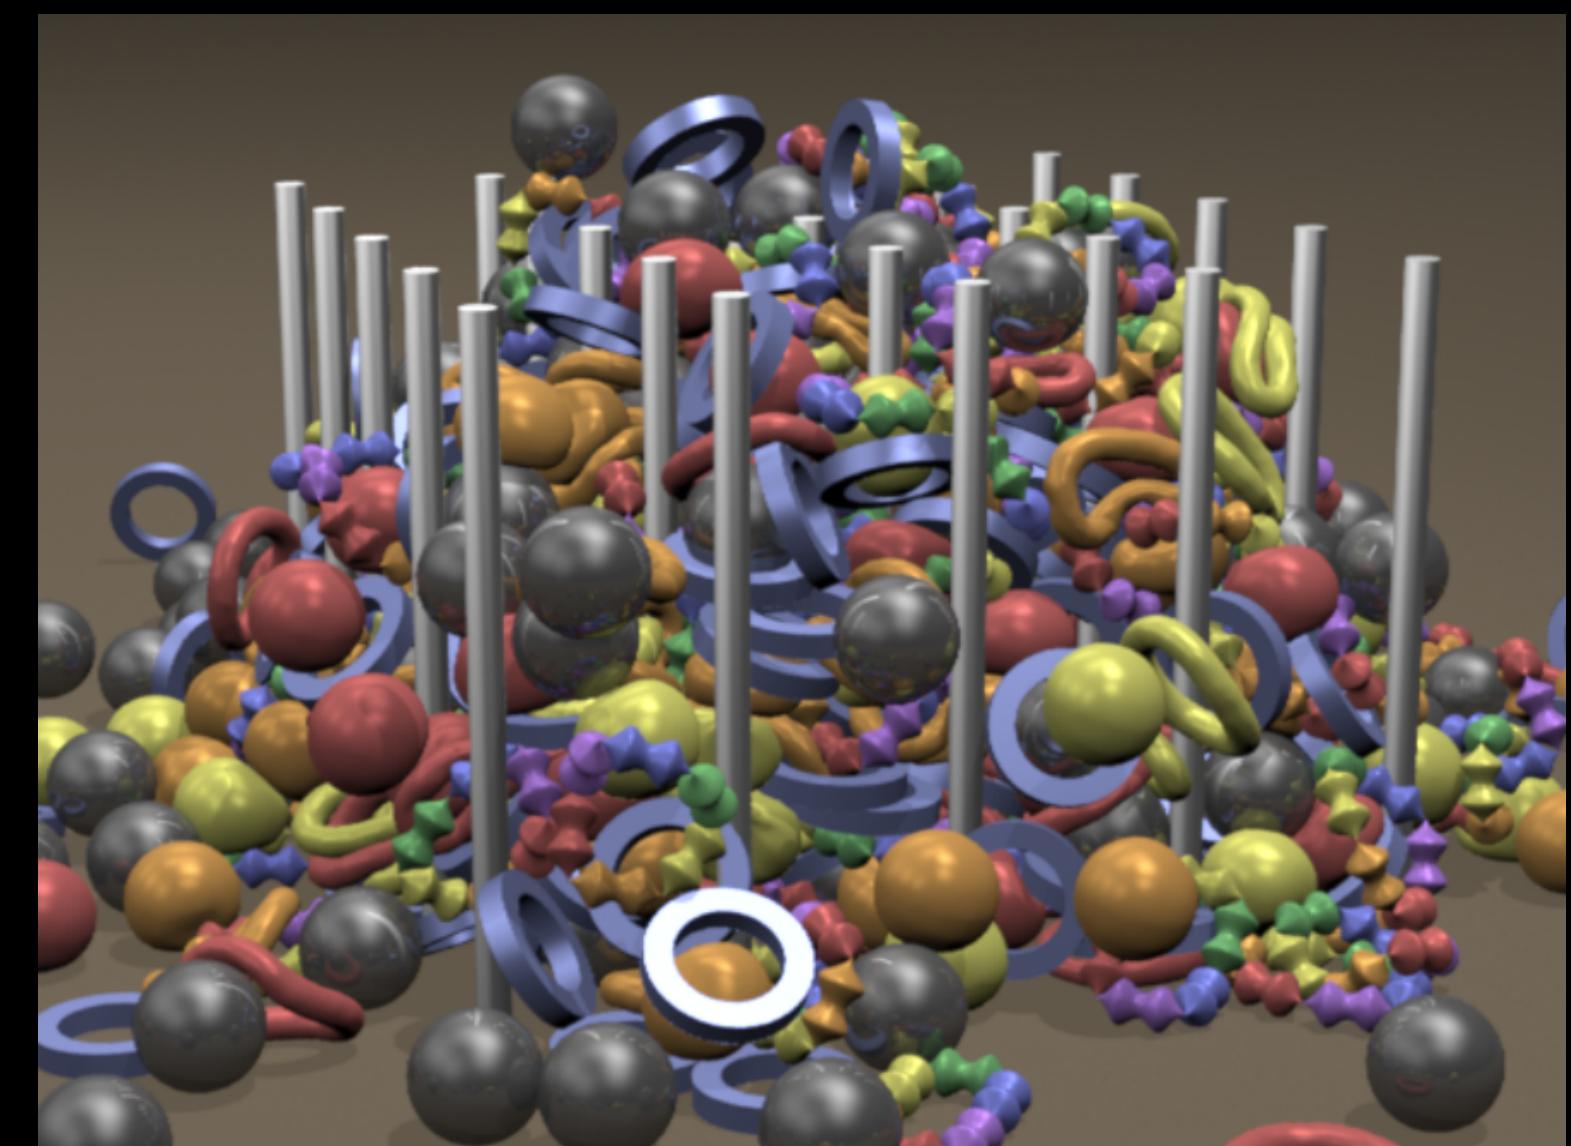


Motivation

- Too many solvers
- Creates redundant work
- Want two-way interaction between all object types



[Robinson-Mosher et al. 2008]



[Shinar et al. 2008]

Core Idea

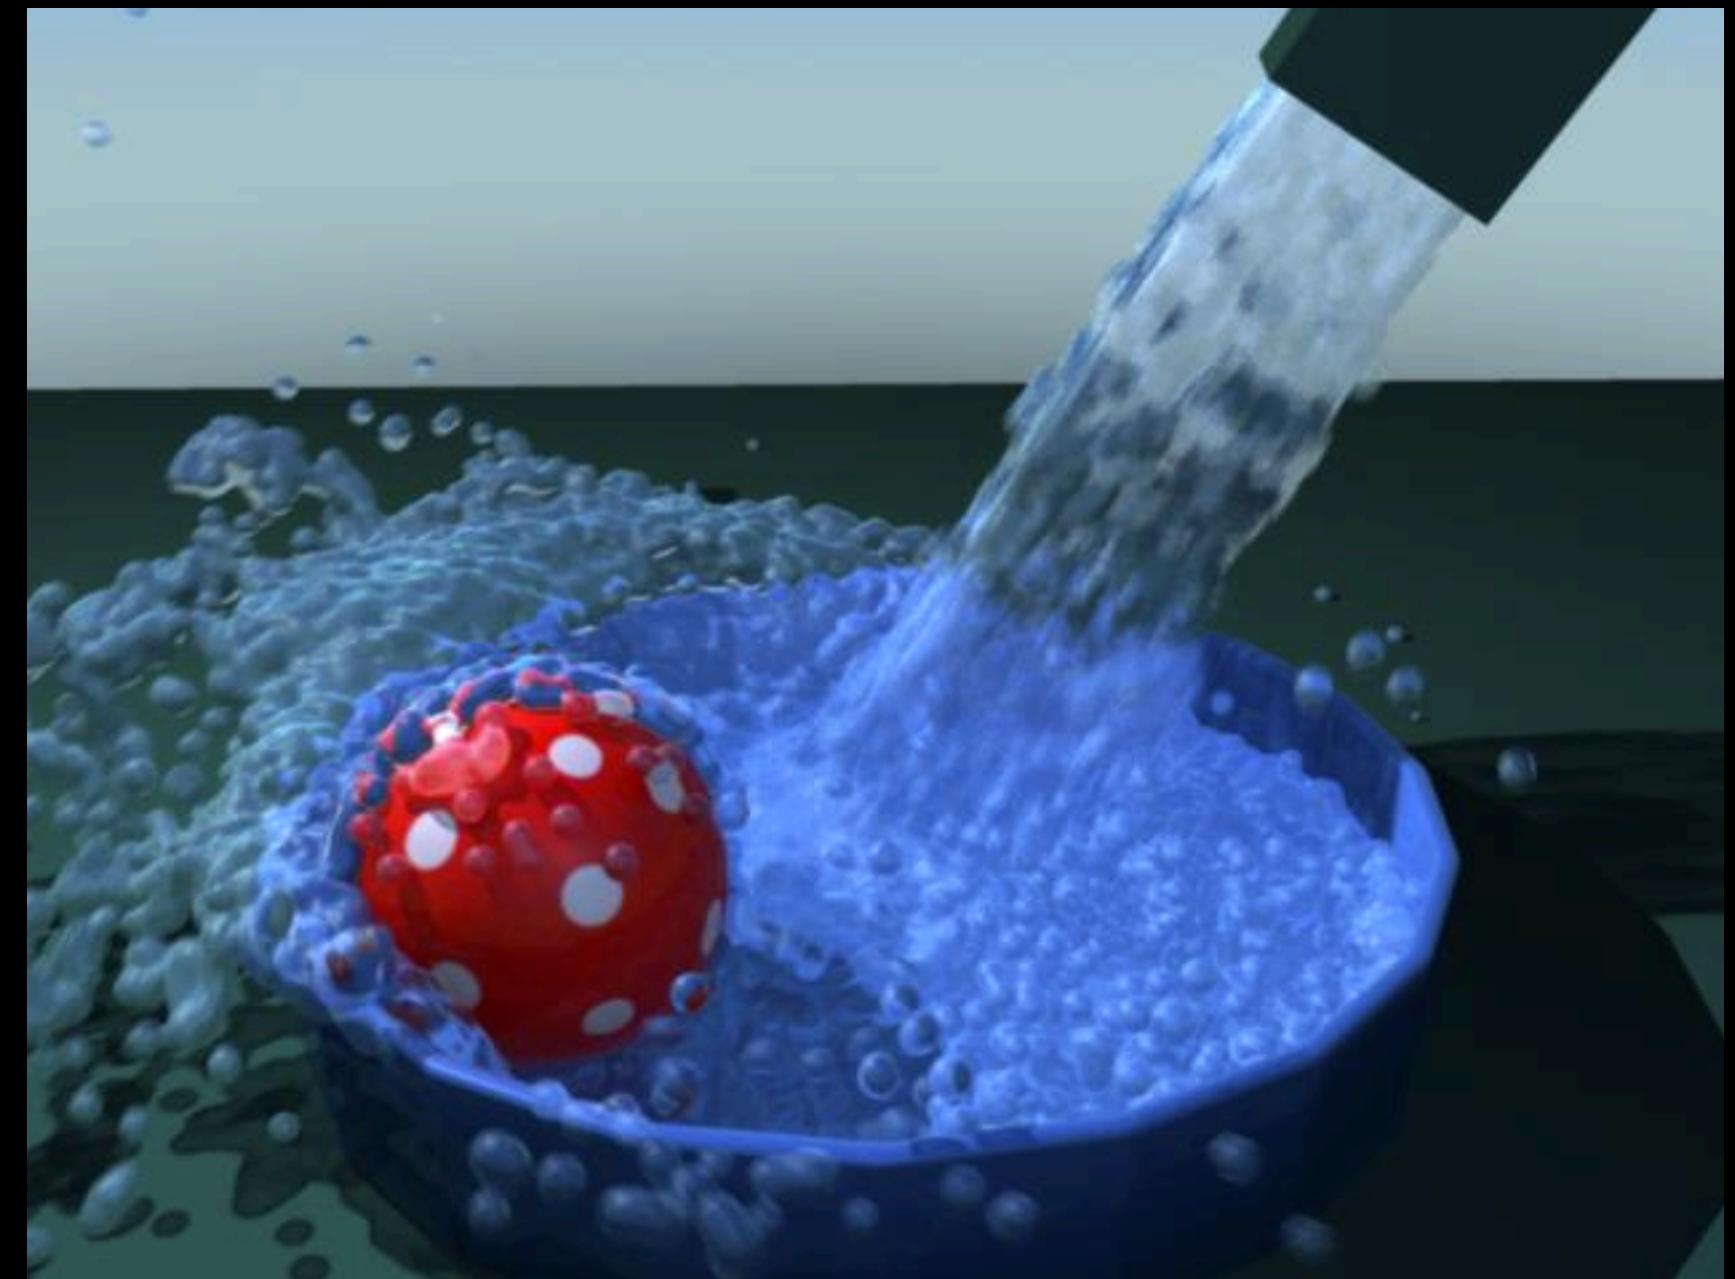
*Everything is a set of particles connected
by constraints*

Advantages

- Simplifies collision detection
- Stable two-way interaction of all object types:
 - ▶ Cloth
 - ▶ Deformables
 - ▶ Liquids
 - ▶ Gases
 - ▶ Rigid Bodies
- Fits well on the GPU

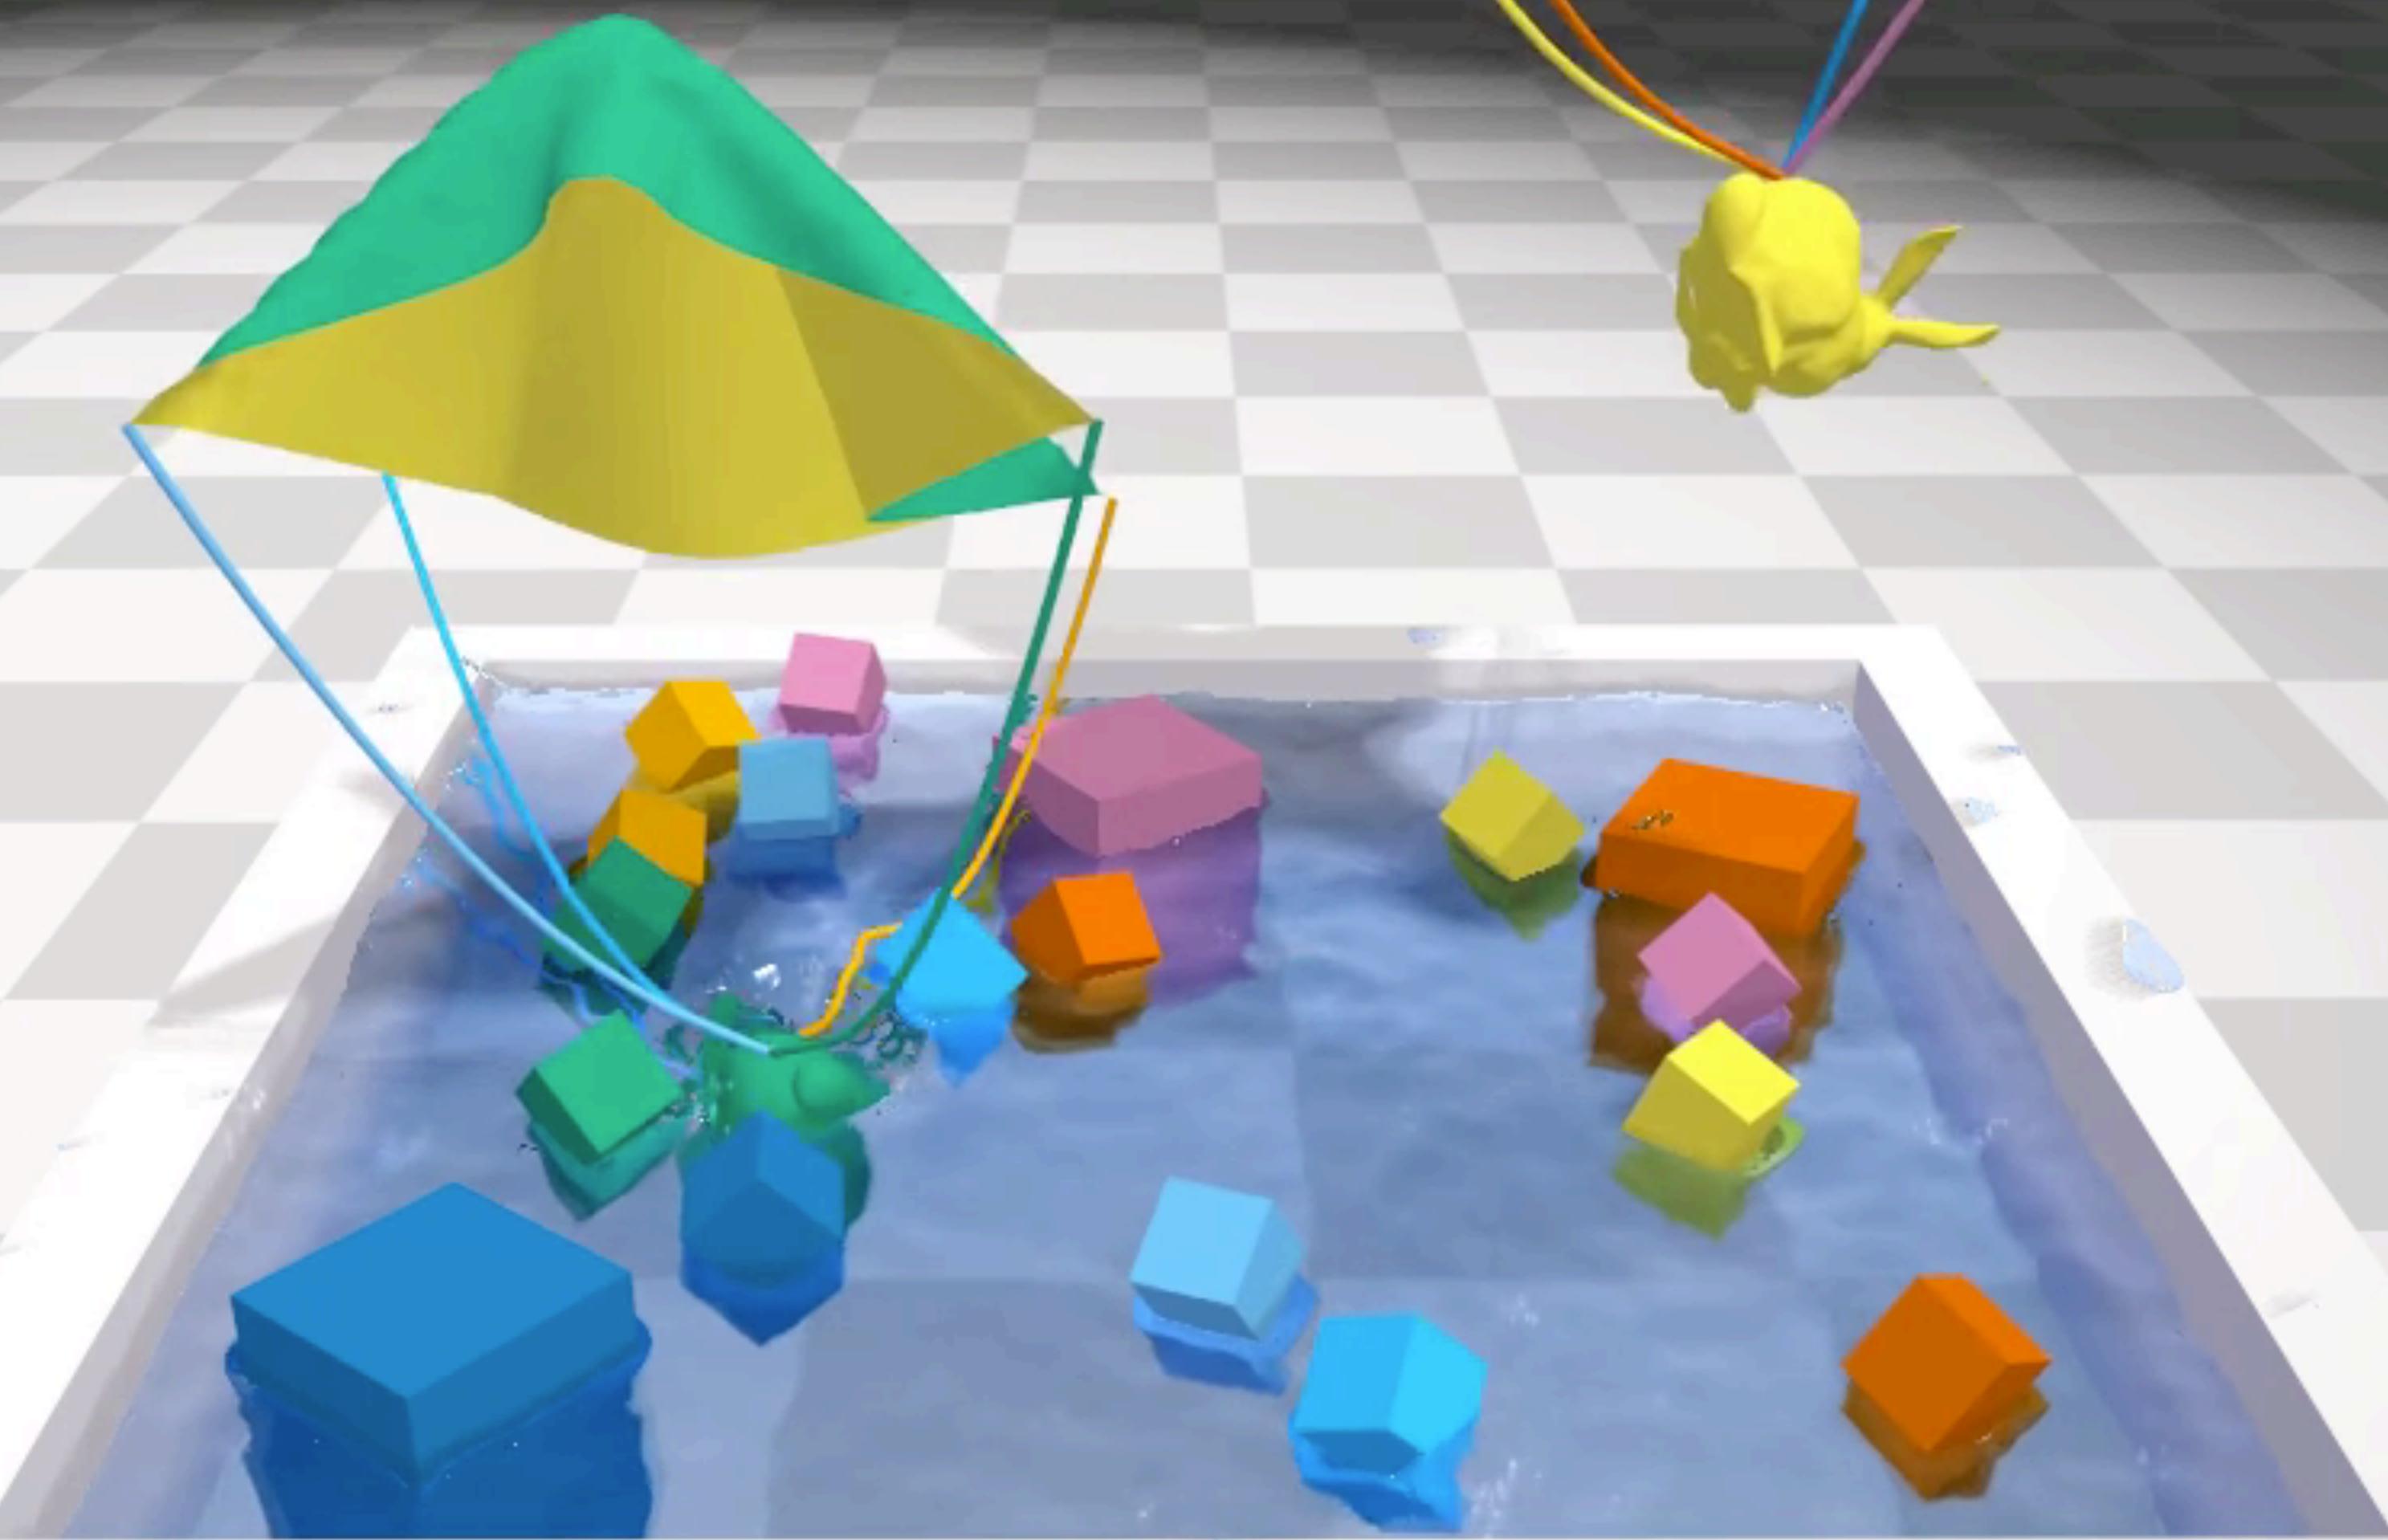
Related Work

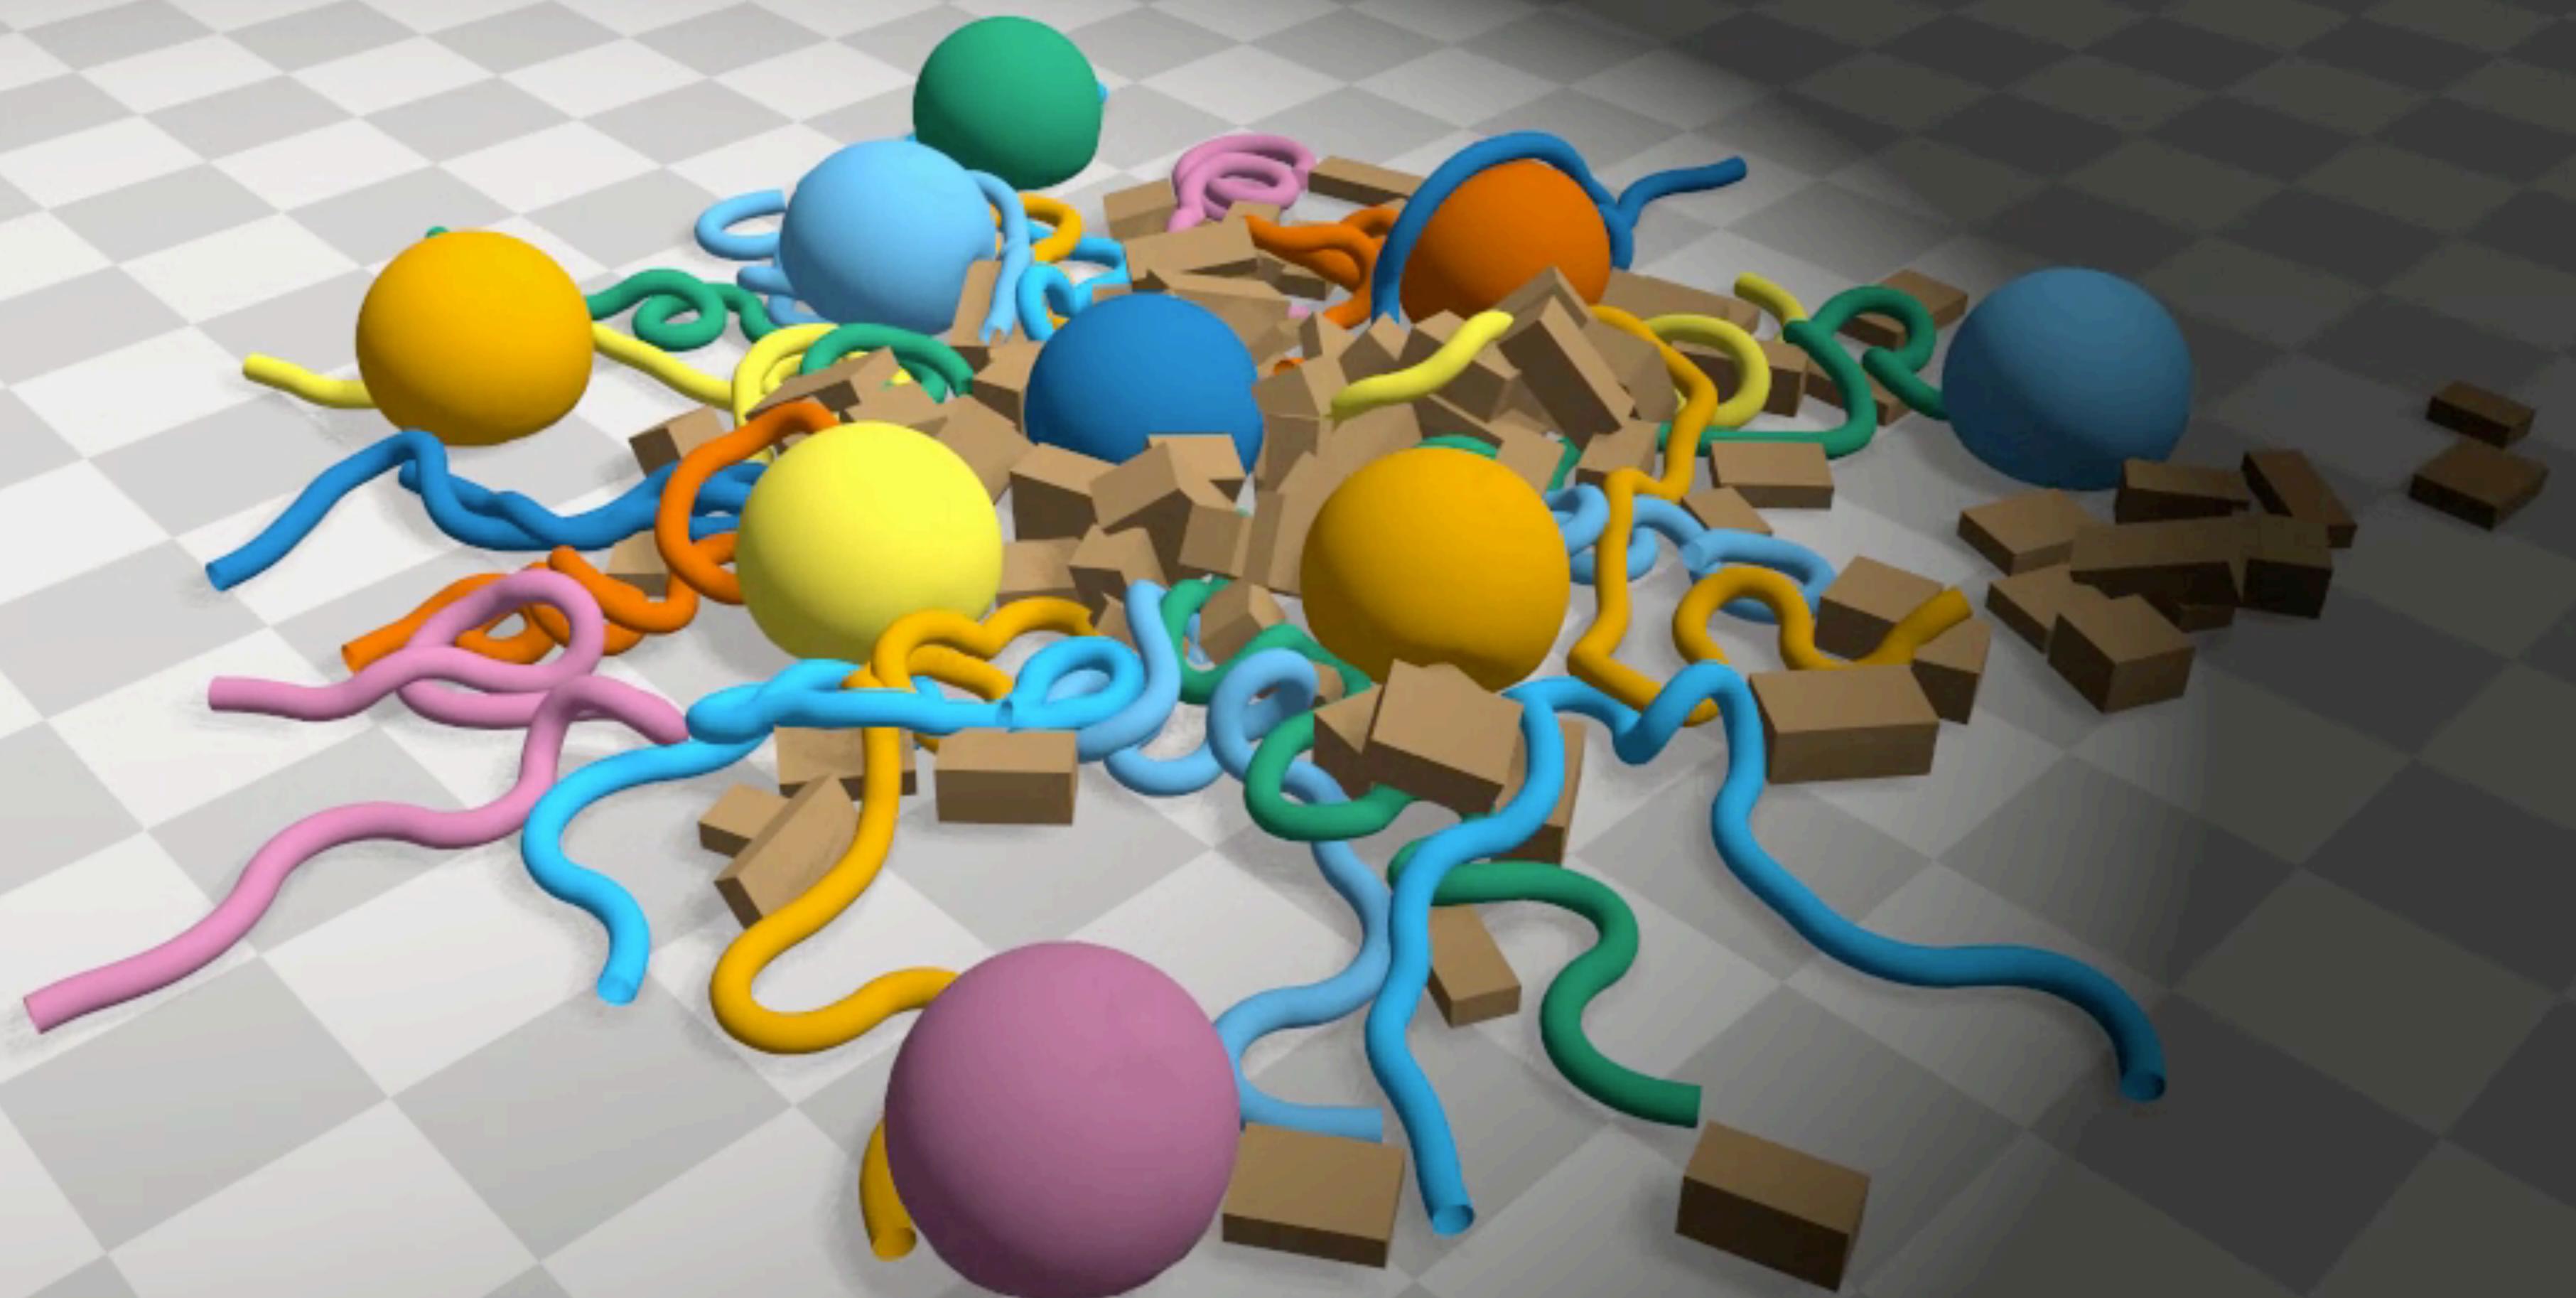
- Unified solvers popular in offline visual effects, e.g.:
 - ▶ Maya's Nucleus solver (nCloth, nParticles) [Stam09]
 - ▶ Softimage's Lagoa (fluids, elastics, granular materials)
- Goal: recreate these packages in real-time

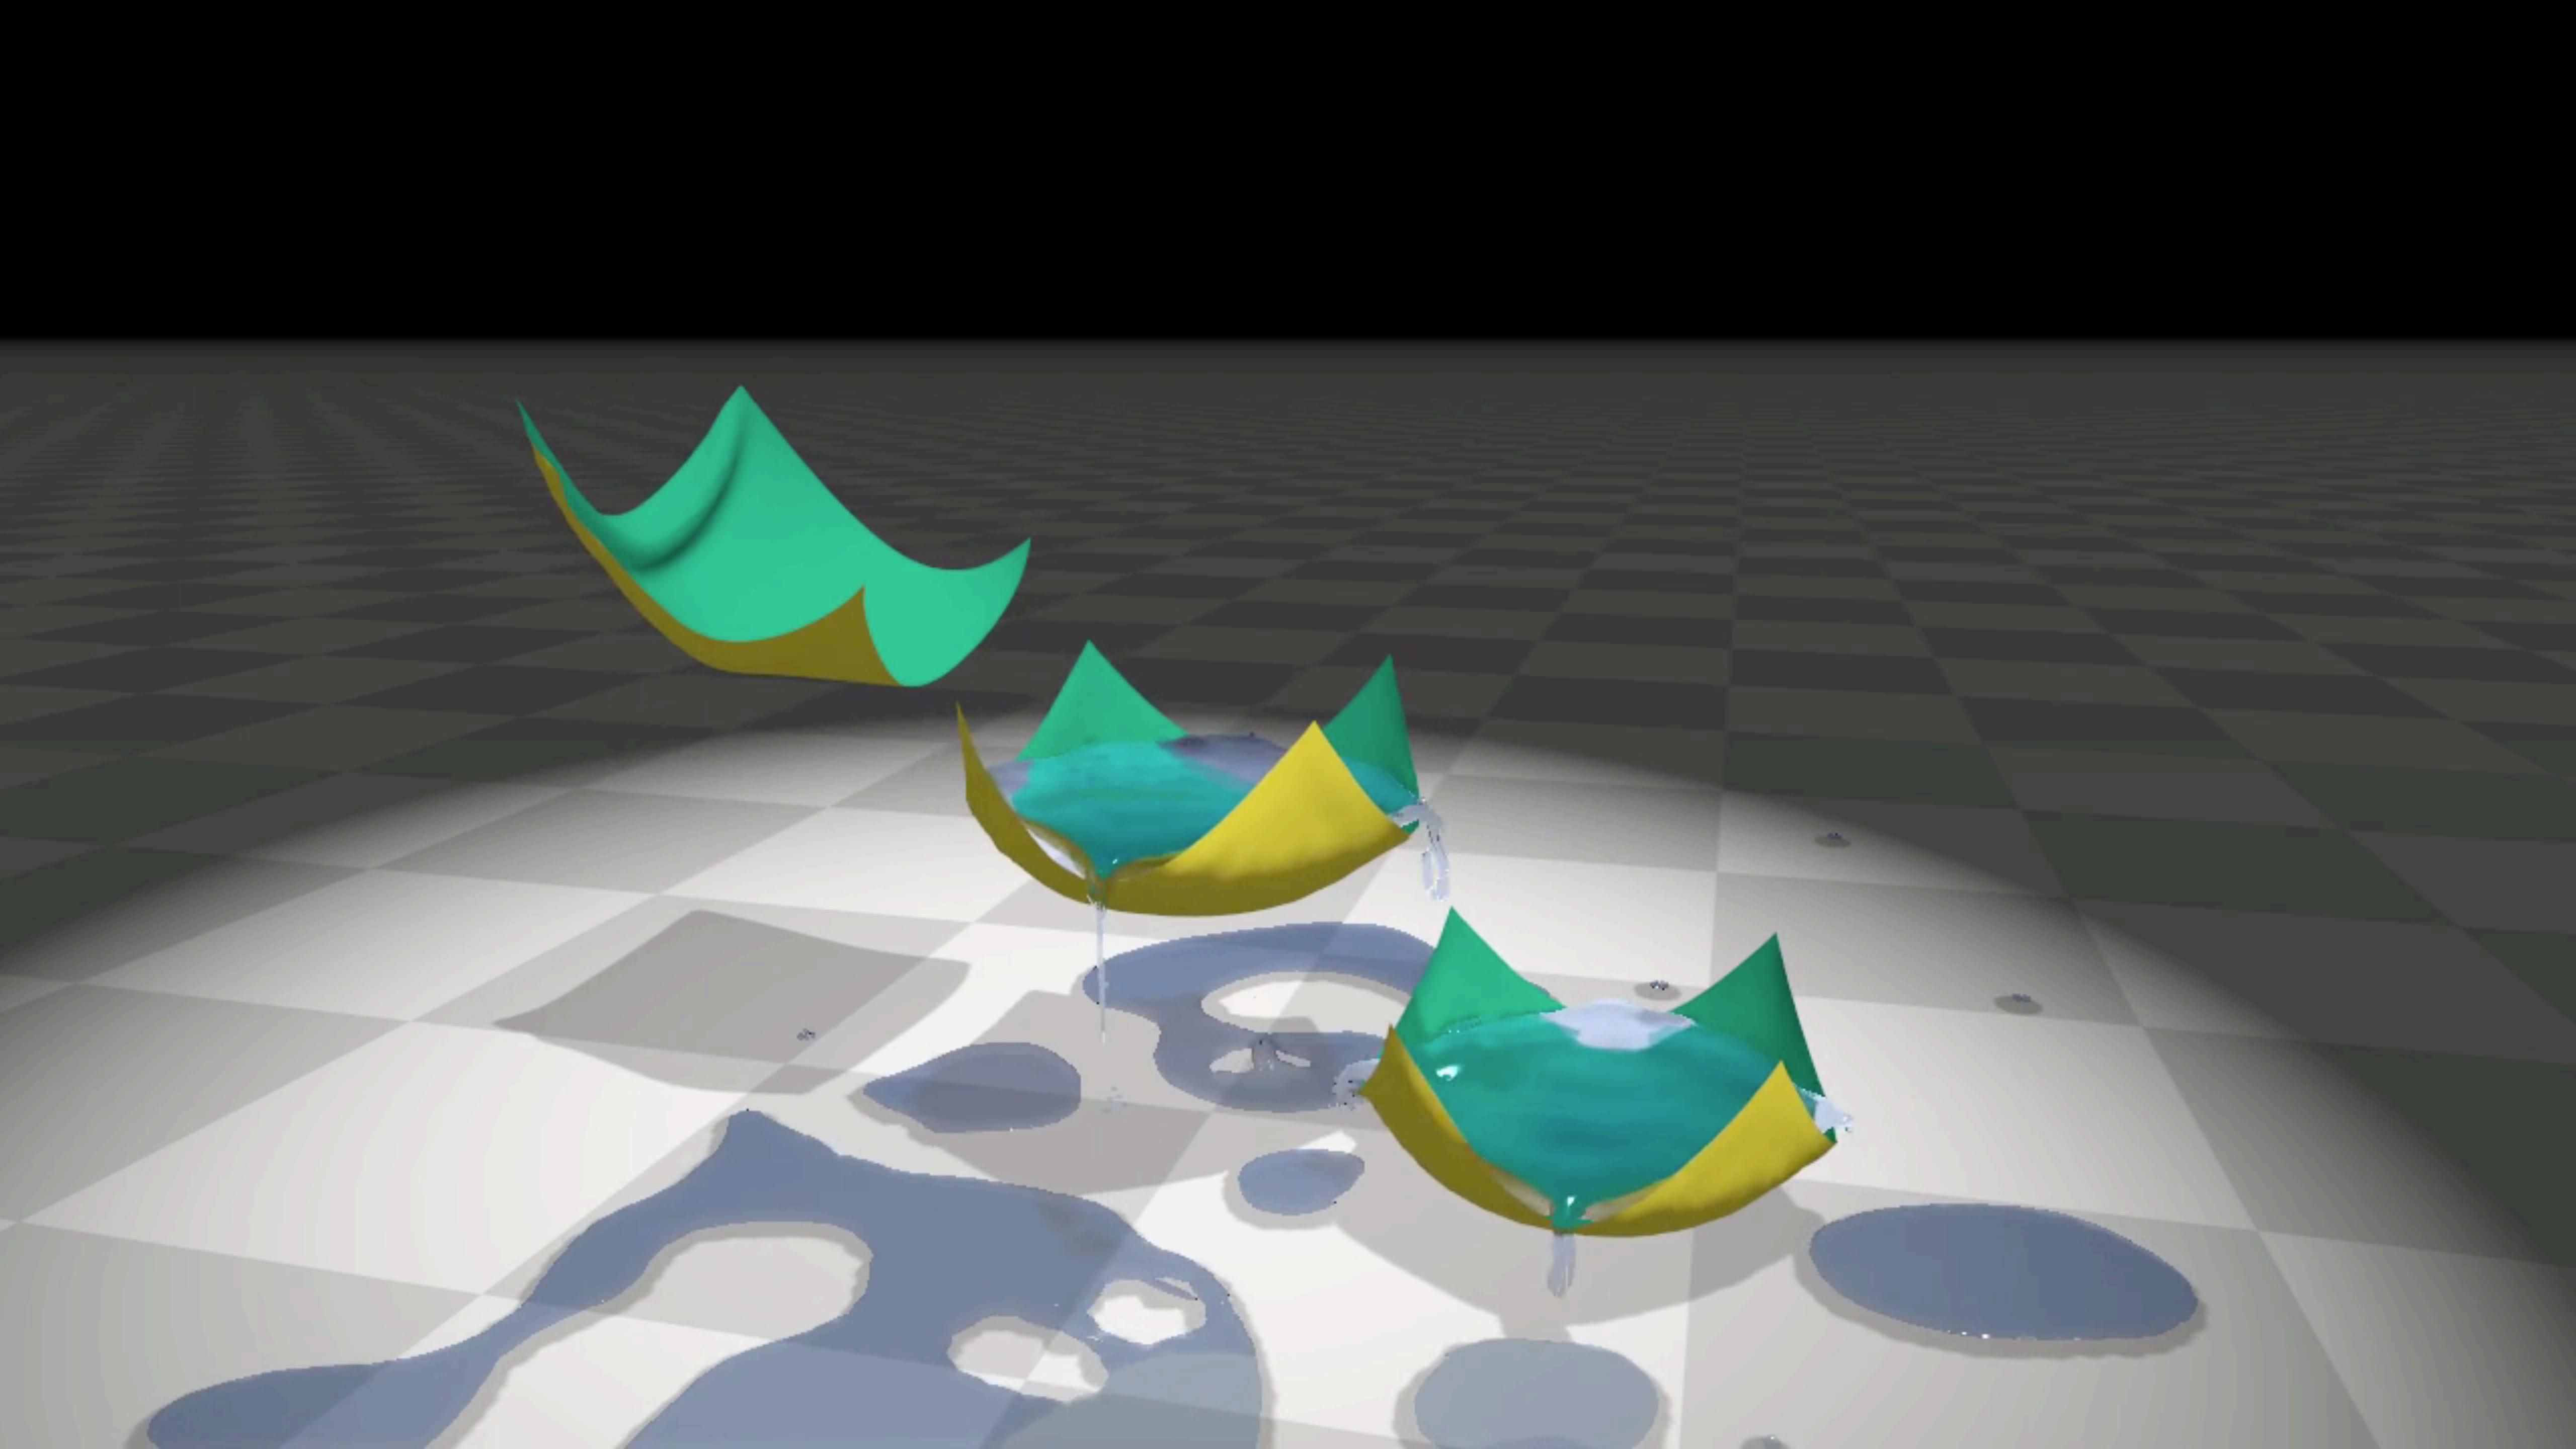


Maya nDynamics

Examples



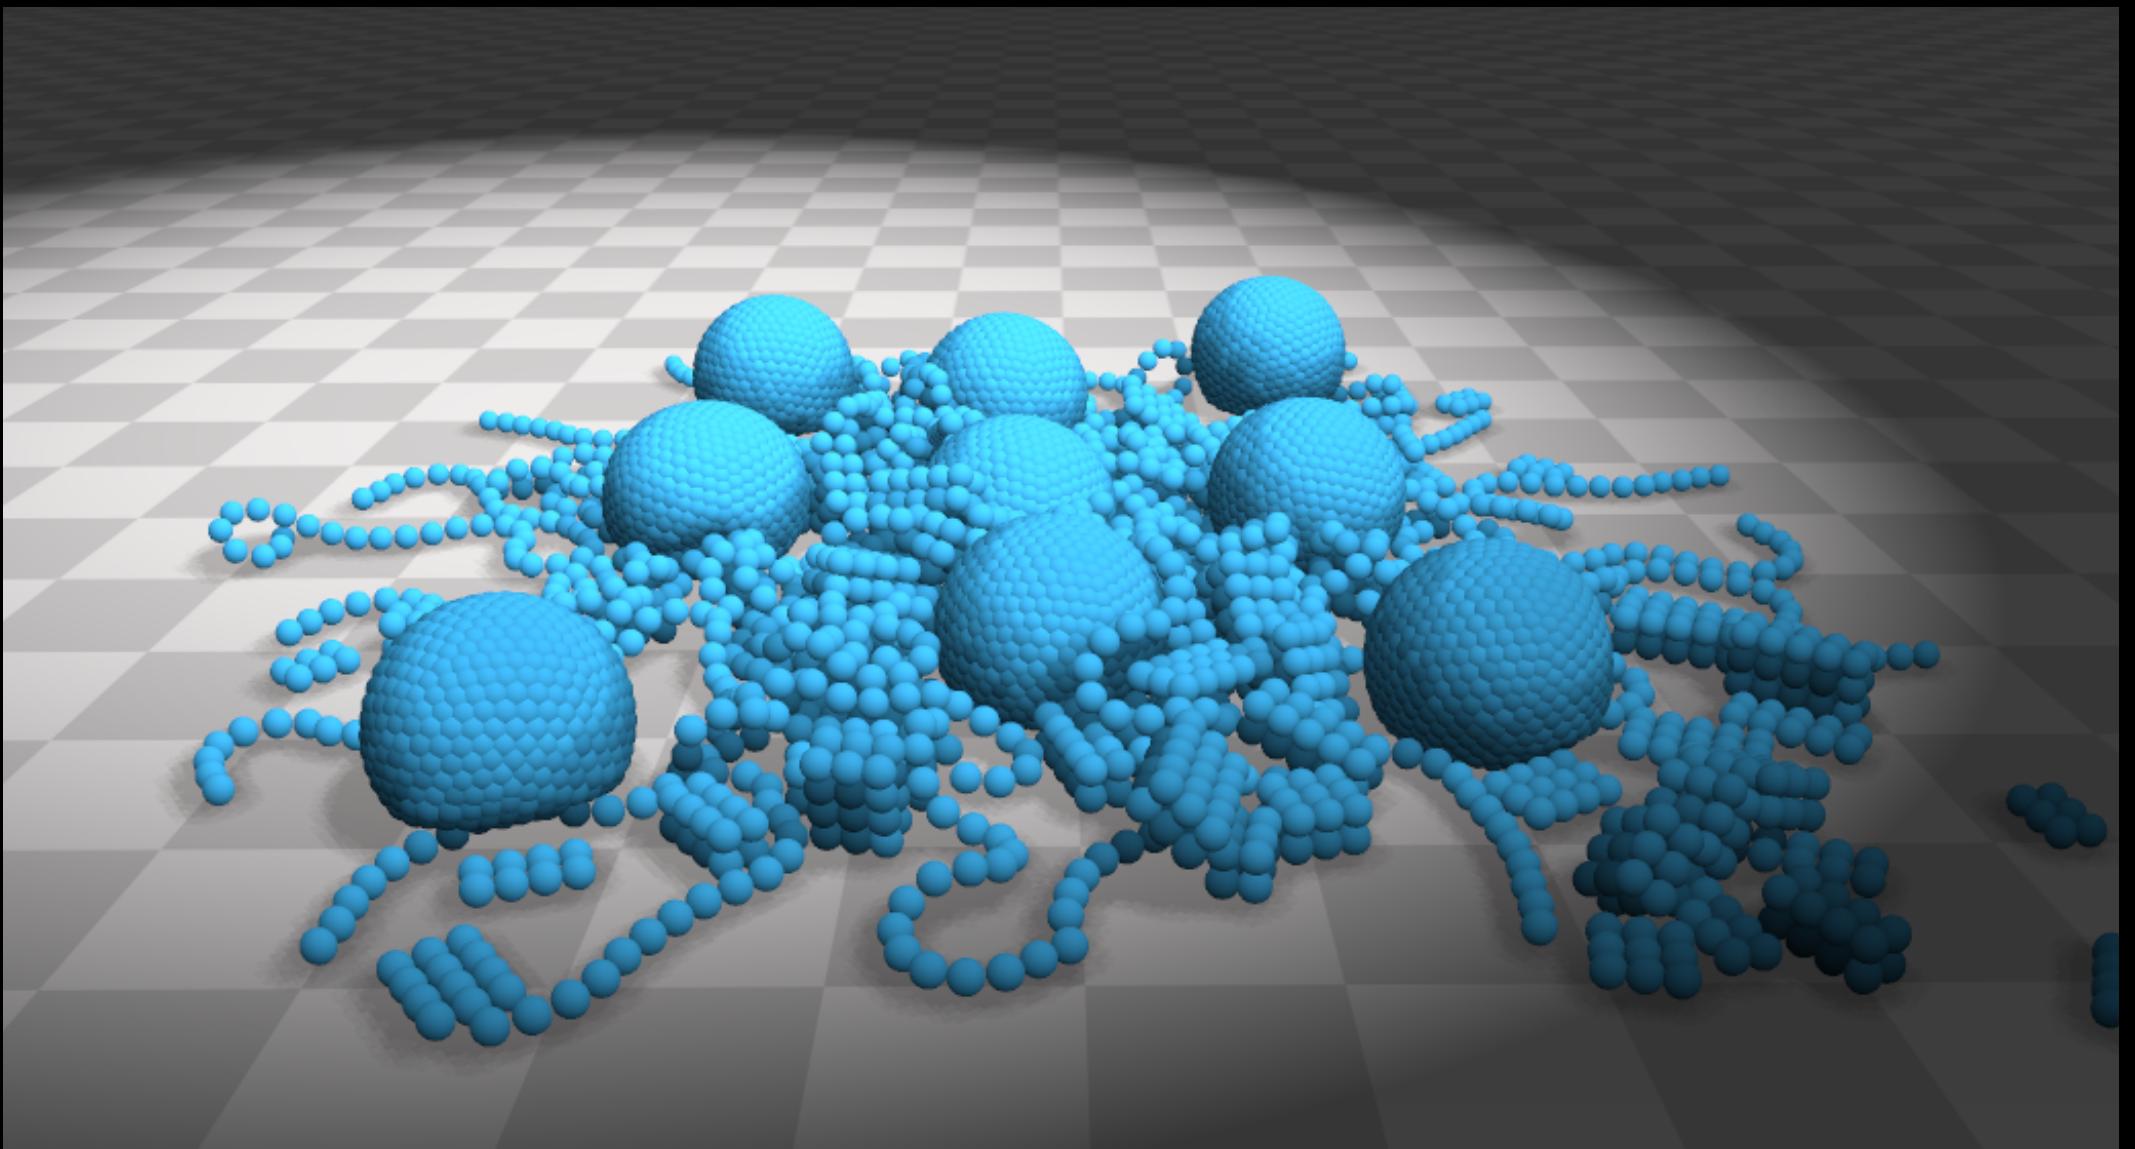
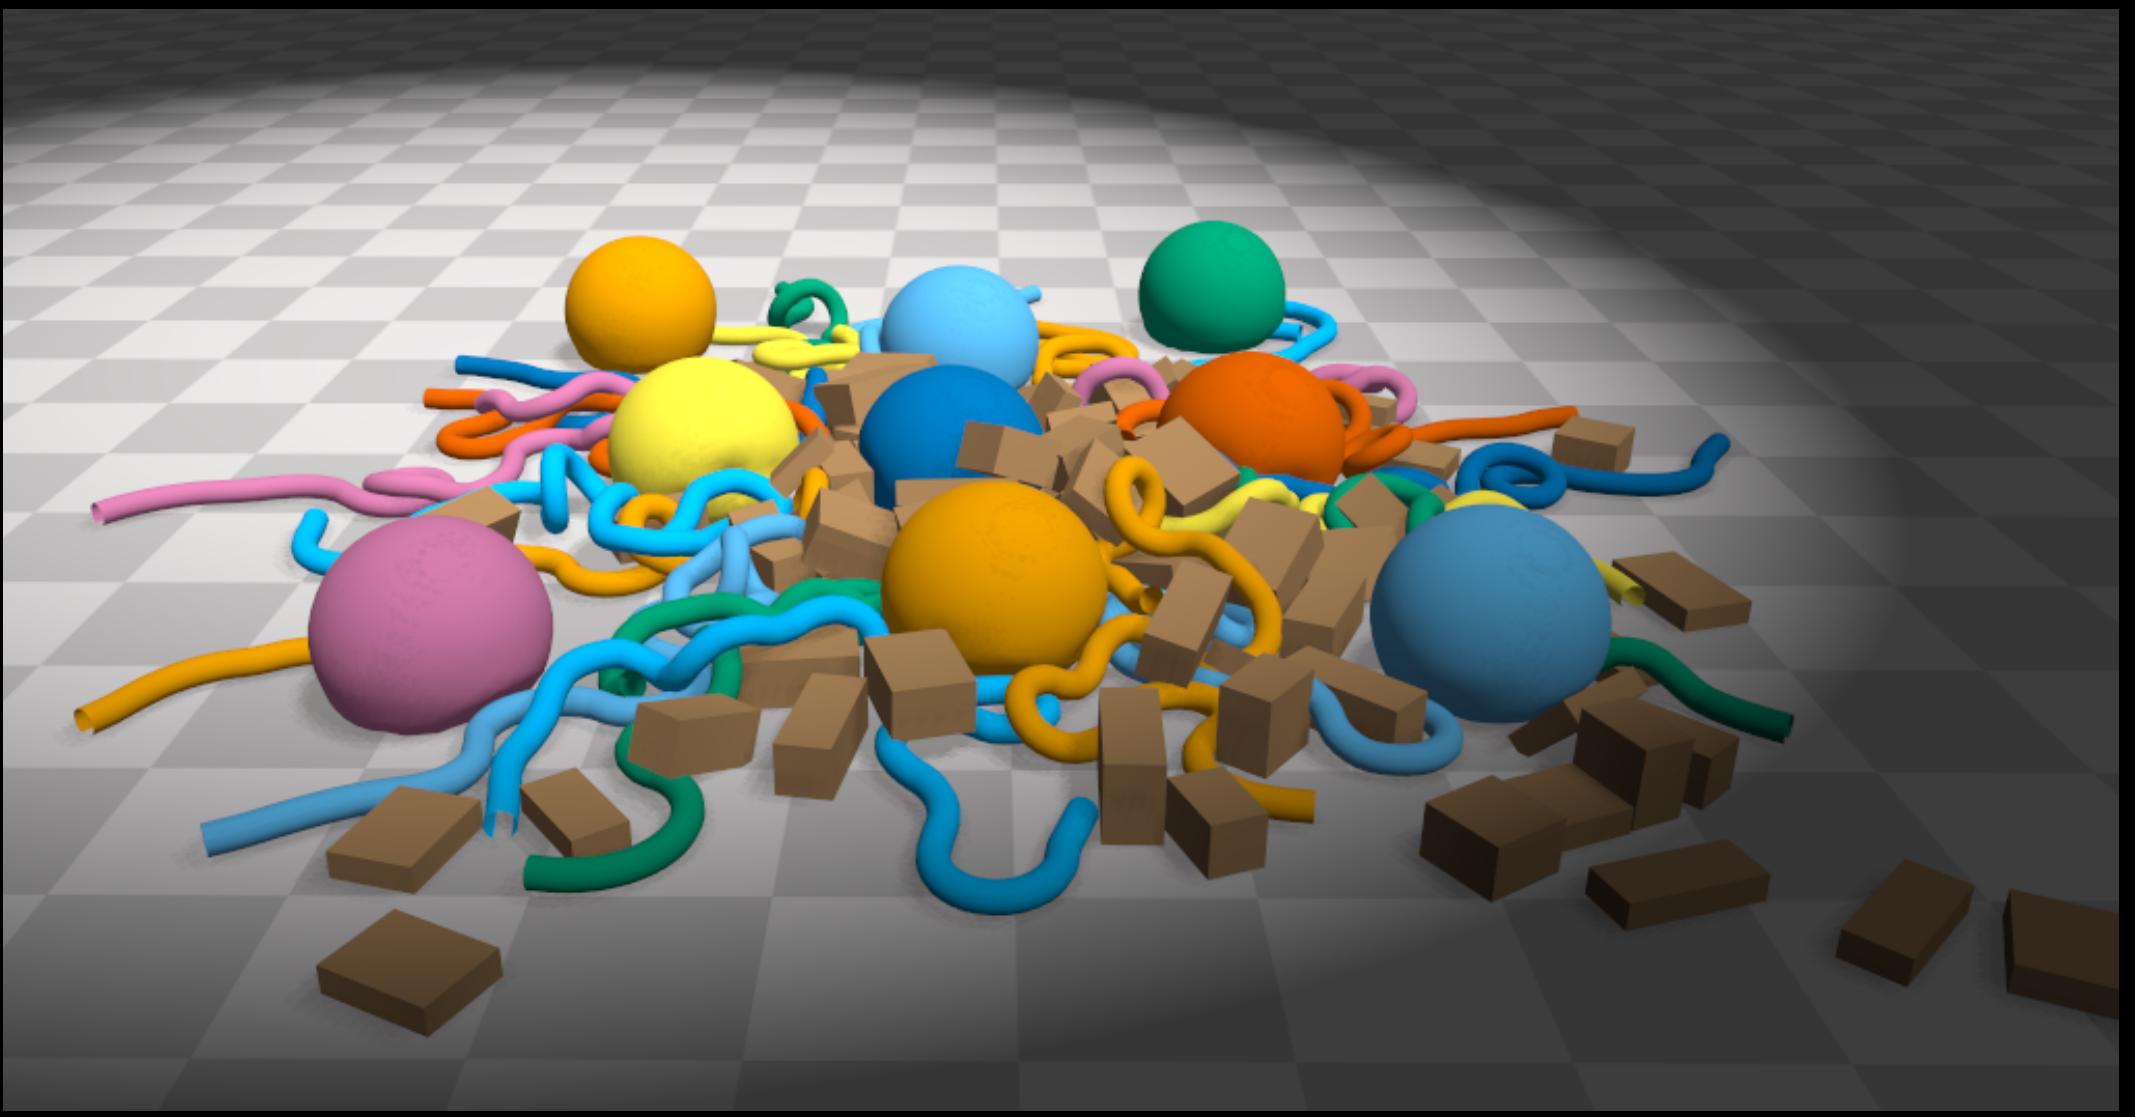




Particles

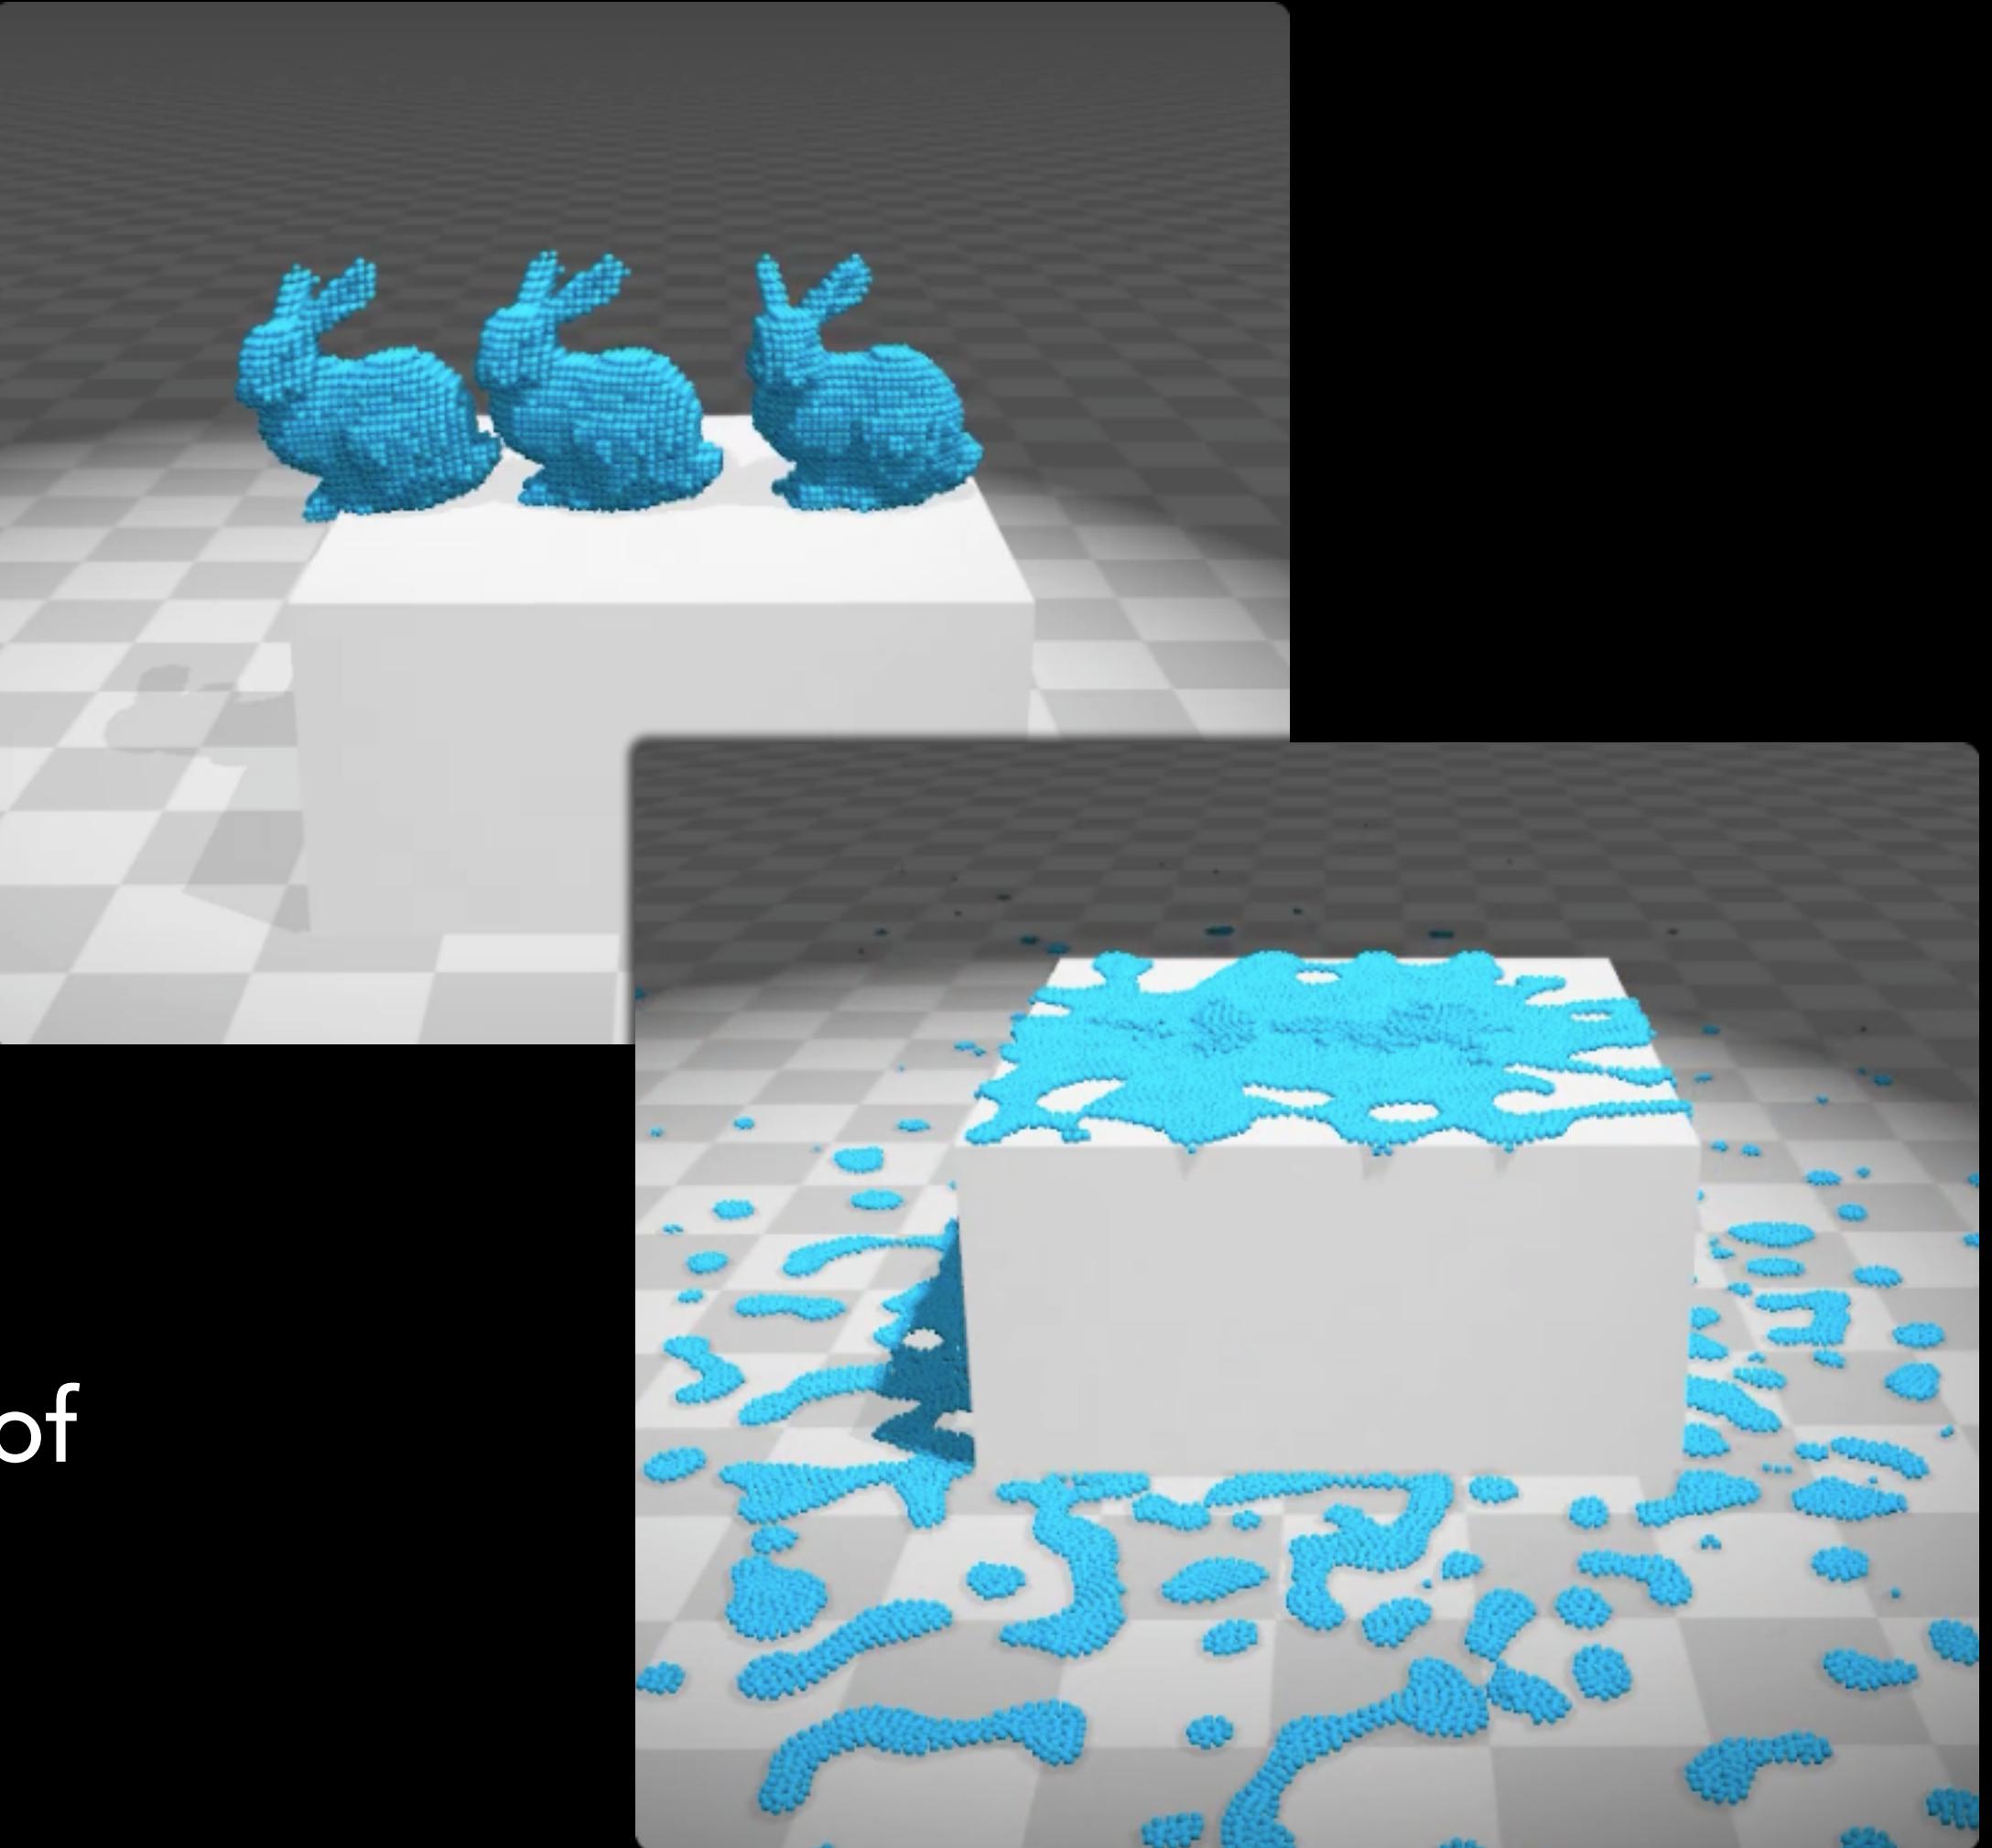
```
struct Particle
{
    float pos[3];
    float vel[3];
    float invMass;
    int phase;
};
```

- Phase-ID used to control collision filtering
- Particles do not belong to a particular object
- Single collision radius



Constraints

- Constraint types:
 - ▶ Distance (clothing)
 - ▶ Shape (rigids, plastics)
 - ▶ Density (fluids)
 - ▶ Volume (inflatables)
 - ▶ Contact (non-penetration, friction)
- Combine constraints to create wide variety of effects
 - ▶ Melting, phase-changes
 - ▶ Stiff cloth, bent metal



Talk Outline

1. Parallel Solver
2. Contact and friction
3. Rigid bodies
4. Gases

Position-Based Dynamics (PBD)

- Predict
- For $k=0$ to solver iterations

- Project

- or Minimize

$$\tilde{\mathbf{x}} = \mathbf{x}^n + \Delta t \mathbf{v}^n$$

$\tilde{\mathbf{x}}^{k+1} = \text{project } (\tilde{\mathbf{x}}^k, C) \text{ along } \mathbf{M}^{-1} \nabla C_{\mathbf{x}^k}$

$$\begin{aligned} \min_{\tilde{\mathbf{x}}^{k+1}} \quad & \frac{1}{2} (\tilde{\mathbf{x}}^{k+1} - \tilde{\mathbf{x}}^k)^T \mathbf{M} (\tilde{\mathbf{x}}^{k+1} - \tilde{\mathbf{x}}^k) \\ \text{s.t.} \quad & C_i(\tilde{\mathbf{x}}^{k+1}) = 0 \end{aligned}$$

- Velocity Update

$$\mathbf{v}^{n+1} = \mathbf{x}^* - \mathbf{x}^n$$

- Position Update

$$\mathbf{x}^{n+1} = \mathbf{x}^*$$

Relationship to Implicit Euler

- Position level formulation of backwards Euler:

$$\mathbf{M}(\mathbf{x}^{n+1} - 2\mathbf{x}^n + \mathbf{x}^{n-1}) = \Delta t^2 \mathbf{f}(\mathbf{x}^{n+1})$$

- Can be seen as first order optimality condition for the following minimization:

$$\min_{\mathbf{x}^{n+1}} \quad \frac{1}{2} (\mathbf{x}^{n+1} - \tilde{\mathbf{x}})^T \mathbf{M} (\mathbf{x}^{n+1} - \tilde{\mathbf{x}}) + \Delta t^2 E(\mathbf{x}^{n+1})$$

- Predicted (inertial) position: $\tilde{\mathbf{x}} = 2\mathbf{x}^n - \mathbf{x}^{n-1}$
 $= \mathbf{x}^n + \Delta t \mathbf{v}^n$

Backward Euler as Constrained Minimization

- Constraints are infinitely stiff potentials

$$\min_{\mathbf{x}^{n+1}} \quad \frac{1}{2}(\mathbf{x}^{n+1} - \tilde{\mathbf{x}})^T \mathbf{M}(\mathbf{x}^{n+1} - \tilde{\mathbf{x}}) + \Delta t E(\mathbf{x}^{n+1})$$

- Produces the following constrained optimization:

$$\begin{aligned} \min_{\mathbf{x}^{n+1}} \quad & \frac{1}{2}(\mathbf{x}^{n+1} - \tilde{\mathbf{x}})^T \mathbf{M}(\mathbf{x}^{n+1} - \tilde{\mathbf{x}}) \\ \text{s.t.} \quad & C_i(\mathbf{x}^{n+1}) = 0 \end{aligned}$$

- Searching for the point closest to the predicted (inertial) position that lies on the constraint manifold

Implicit Euler

- Predict
- For k=0 to solver iterations

- Project

- or Minimize

$$\tilde{\mathbf{x}} = \mathbf{x}^n + \Delta t \mathbf{v}^n$$

$\tilde{\mathbf{x}}^{k+1} = \text{project } (\tilde{\mathbf{x}}, C) \text{ along } \mathbf{M}^{-1} \nabla C_{\mathbf{x}^k}$

$$\min_{\tilde{\mathbf{x}}^{k+1}} \frac{1}{2} (\tilde{\mathbf{x}}^{k+1} - \tilde{\mathbf{x}})^T \mathbf{M} (\tilde{\mathbf{x}}^{k+1} - \tilde{\mathbf{x}})$$

$$\text{s.t. } C_i(\tilde{\mathbf{x}}^{k+1}) = 0$$

- Velocity Update
- Position Update

$$\mathbf{v}^{n+1} = \mathbf{x}^* - \mathbf{x}^n$$

[Martin et al. 2011]

$$\mathbf{x}^{n+1} = \mathbf{x}^*$$

Optimality Conditions for Implicit Euler

- Applying Newton's method to the optimality conditions leads to the following KKT matrix for each QP sub-problem

$$\begin{bmatrix} \mathbf{M} & \nabla C(\mathbf{x}_i) \\ \nabla C(\mathbf{x}_i)^T & 0 \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \lambda \end{bmatrix} = \begin{bmatrix} \mathbf{M}\tilde{\mathbf{x}} \\ -\mathbf{b} \end{bmatrix}$$

- Eliminate \mathbf{x} to obtain backward Euler update:

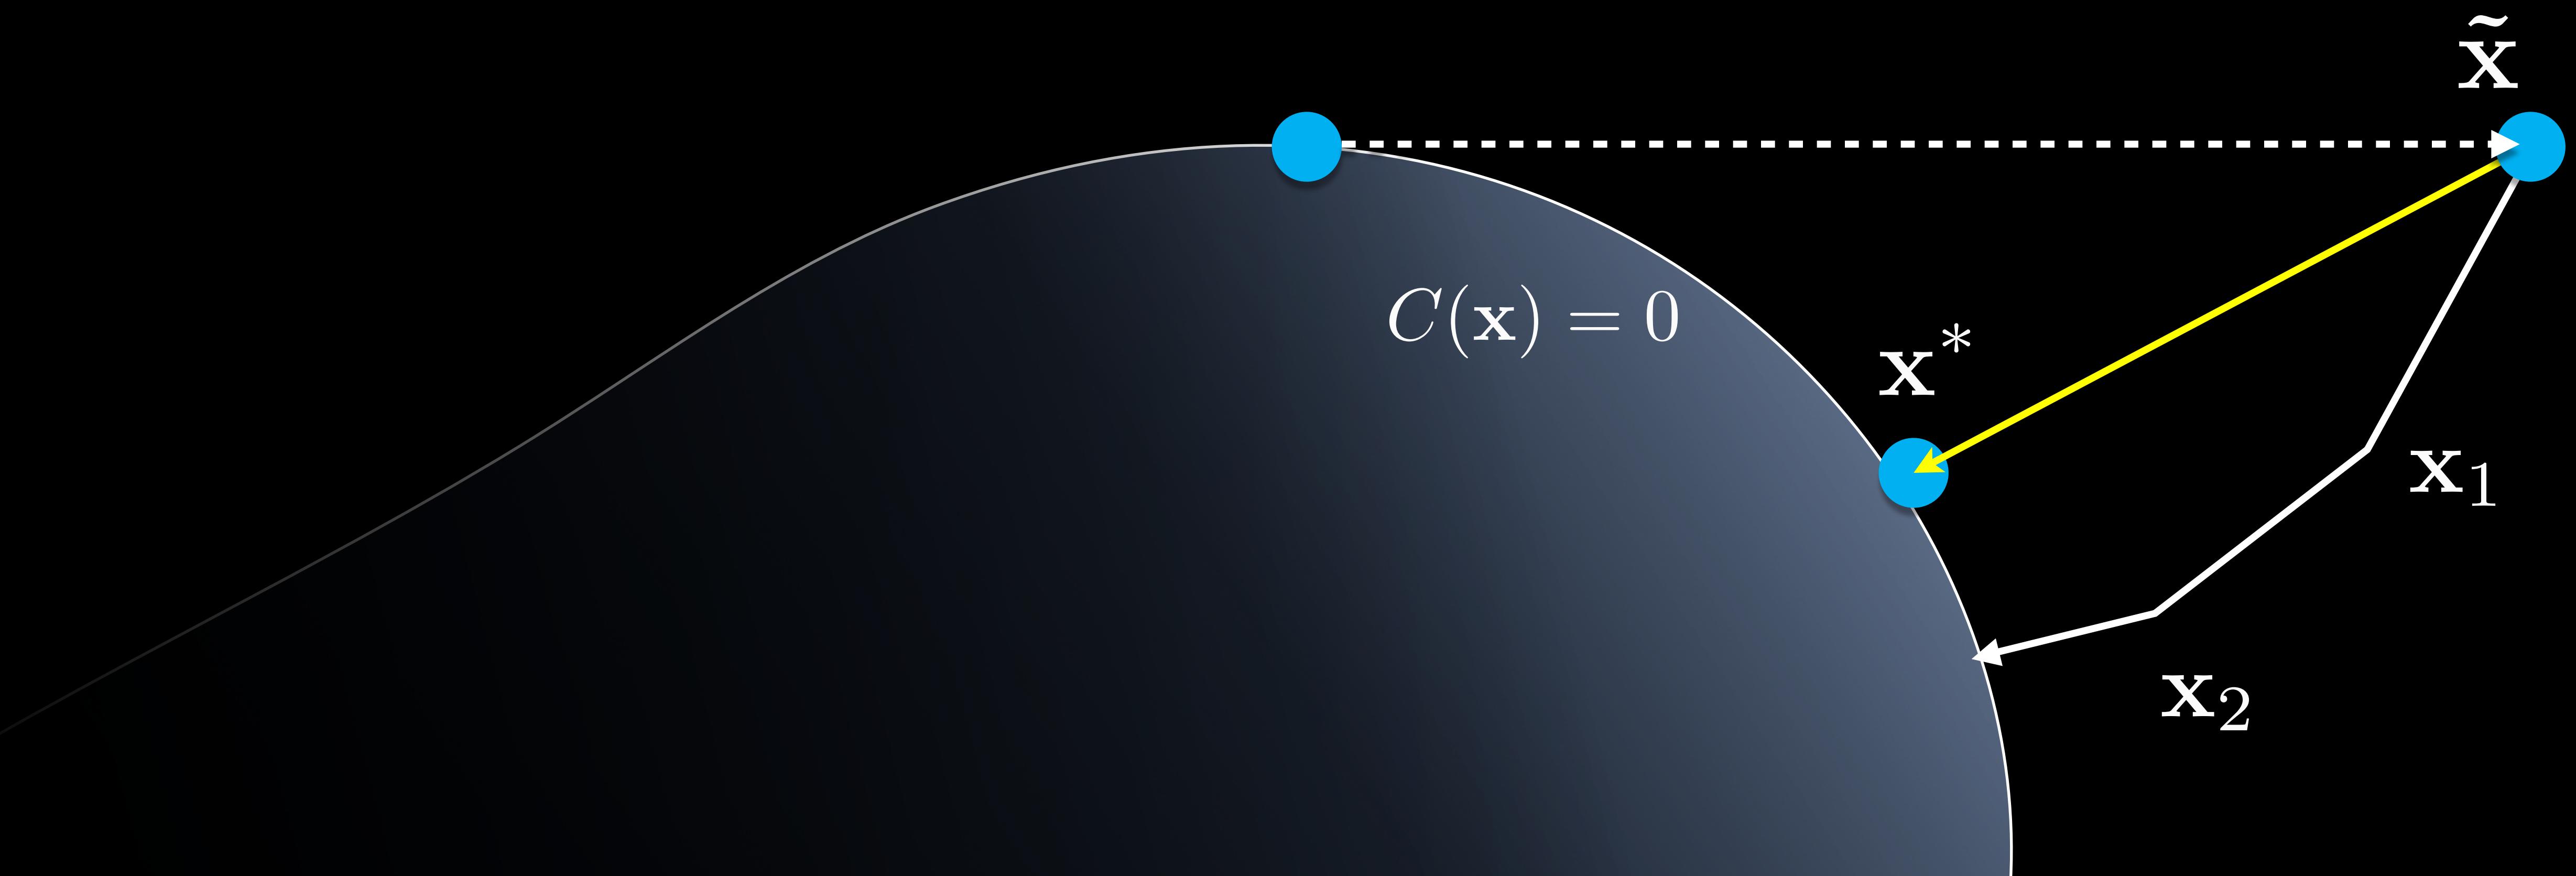
$$[\nabla C(\mathbf{x}_i)^T \mathbf{M}^{-1} \nabla C(\mathbf{x}_i)] \lambda = -\mathbf{C}(\tilde{\mathbf{x}})$$

- The same as PBD? Not quite, different right-hand side:

$$[\nabla C(\mathbf{x}_i)^T \mathbf{M}^{-1} \nabla C(\mathbf{x}_i)] \lambda = -\mathbf{C}(\mathbf{x}_i)$$

PBD and Implicit Euler

- In practice different minimization makes little visual difference
- Identical for linear constraints



Equivalence to Nucleus

- Nucleus [Stam 09]:

$$C(\mathbf{x}^n + \Delta t \mathbf{v}^n + \Delta t \Delta \mathbf{v}) = 0$$

- Position Based Dynamics [Müller et al. 06]

$$C(\tilde{\mathbf{x}} + \Delta \mathbf{x}) = 0 \quad \tilde{\mathbf{x}} = \mathbf{x}^n + \Delta t \mathbf{v}^n \quad \Delta \mathbf{v} = \frac{\Delta \mathbf{x}}{\Delta t}$$

- PBD converts position changes to impulses applied at the beginning of the time-step

Parallel Position Based Dynamics

- At each iteration we need to solve the following system:

$$[\nabla C(\mathbf{x}_i)^T \mathbf{M}^{-1} \nabla C(\mathbf{x}_i)] \lambda = -\mathbf{C}(\mathbf{x}_i)$$



- Position Based Dynamics (PBD) is typically serial
- Use Gauss-Jacobi for parallelism and simple handling of inequalities
- Problem: system matrix can be indefinite, Jacobi will not converge, e.g.: for redundant constraints (cf. figure)

Constraint Averaging

- Regularized Jacobi iteration via averaging [Bridson et al. 02]
- Sum all constraint deltas together and divide by constraint count for that particle

$$\mathbf{x}_i \leftarrow \mathbf{x}_i + \frac{1}{n_i} \sum_{n_i} \lambda_j \nabla C_j$$

- Successive-over relaxation by user parameter omega [0,2]:

$$\mathbf{x}_i \leftarrow \mathbf{x}_i + \frac{\omega}{n_i} \sum_{n_i} \lambda_j \nabla C_j$$

Constraint Solving on the GPU

- Two ways to solve constraints:

Particle-centric approach
(gather)

```
foreach particle (in parallel)
{
    foreach constraint
    {
        calculate constraint error
        update delta
    }
}
```

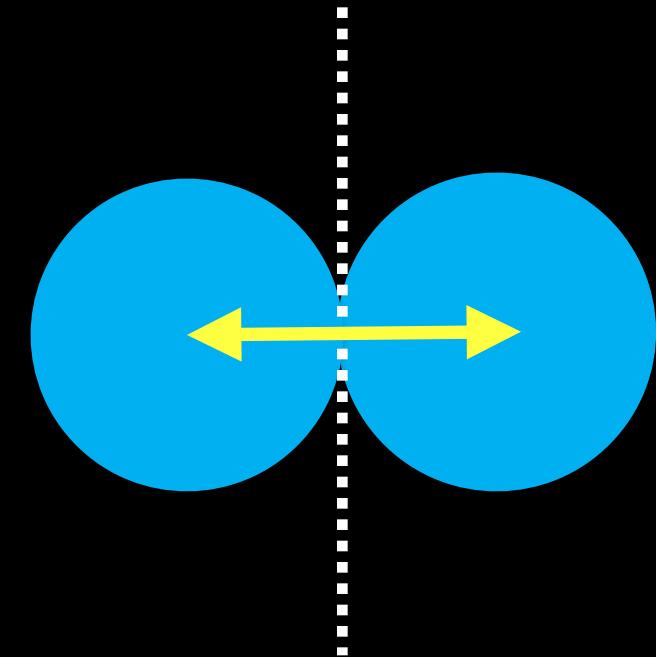
Constraint-centric approach
(scatter)

```
foreach constraint (in parallel)
{
    calculate constraint error
    foreach particle
    {
        update delta (atomically)
    }
}
```

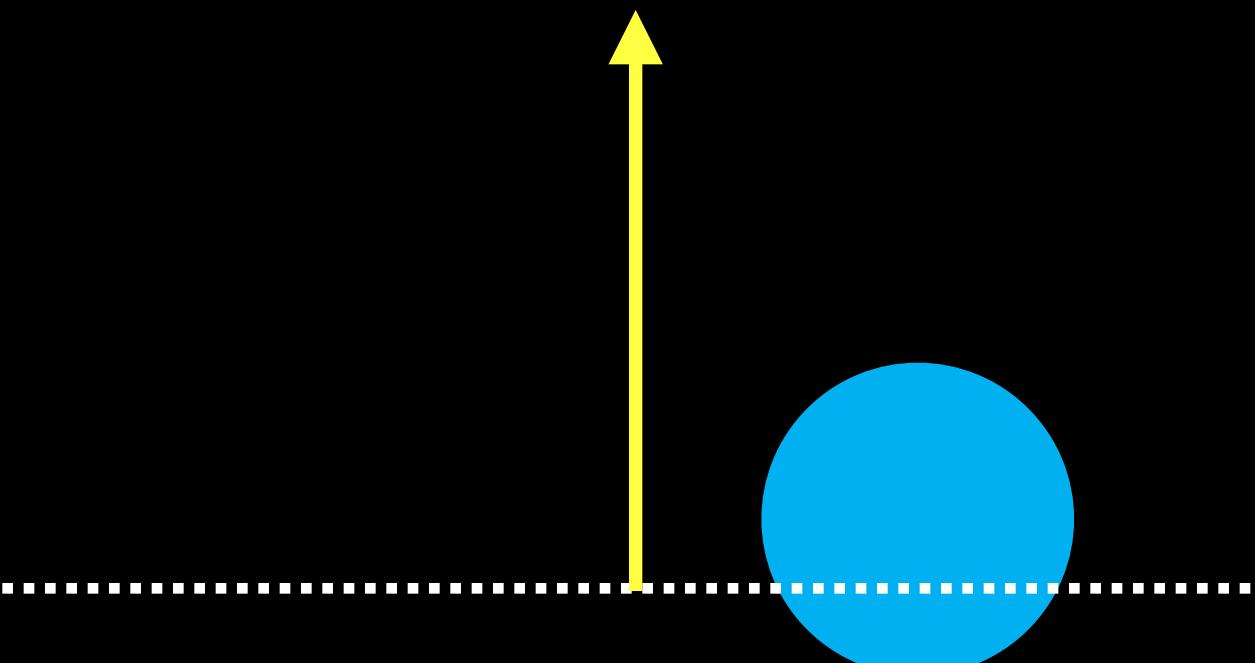
Contact and friction

Collision Detection

- All dynamics represented as particles
- Kinematic objects represented as meshes
- Two types of collision detection:
 - ▶ Particle-Particle
 - ▶ Particle-Mesh



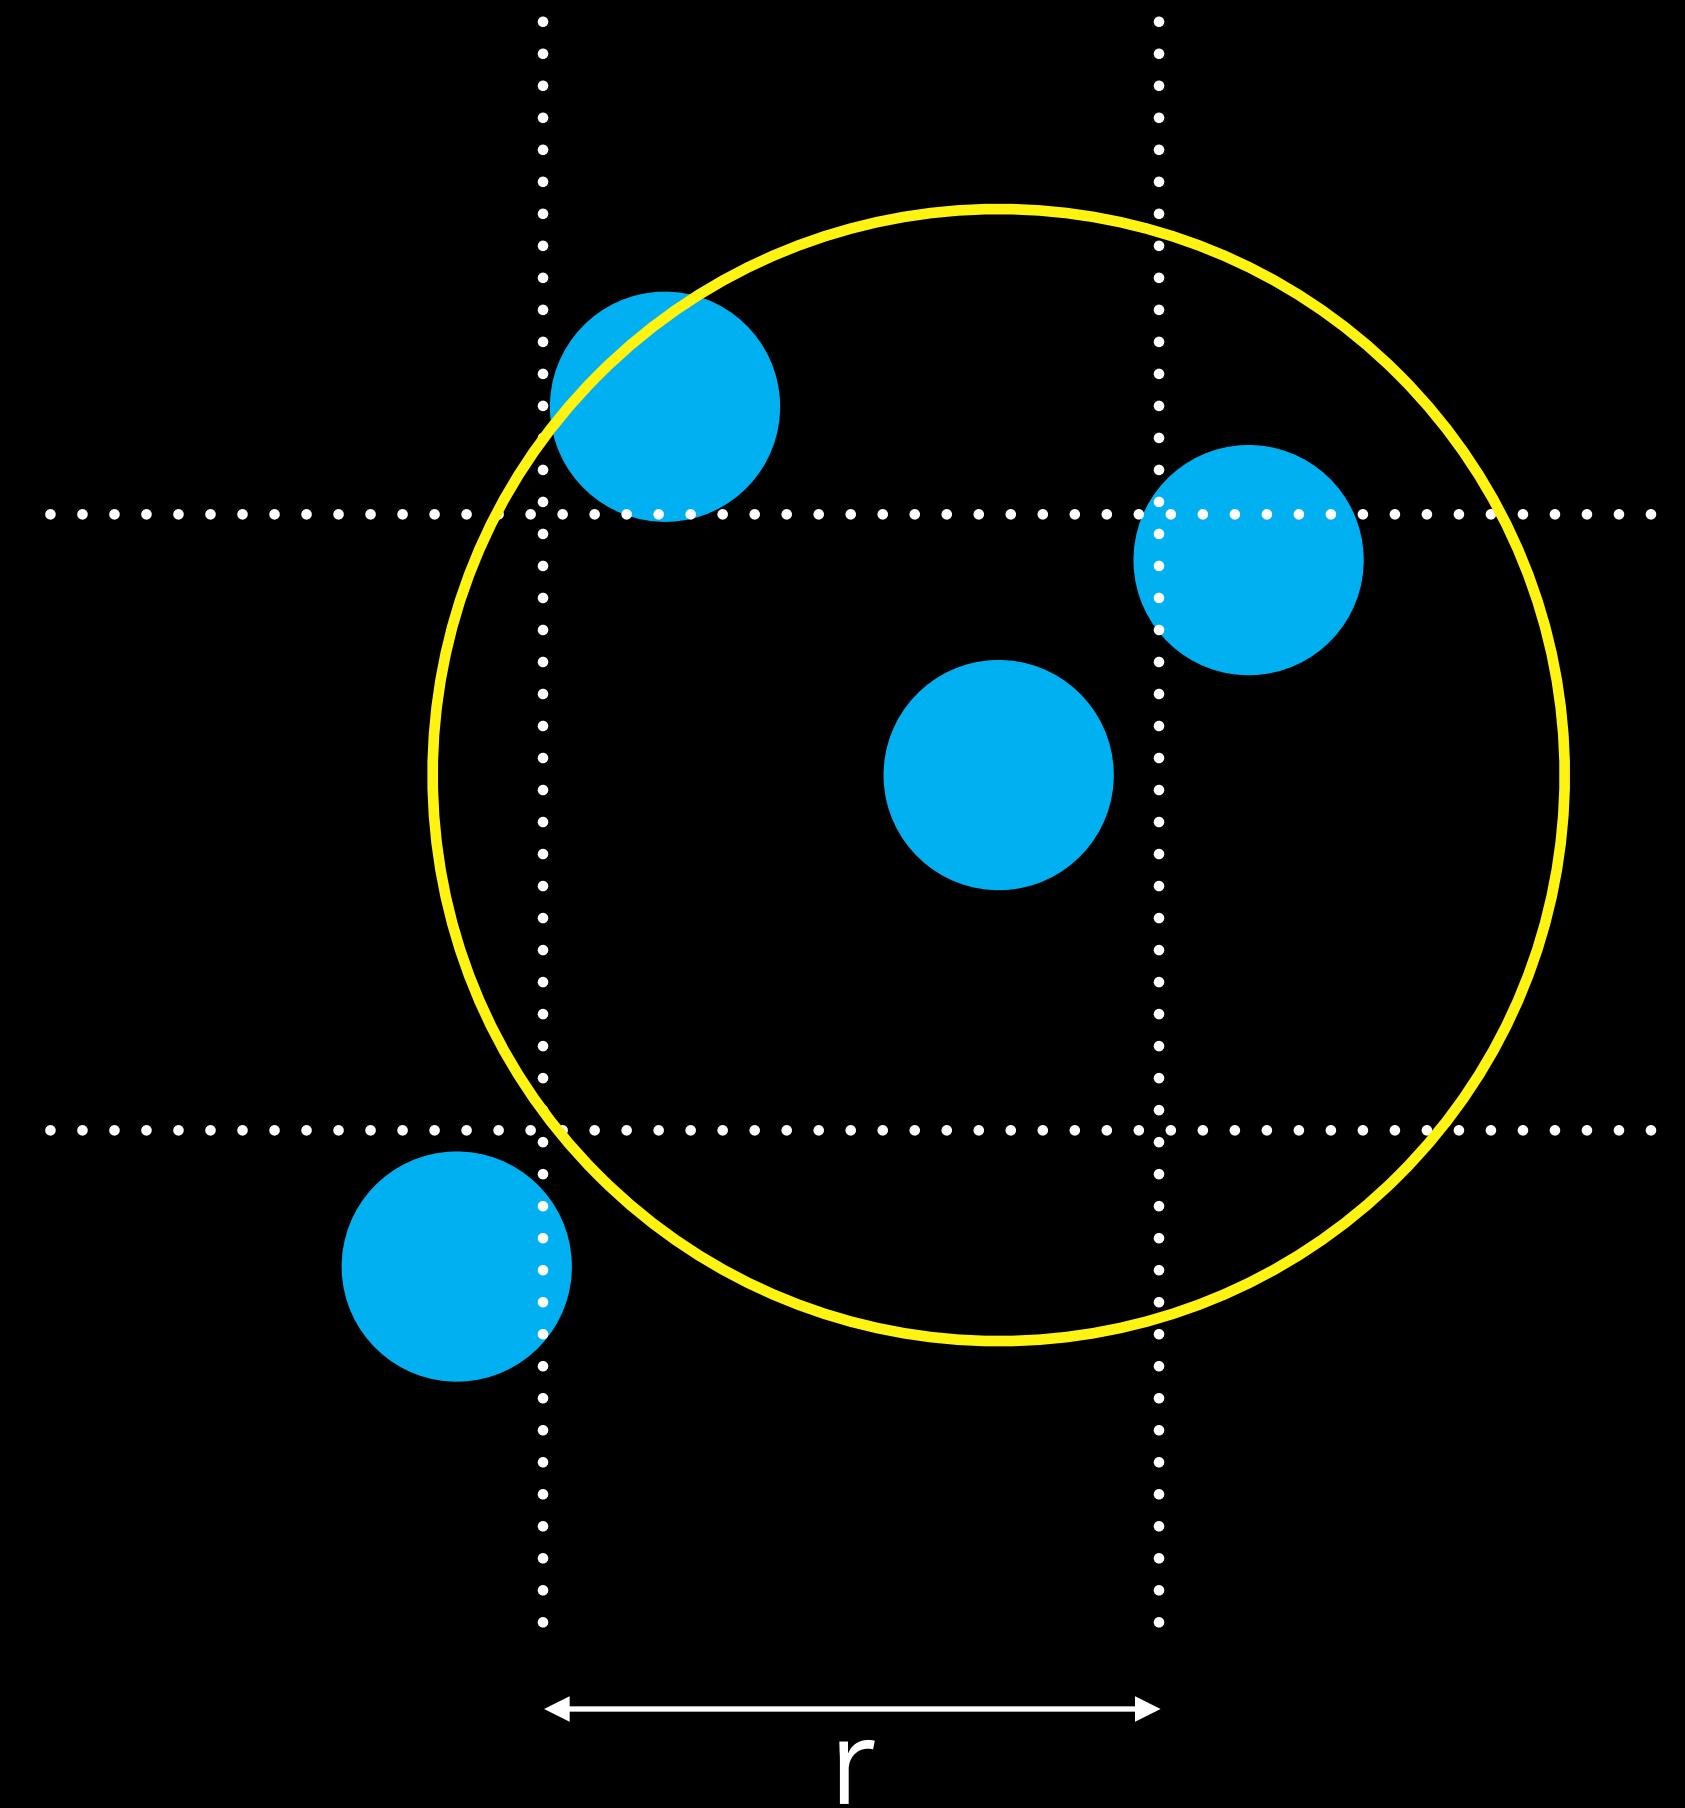
$$C_{contact} = |\mathbf{x}_i - \mathbf{x}_j| - 2r \geq 0$$



$$C_{contact} = \mathbf{n} \cdot \mathbf{x} - r \geq 0$$

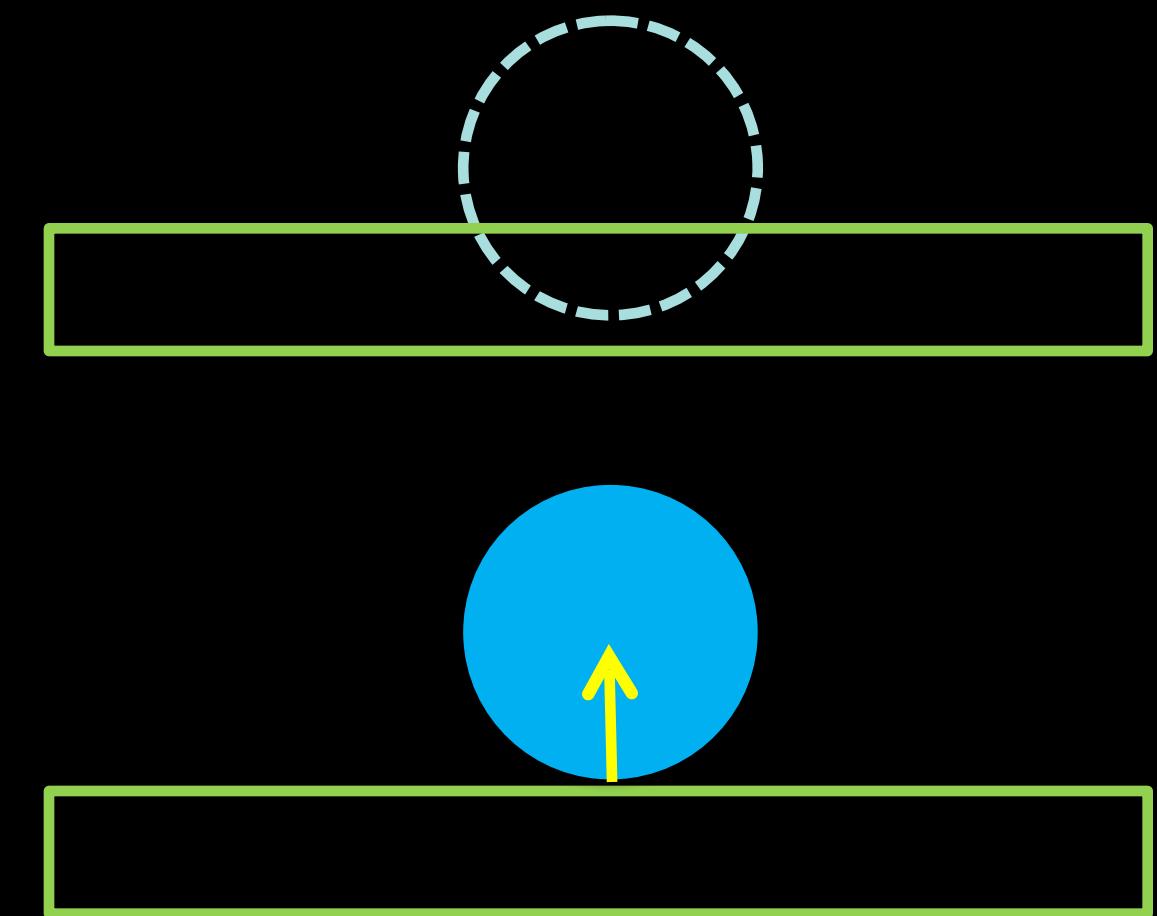
Collision Detection

- Particle-Particle
 - ▶ Tiled uniform grid
 - ▶ Fixed maximum radius
 - ▶ Built using `cub::DeviceRadixSort`
 - ▶ Re-order particle data according to cell index to improve memory locality



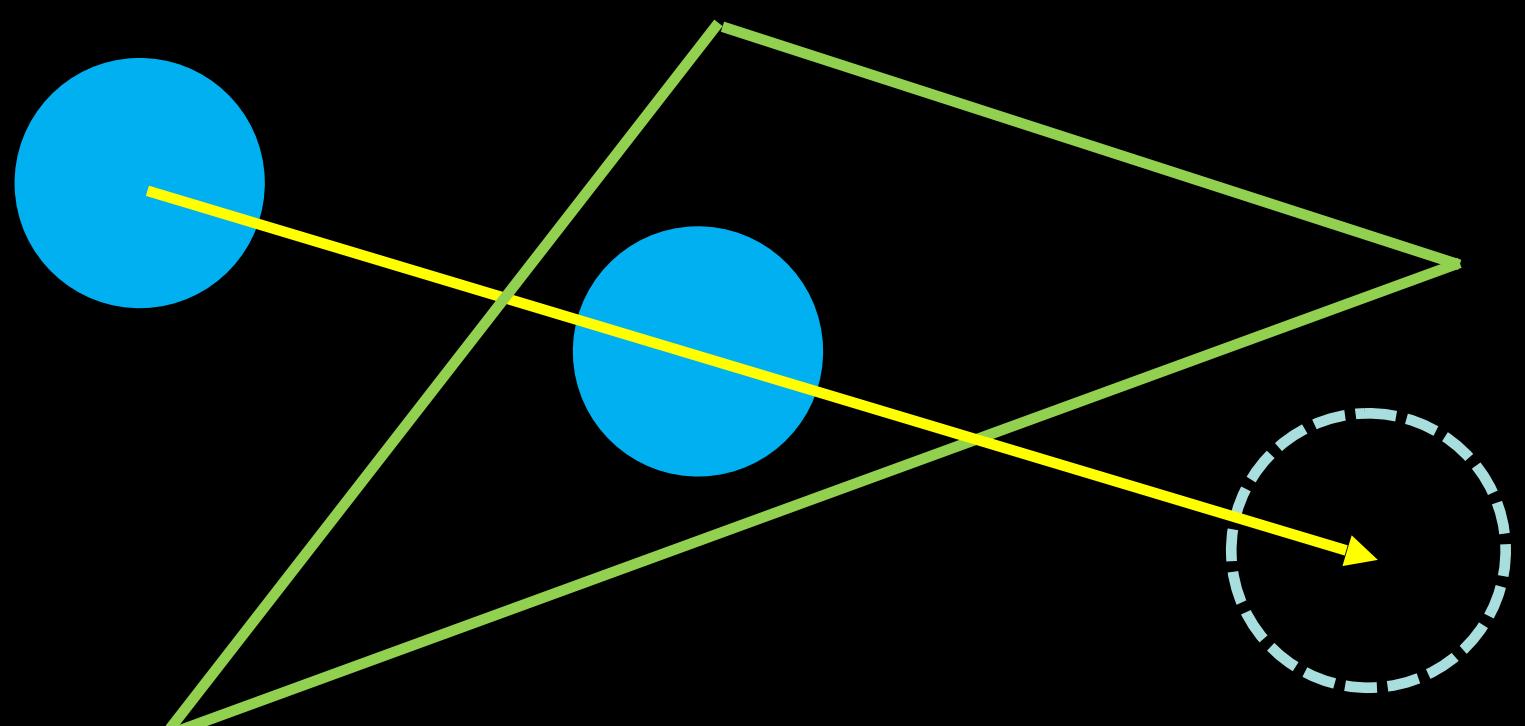
Collision Detection

- Particle-Convex
 - ▶ 2D hash-grid
 - ▶ Built on GPU
 - ▶ 1 warp-per shape, rasterizes projected bounds to grid (~1500 shapes / ms)



Convex Collision (MTD)

- Particle-Triangle Mesh
 - ▶ 3D hash-grid
 - ▶ Rasterized in CUDA
 - ▶ Lollipop test (CCD)

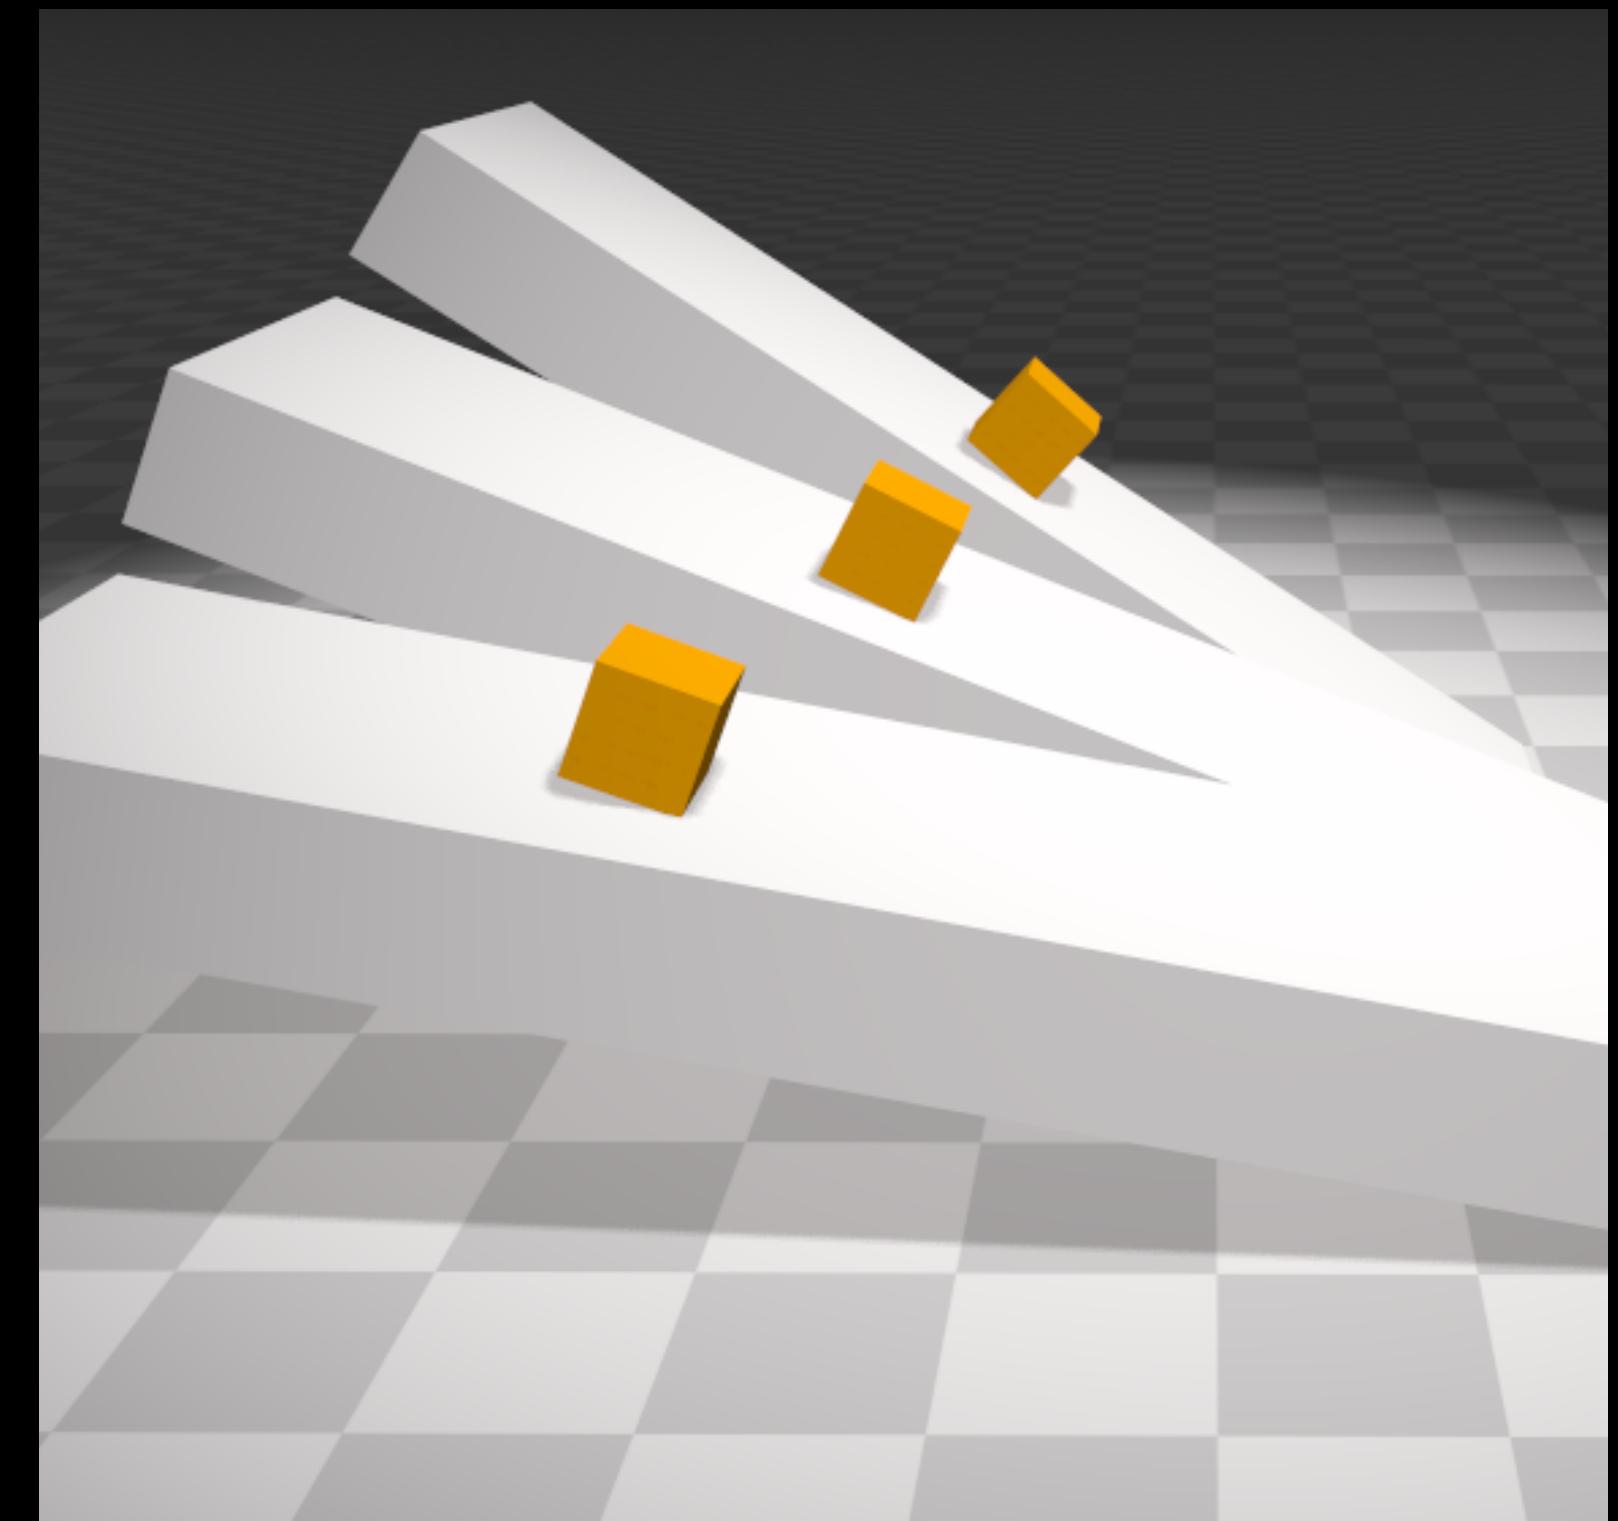


Triangle Collision (TOI)

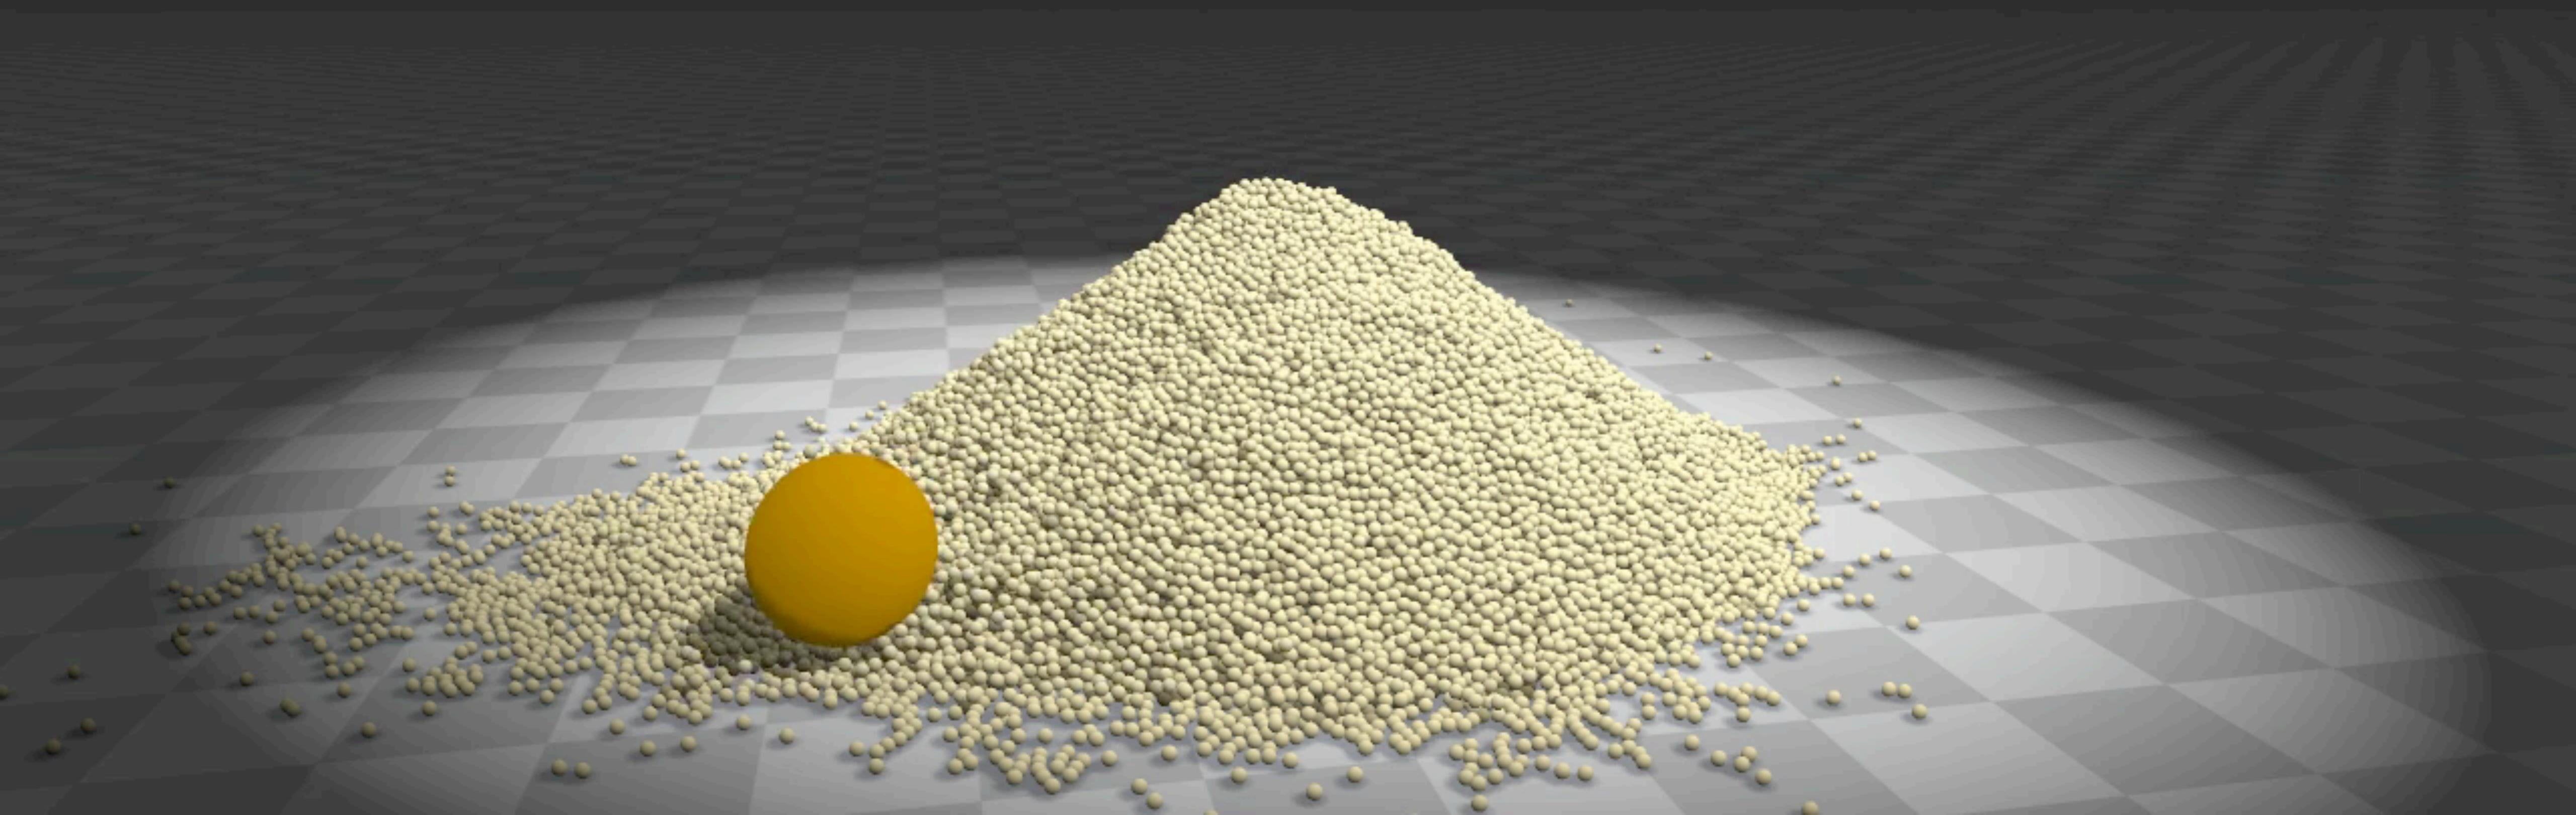
Friction

- Friction in PBD traditionally applied using a velocity filter
- We introduce a position-level **frictional constraint**

$$C_{friction} = |(\mathbf{x} - \mathbf{x}_0)_{\perp} \mathbf{n}|$$



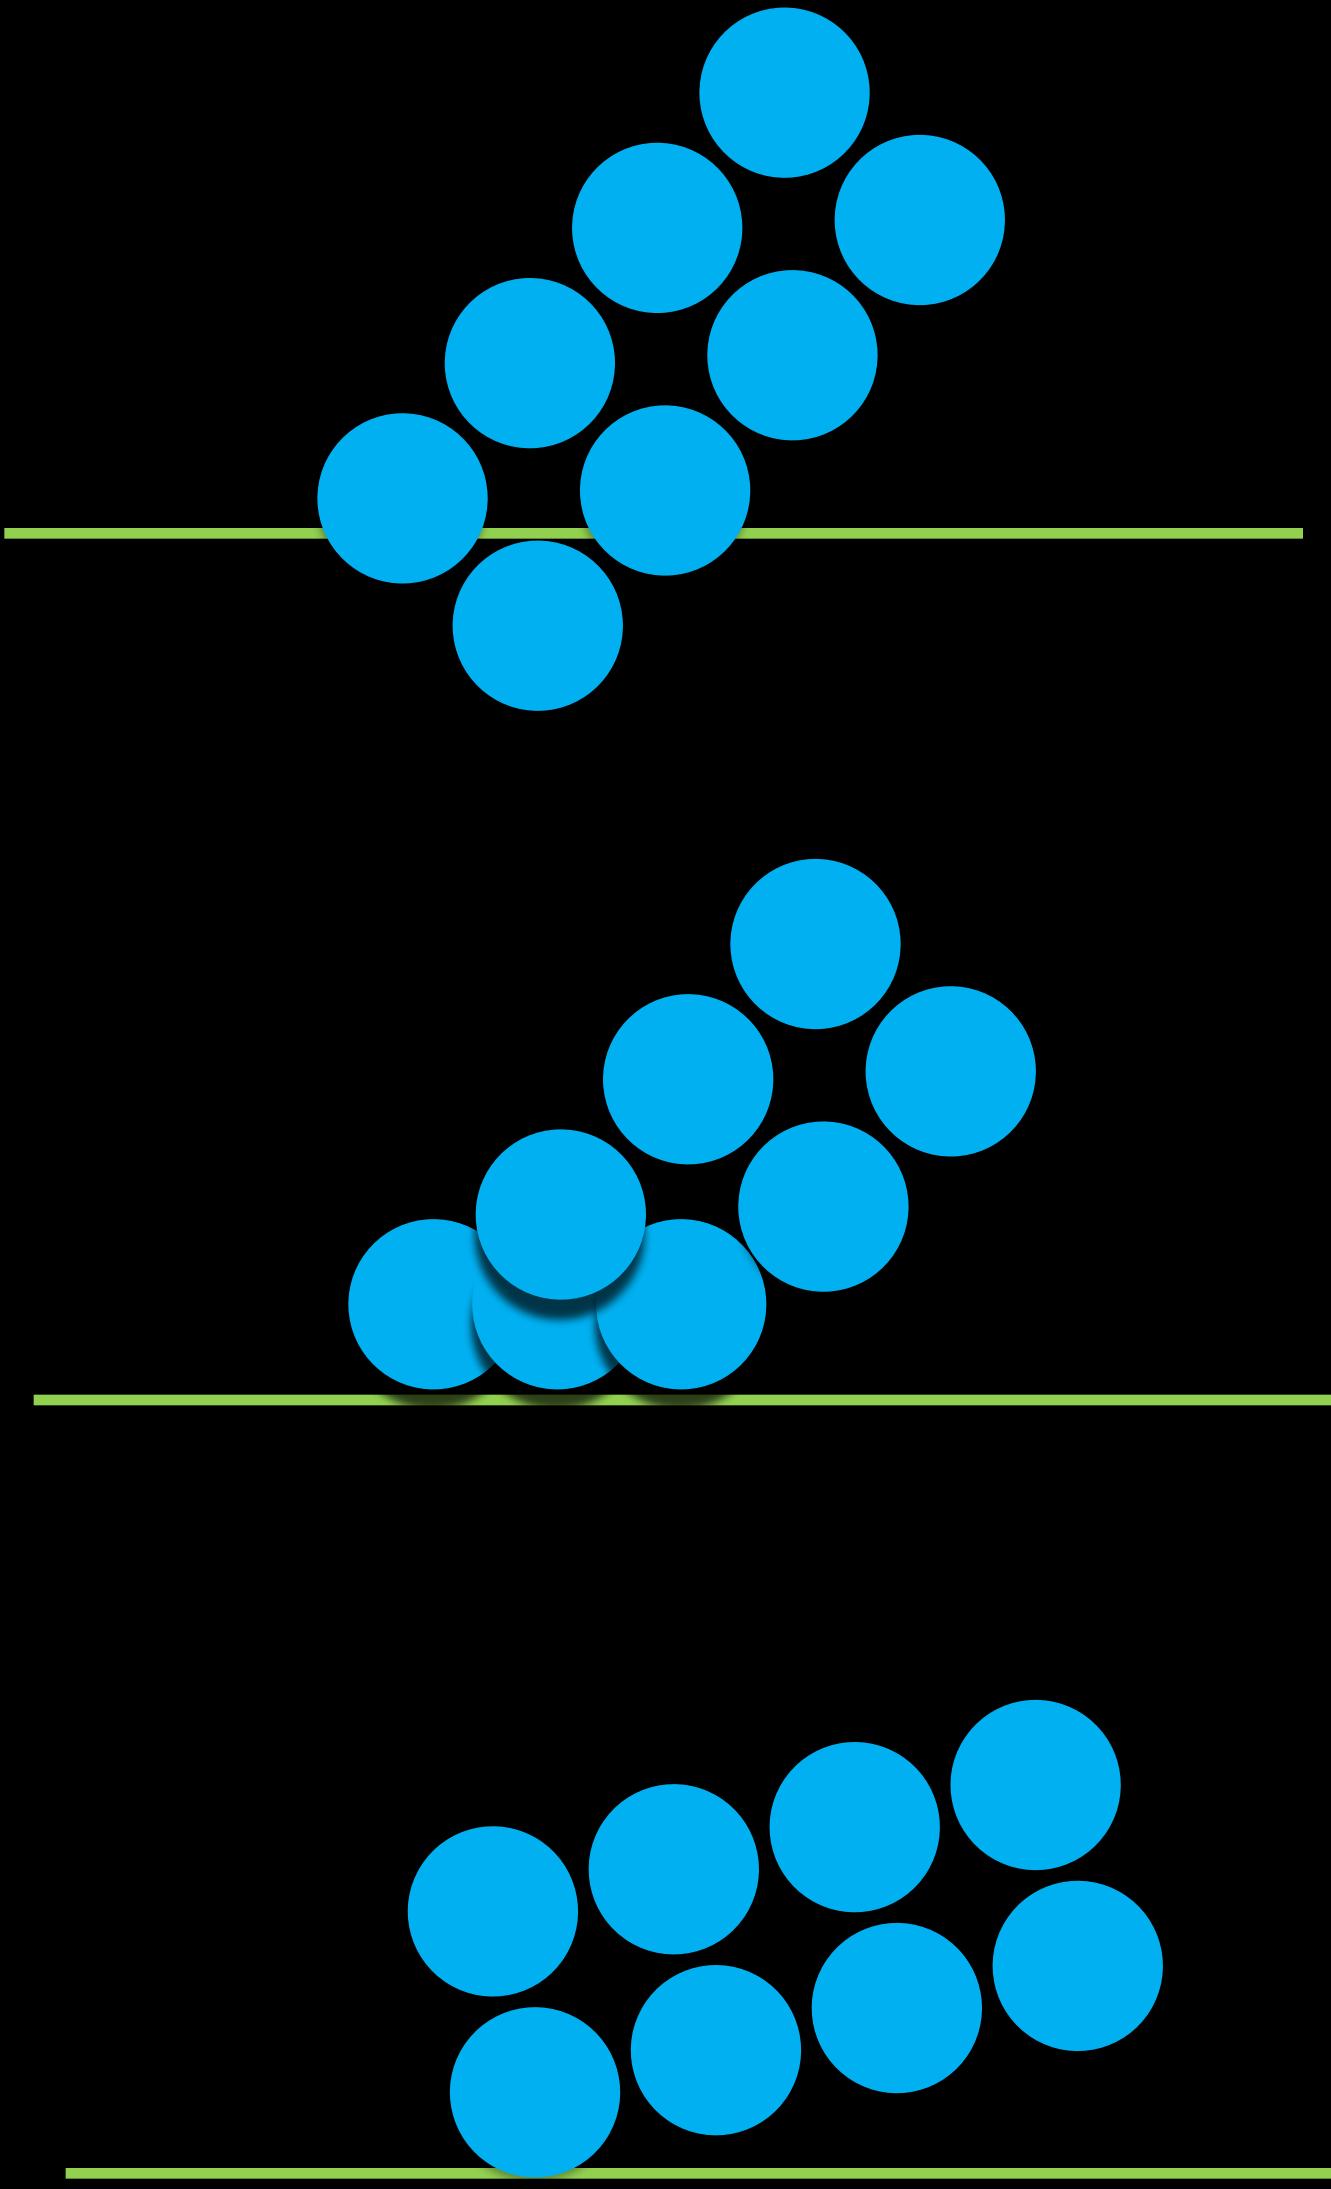
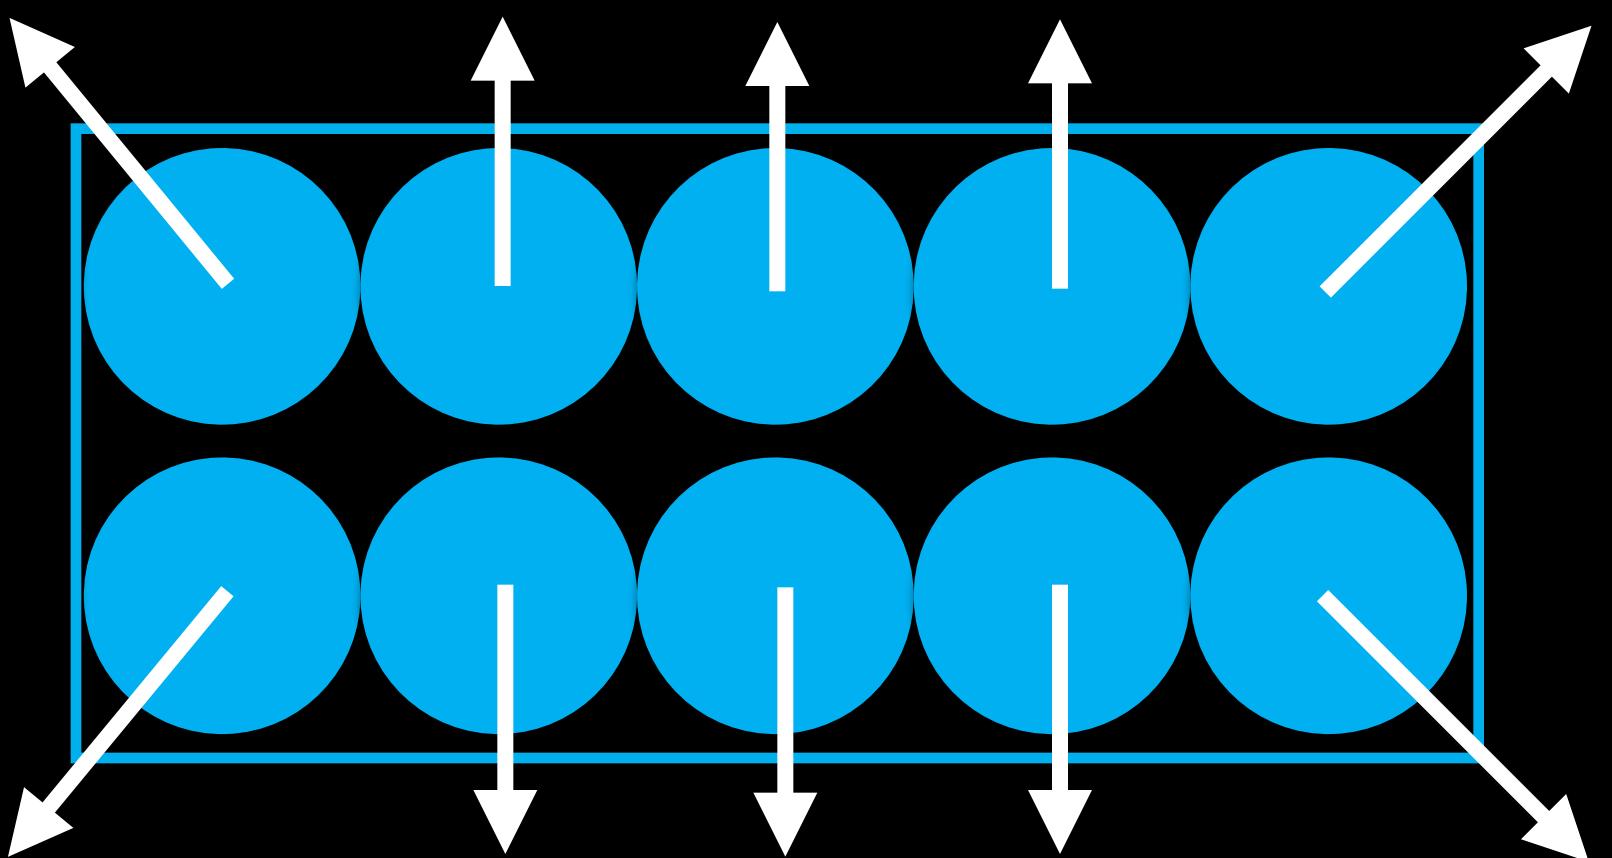
- Approximate Coulomb friction using penetration depth to limit lambda



Rigid Bodies

Rigid Bodies

- Convert mesh->SDF
- Place particles in interior
- Add **shape-matching** constraint
- Store SDF dist + gradient on particles:



Shape matching on the GPU

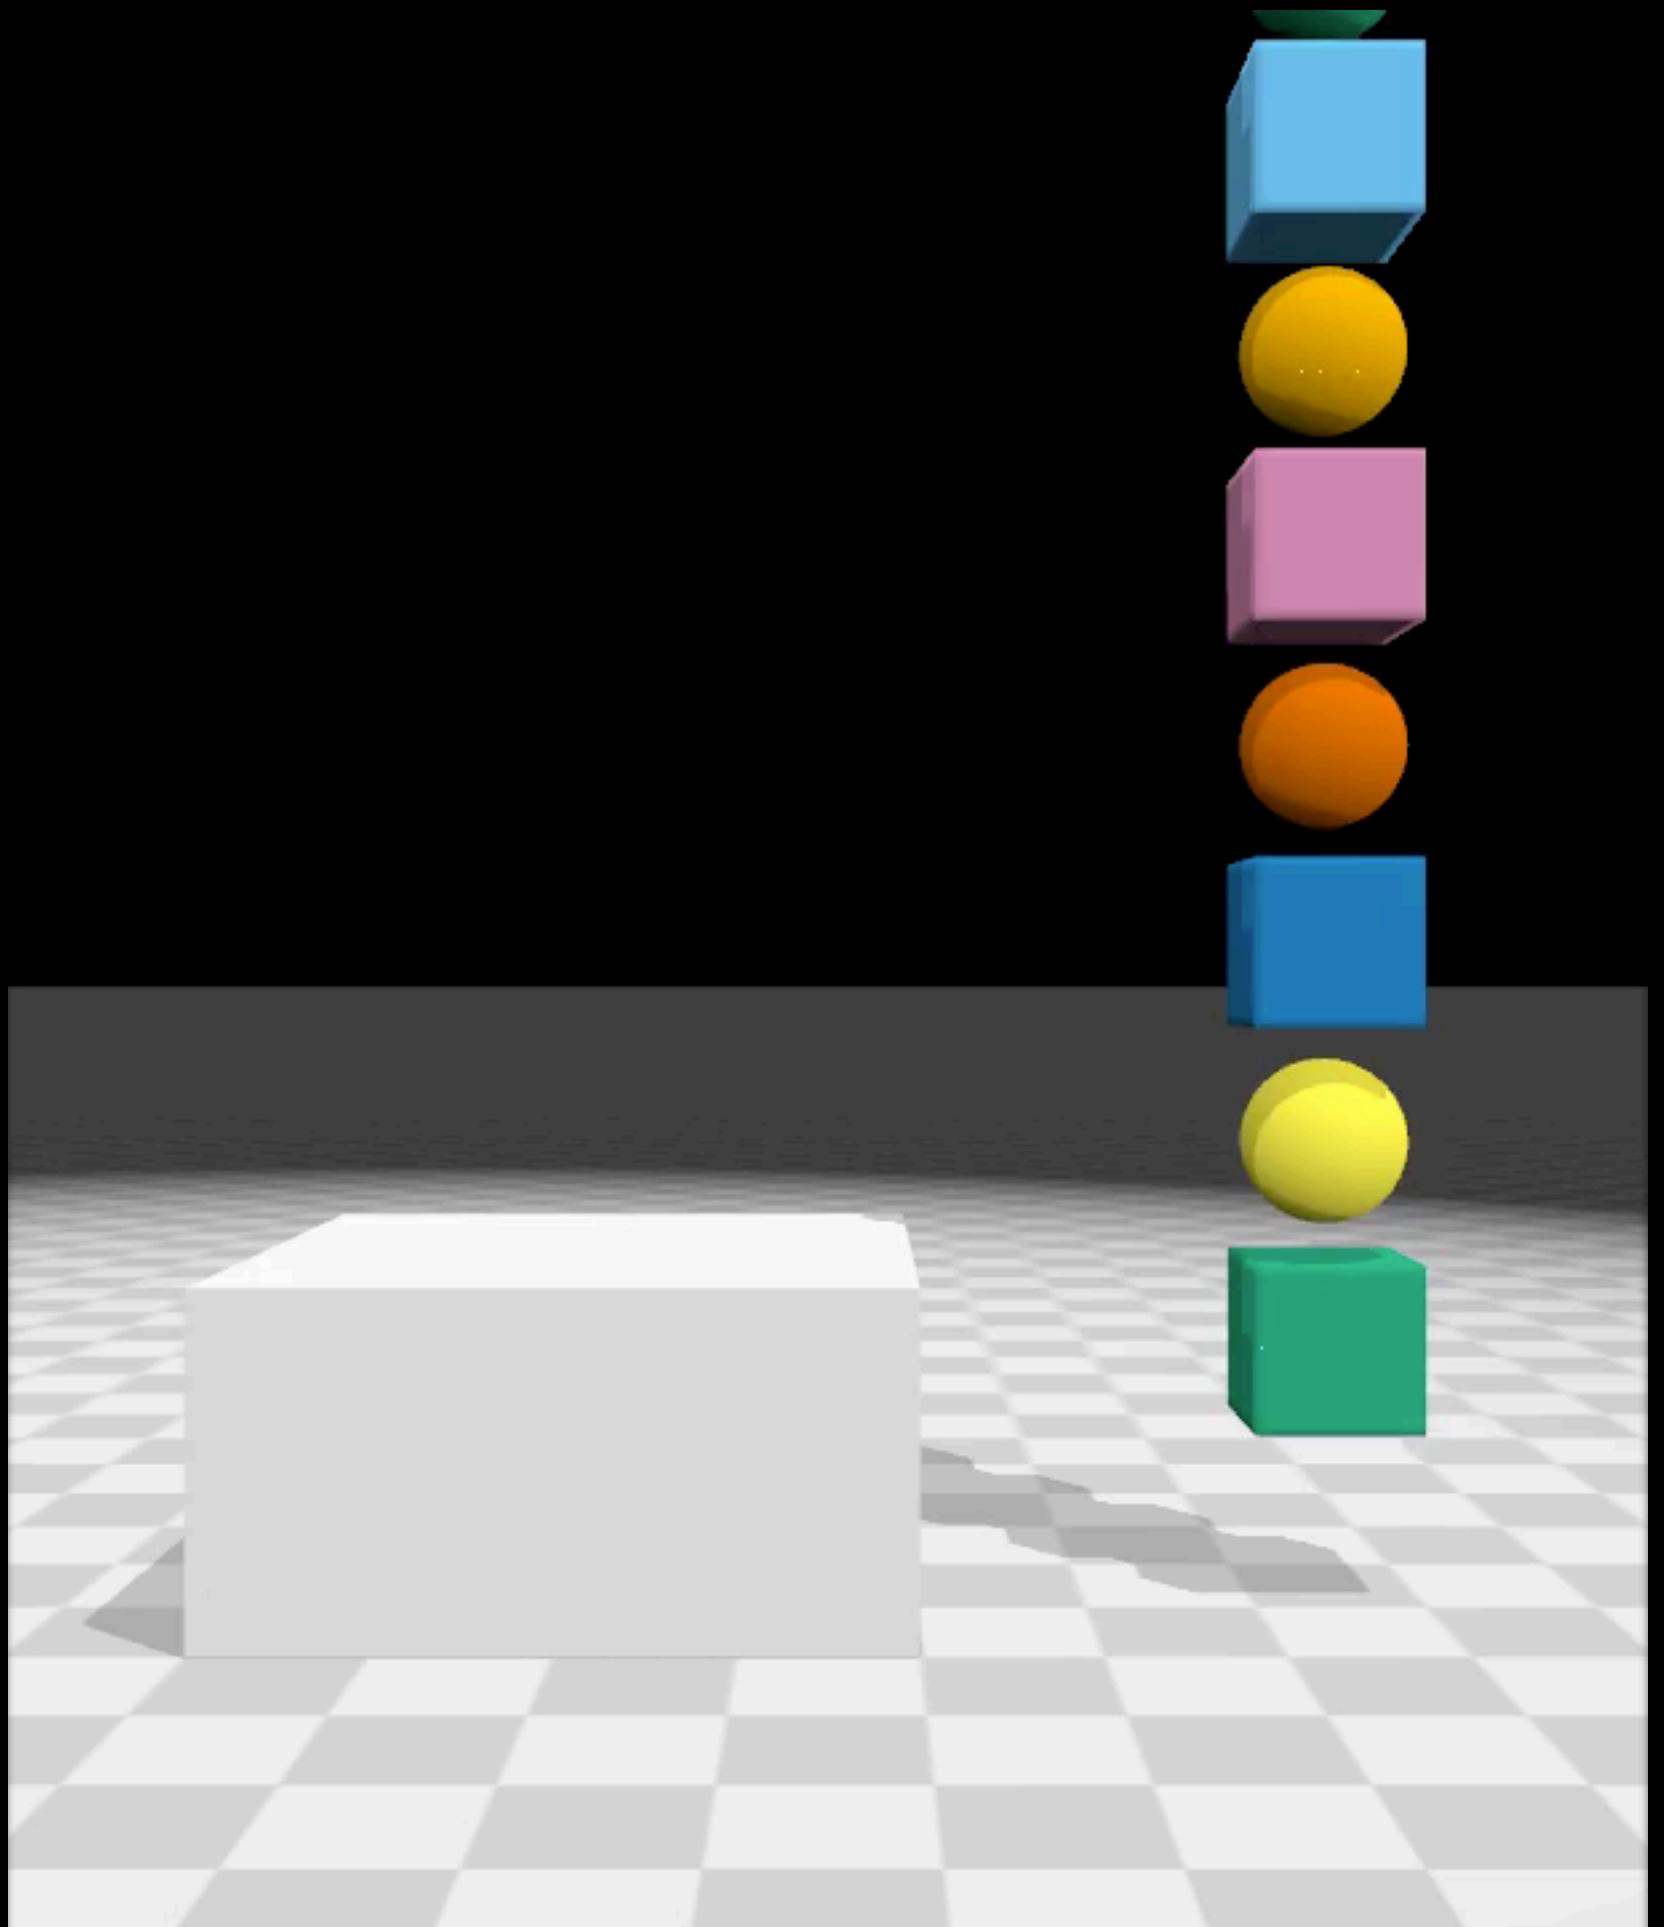
- Shape matching requires computing centre of mass and the moment matrix for particles:

$$\mathbf{c} = \sum_i m_i \mathbf{x}_i / \sum_i m_i \quad \mathbf{A} = \sum_i m_i (\mathbf{x}_i - \mathbf{c})(\bar{\mathbf{x}}_i - \bar{\mathbf{c}})^T$$

- Large summations, not immediately parallel friendly
- Optimized using two parallel cub::BlockReduce calls
- O(N) -> O(logN) (18ms -> 0.6ms)
- 1 block per-rigid shape (64 threads, heuristic, irregular workload problem)
- Polar decomposition still single threaded

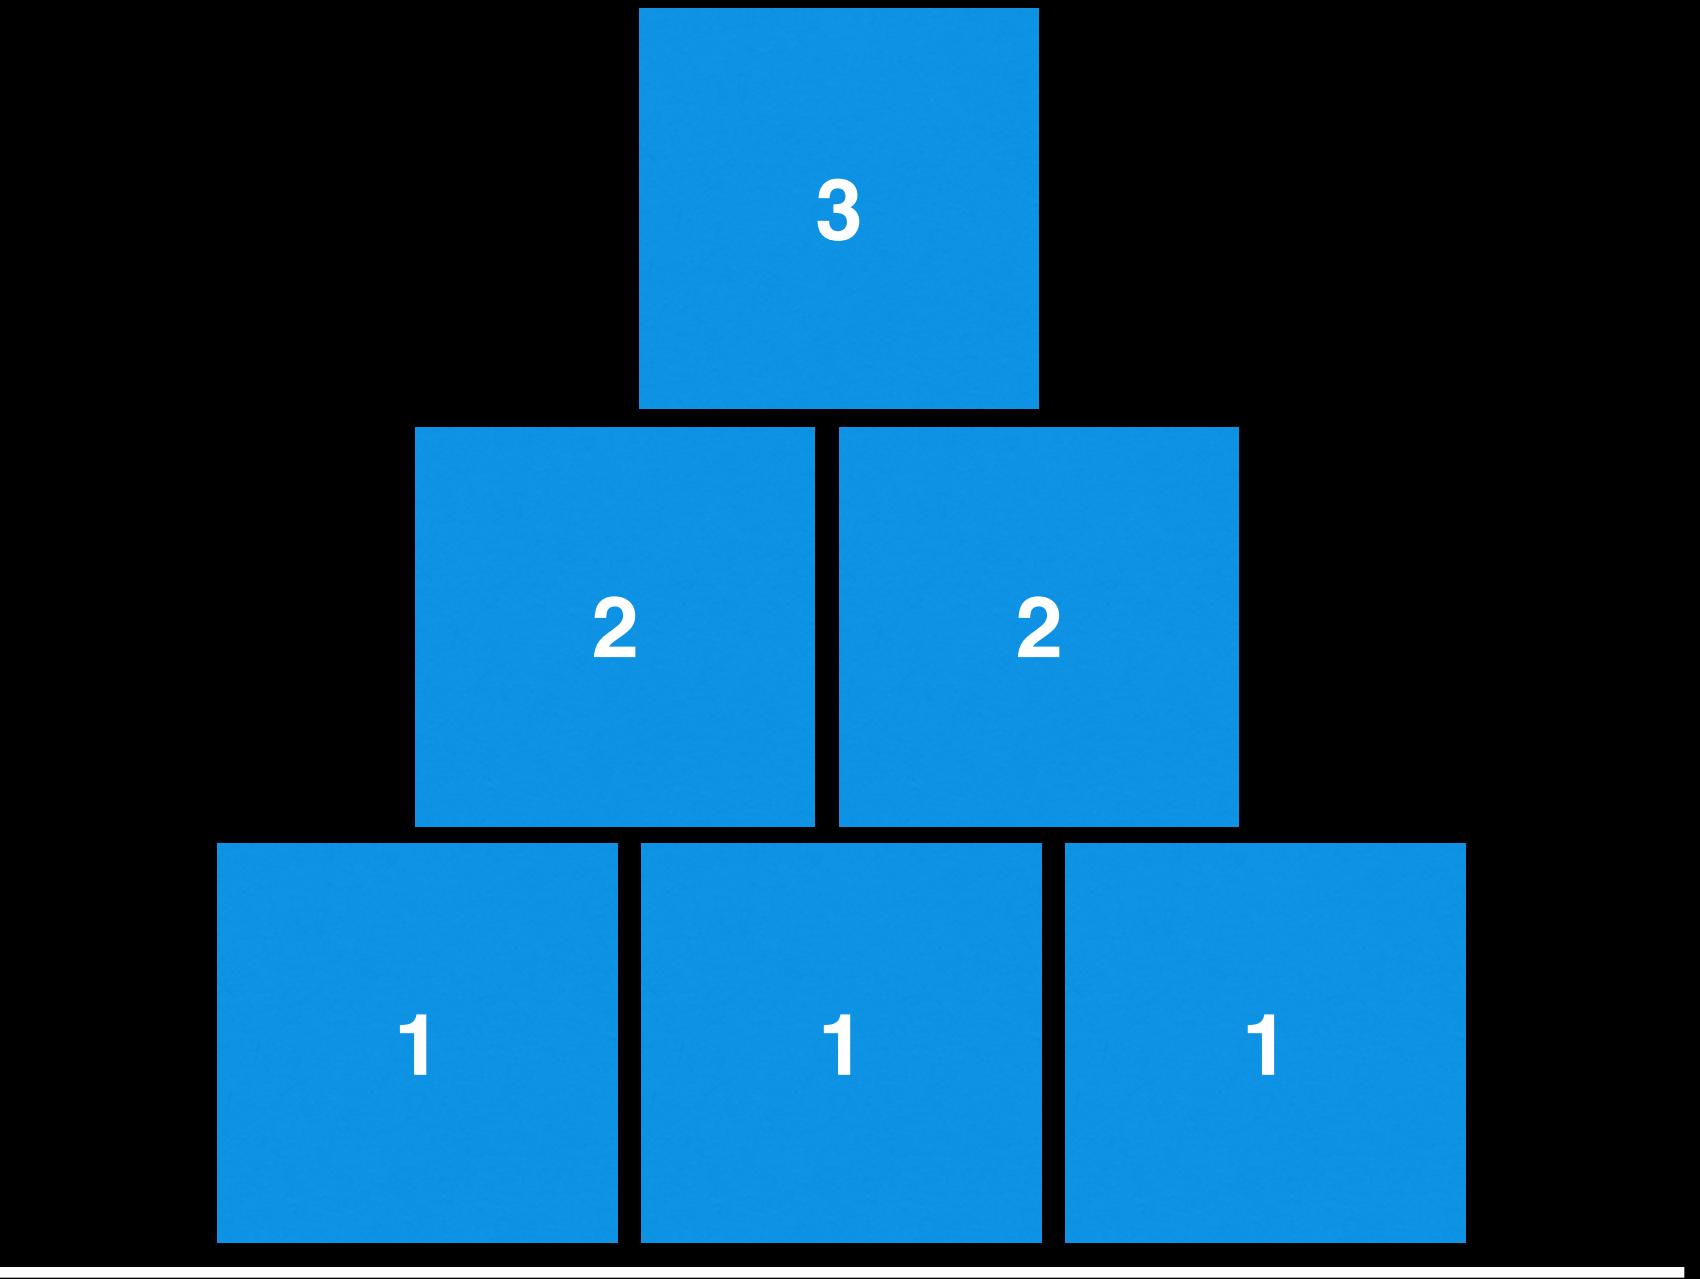
Plastic Deformation

- Detect when deformation exceeds a threshold
- Simply change rest-configuration of particles
- Adjust visual mesh (linear skinning)



Rigid Bodies - Piles and Stacks

- Piles of objects can take many iterations to appear stiff
- Common solution: **shock propagation** [Guendelman et al. 03]
 - ▶ Re-orders constraint solve bottom->top
 - ▶ Sets mass of each layer = ∞
 - ▶ Problem: limited parallelism

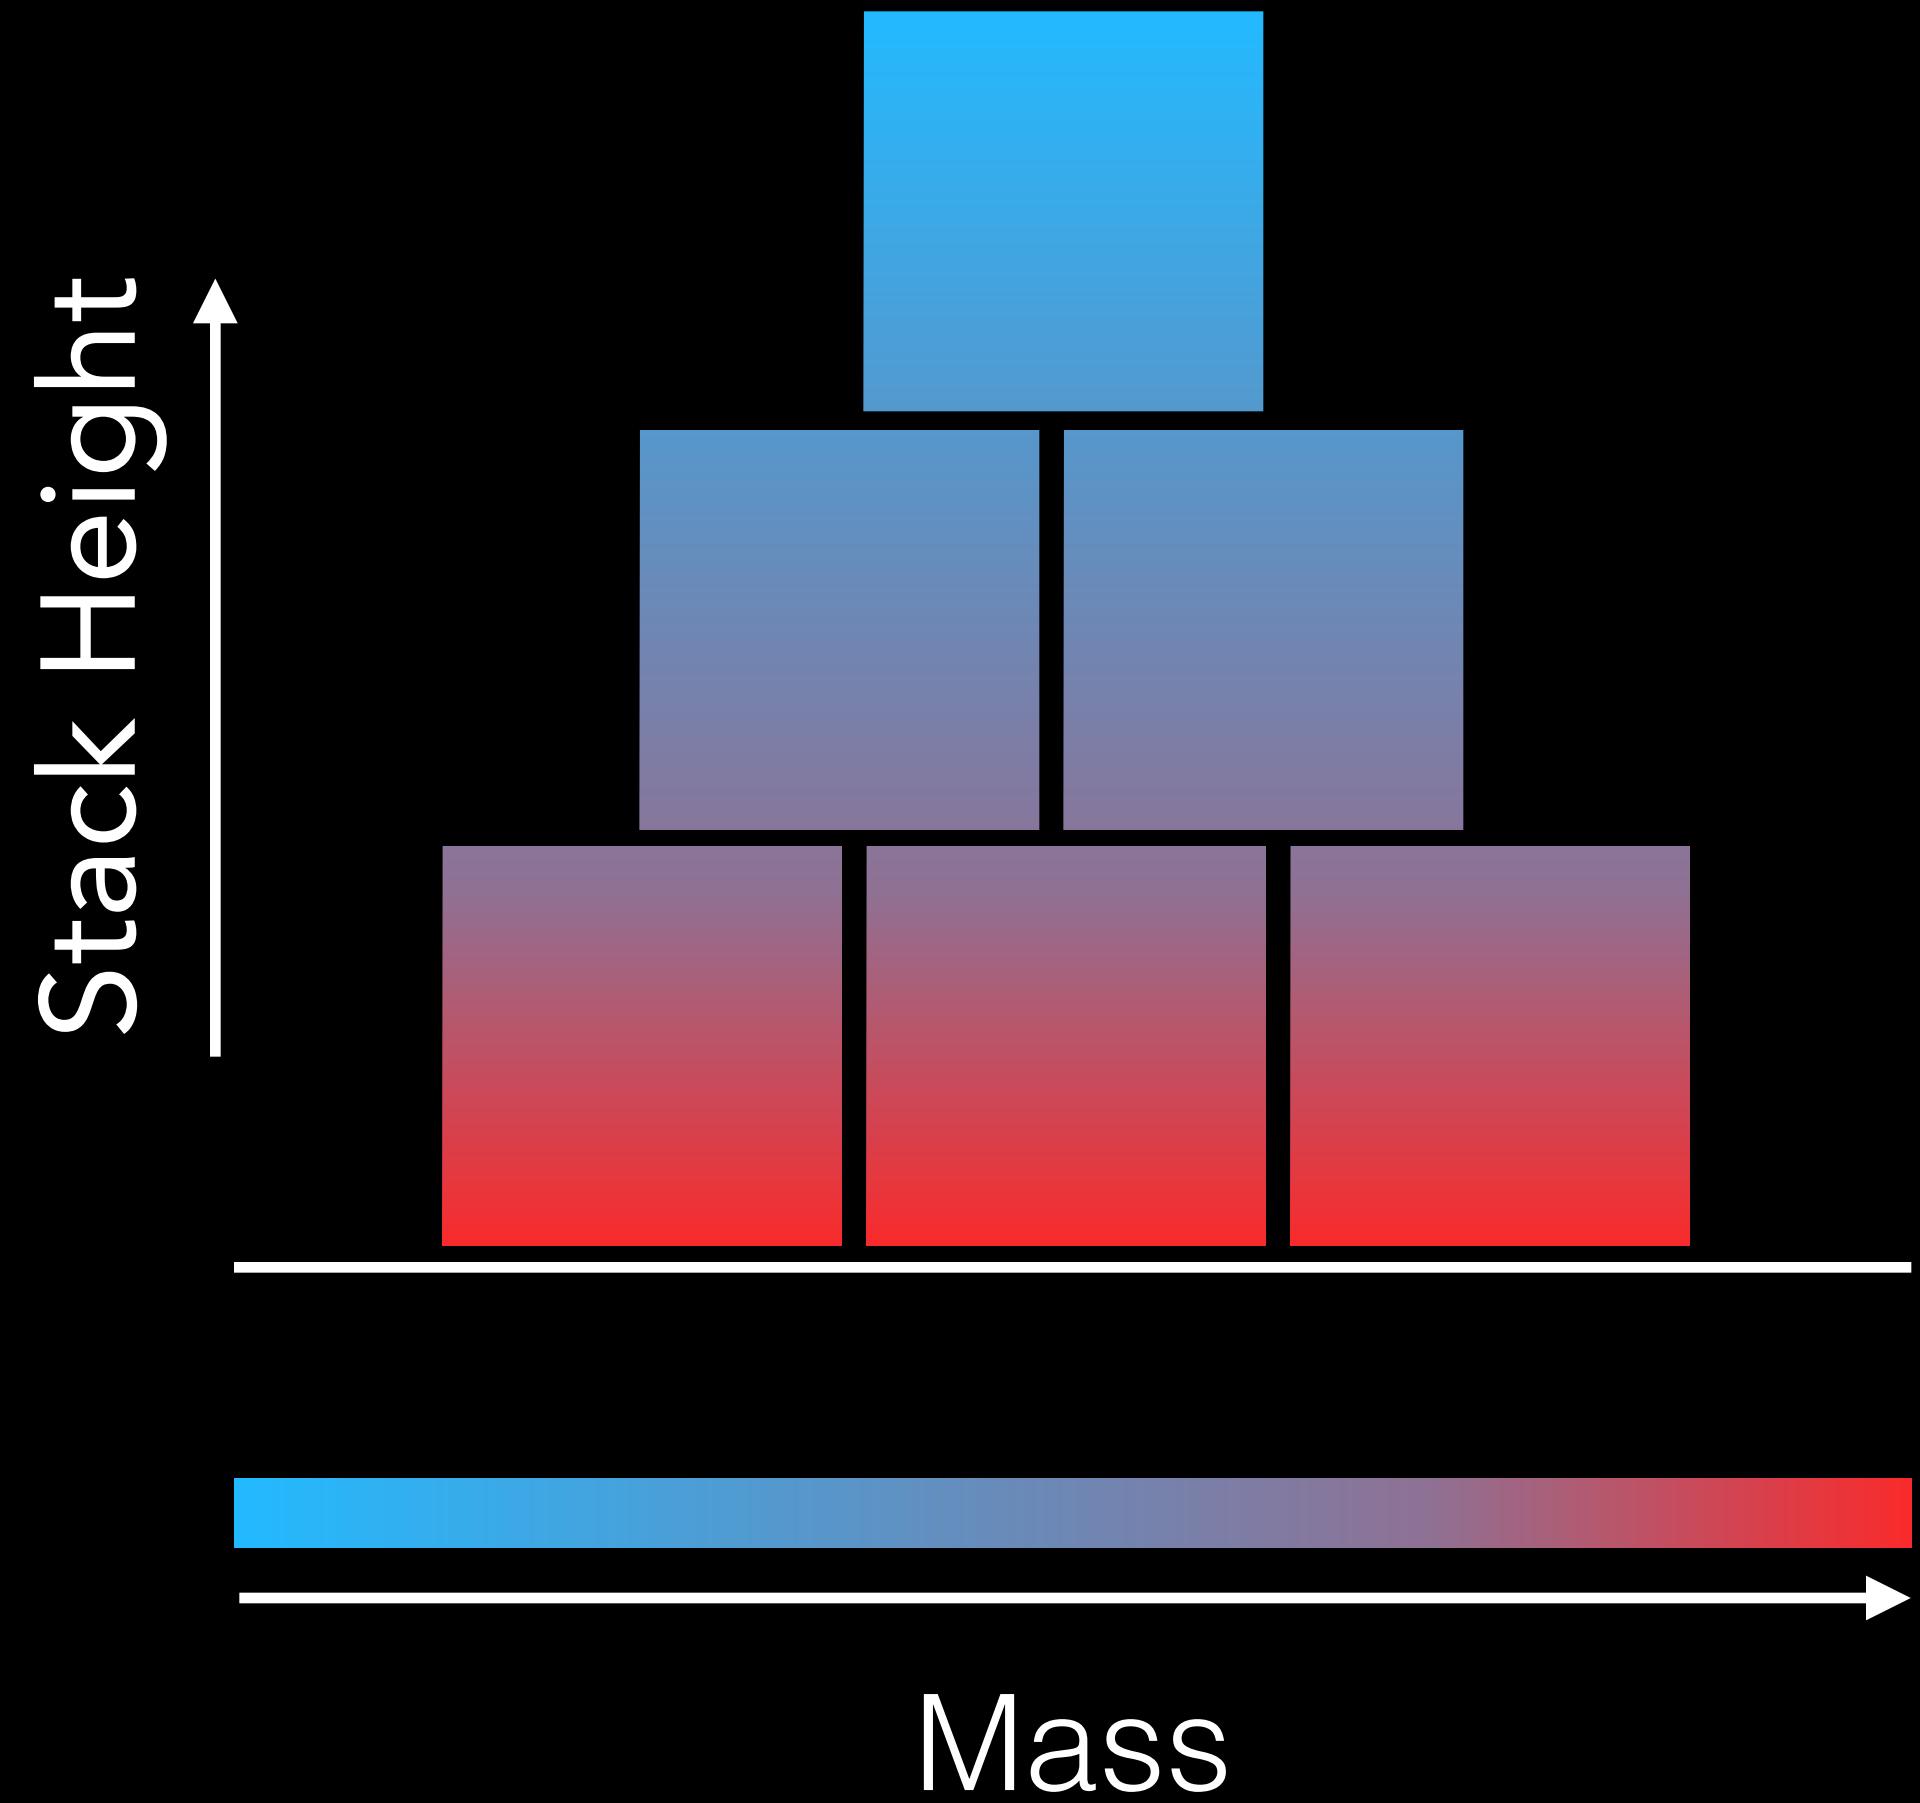


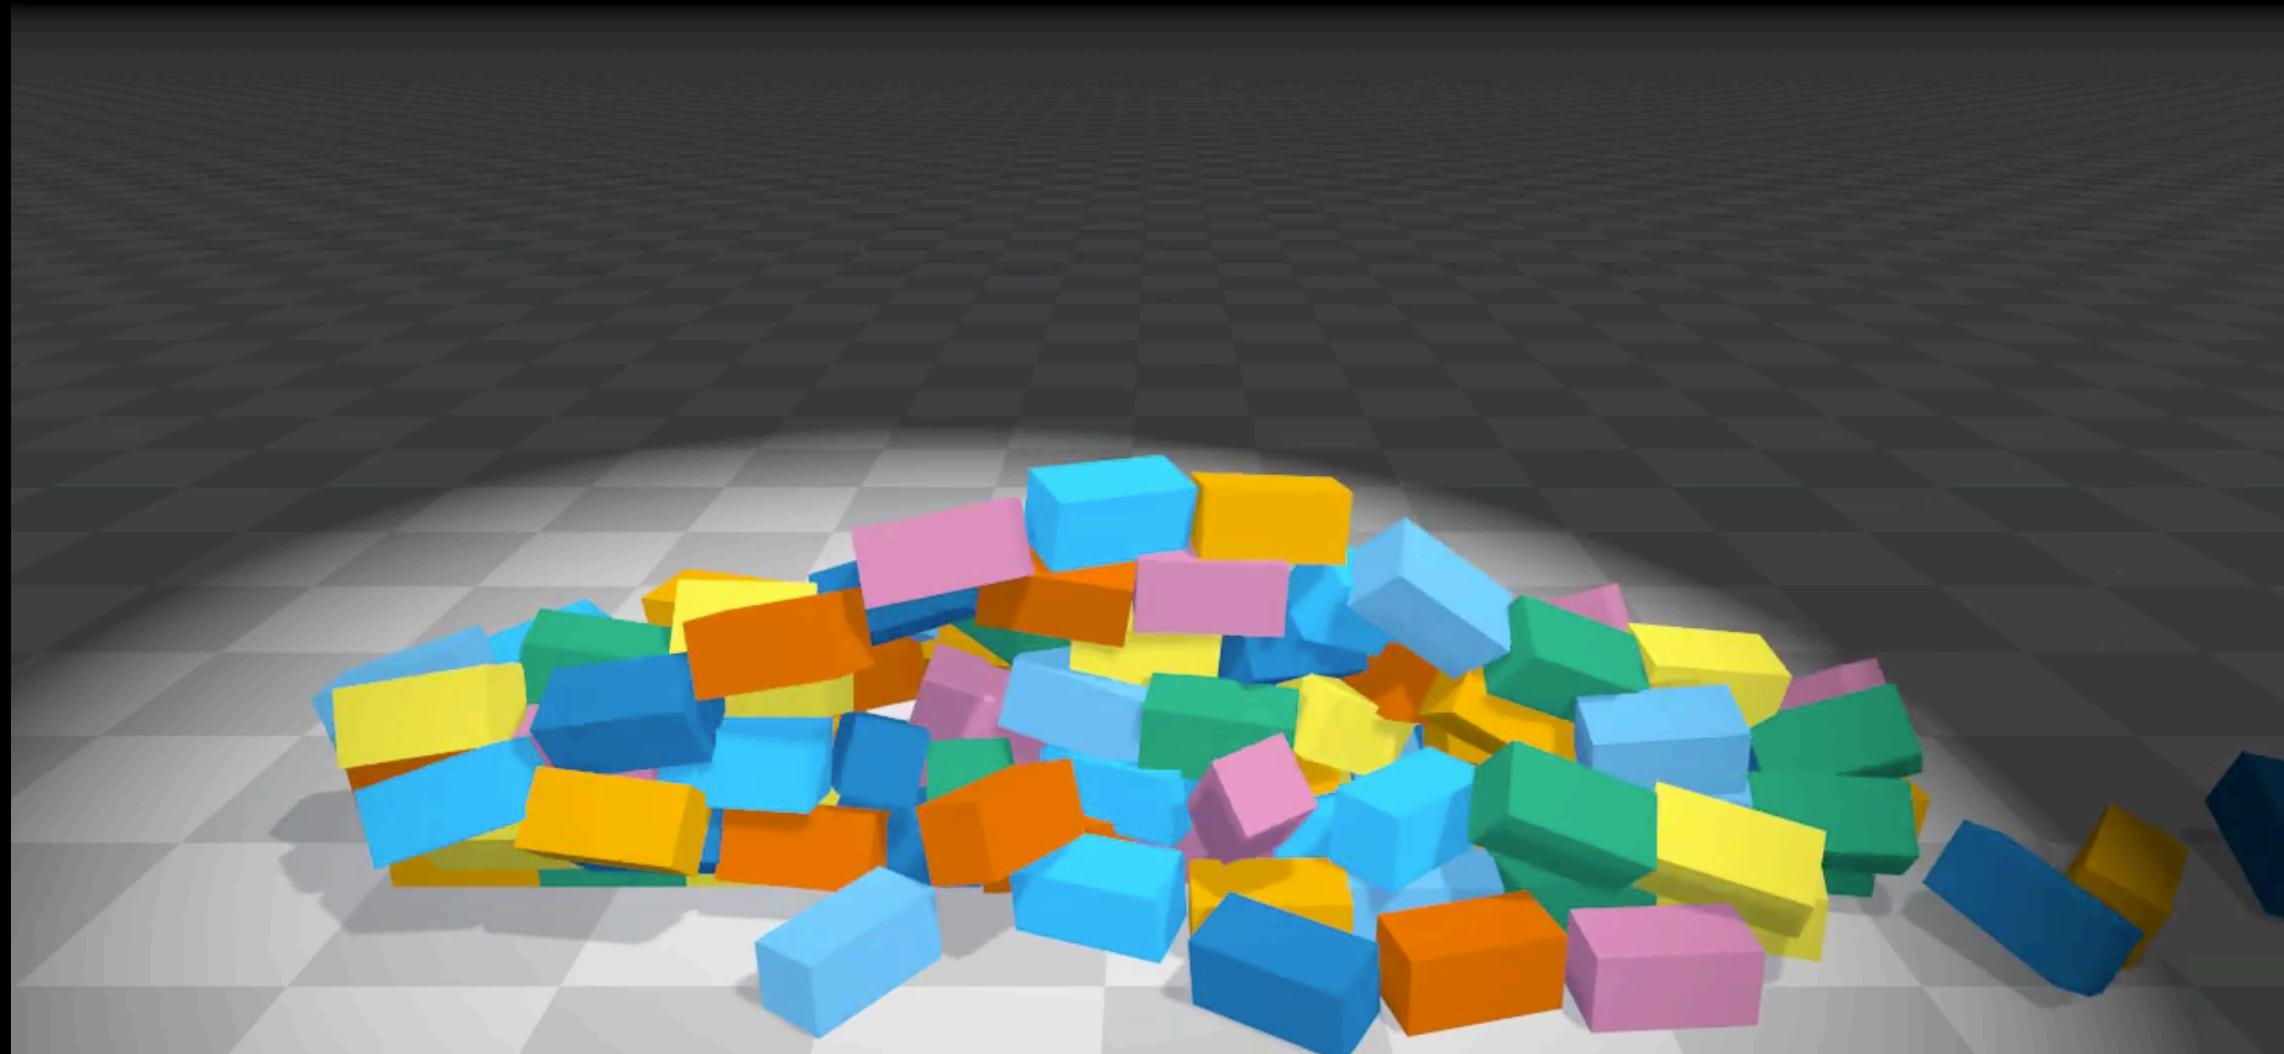
Approximate Shock Propagation

- A parallel friendly solution
- Instead of discrete layers with infinite mass, we **modify mass continuously**
- Choose 'stack height' function and evaluate for each particle:

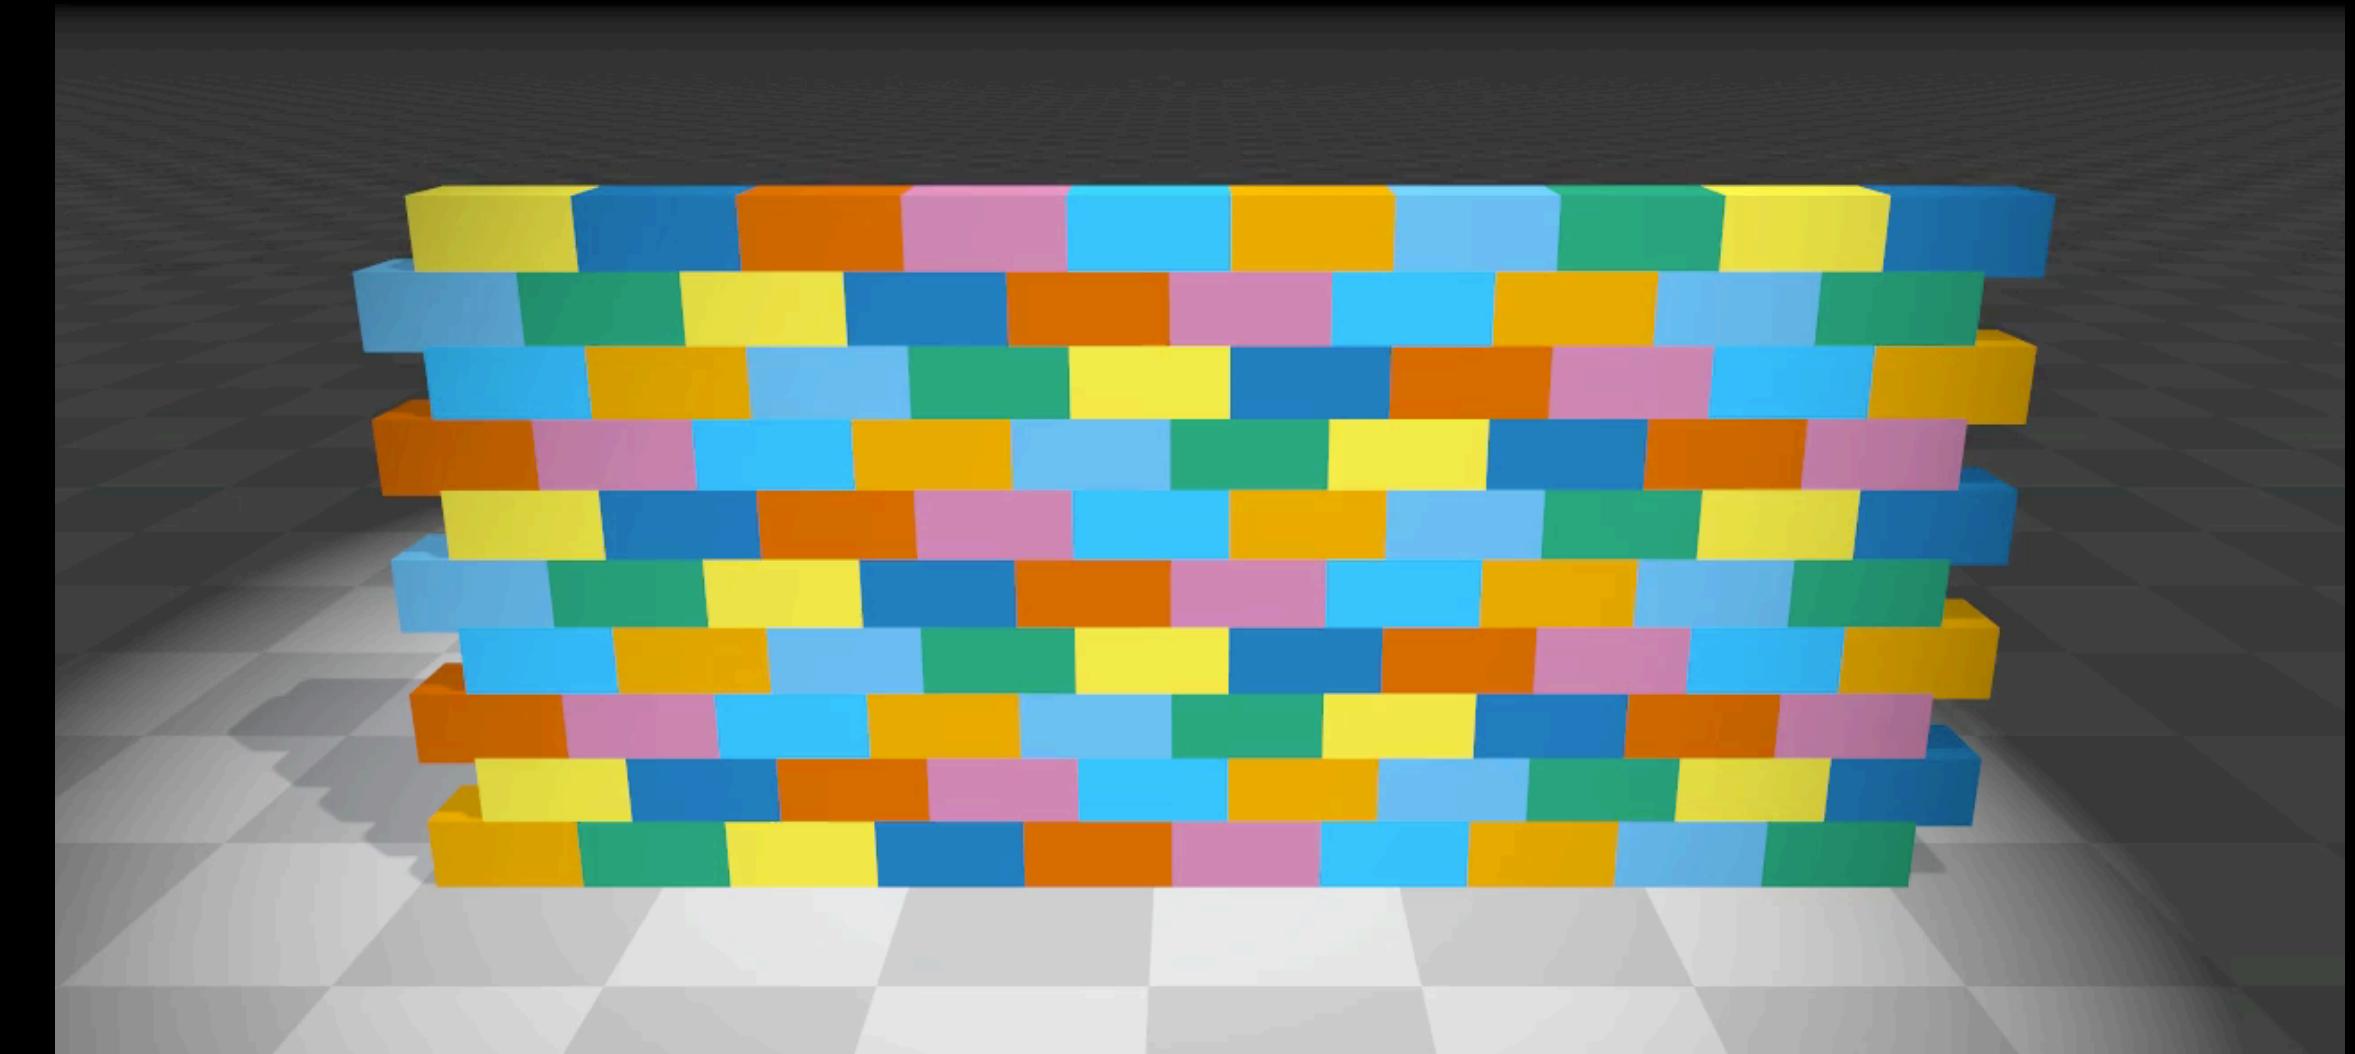
$$height(\mathbf{x}) = |\mathbf{x} - \mathbf{x}_{boundary}|$$

- Temporarily scale particle mass inversely with height





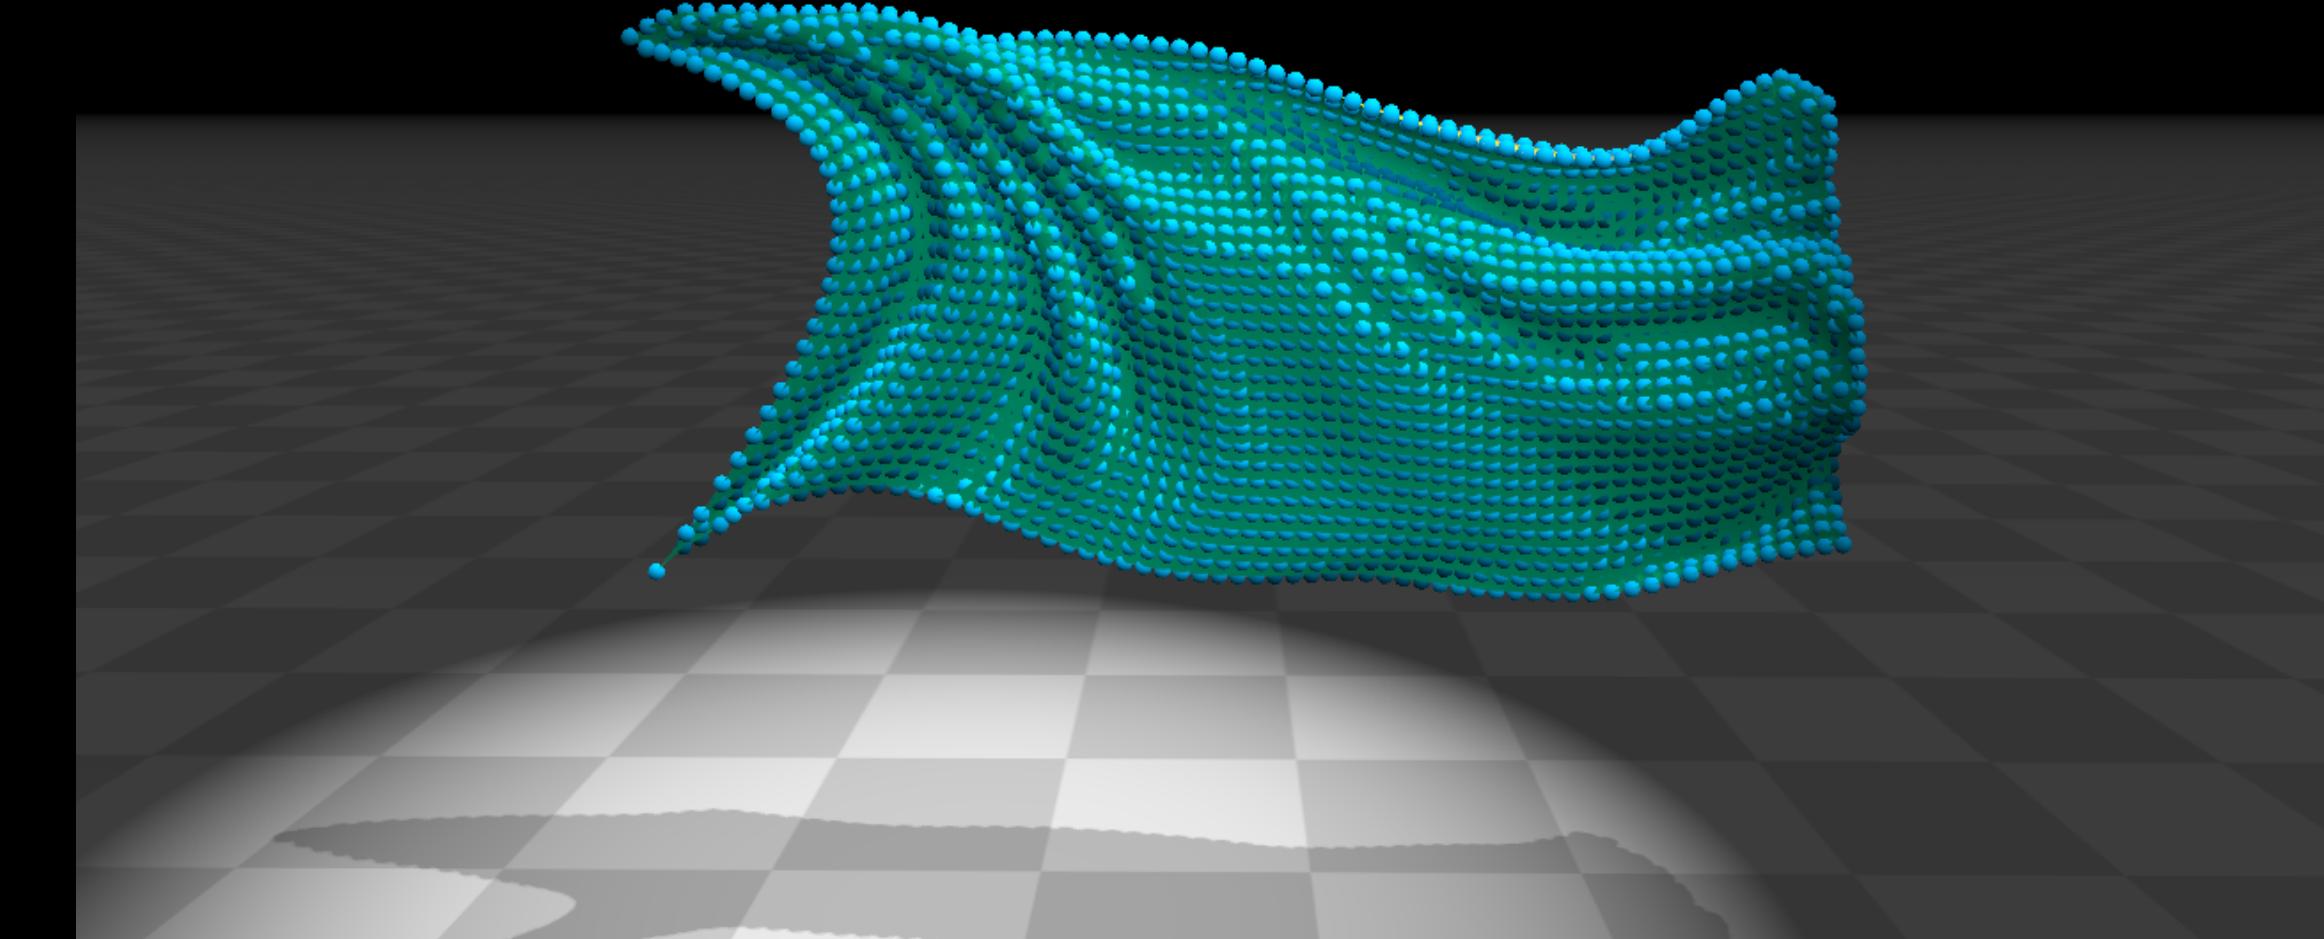
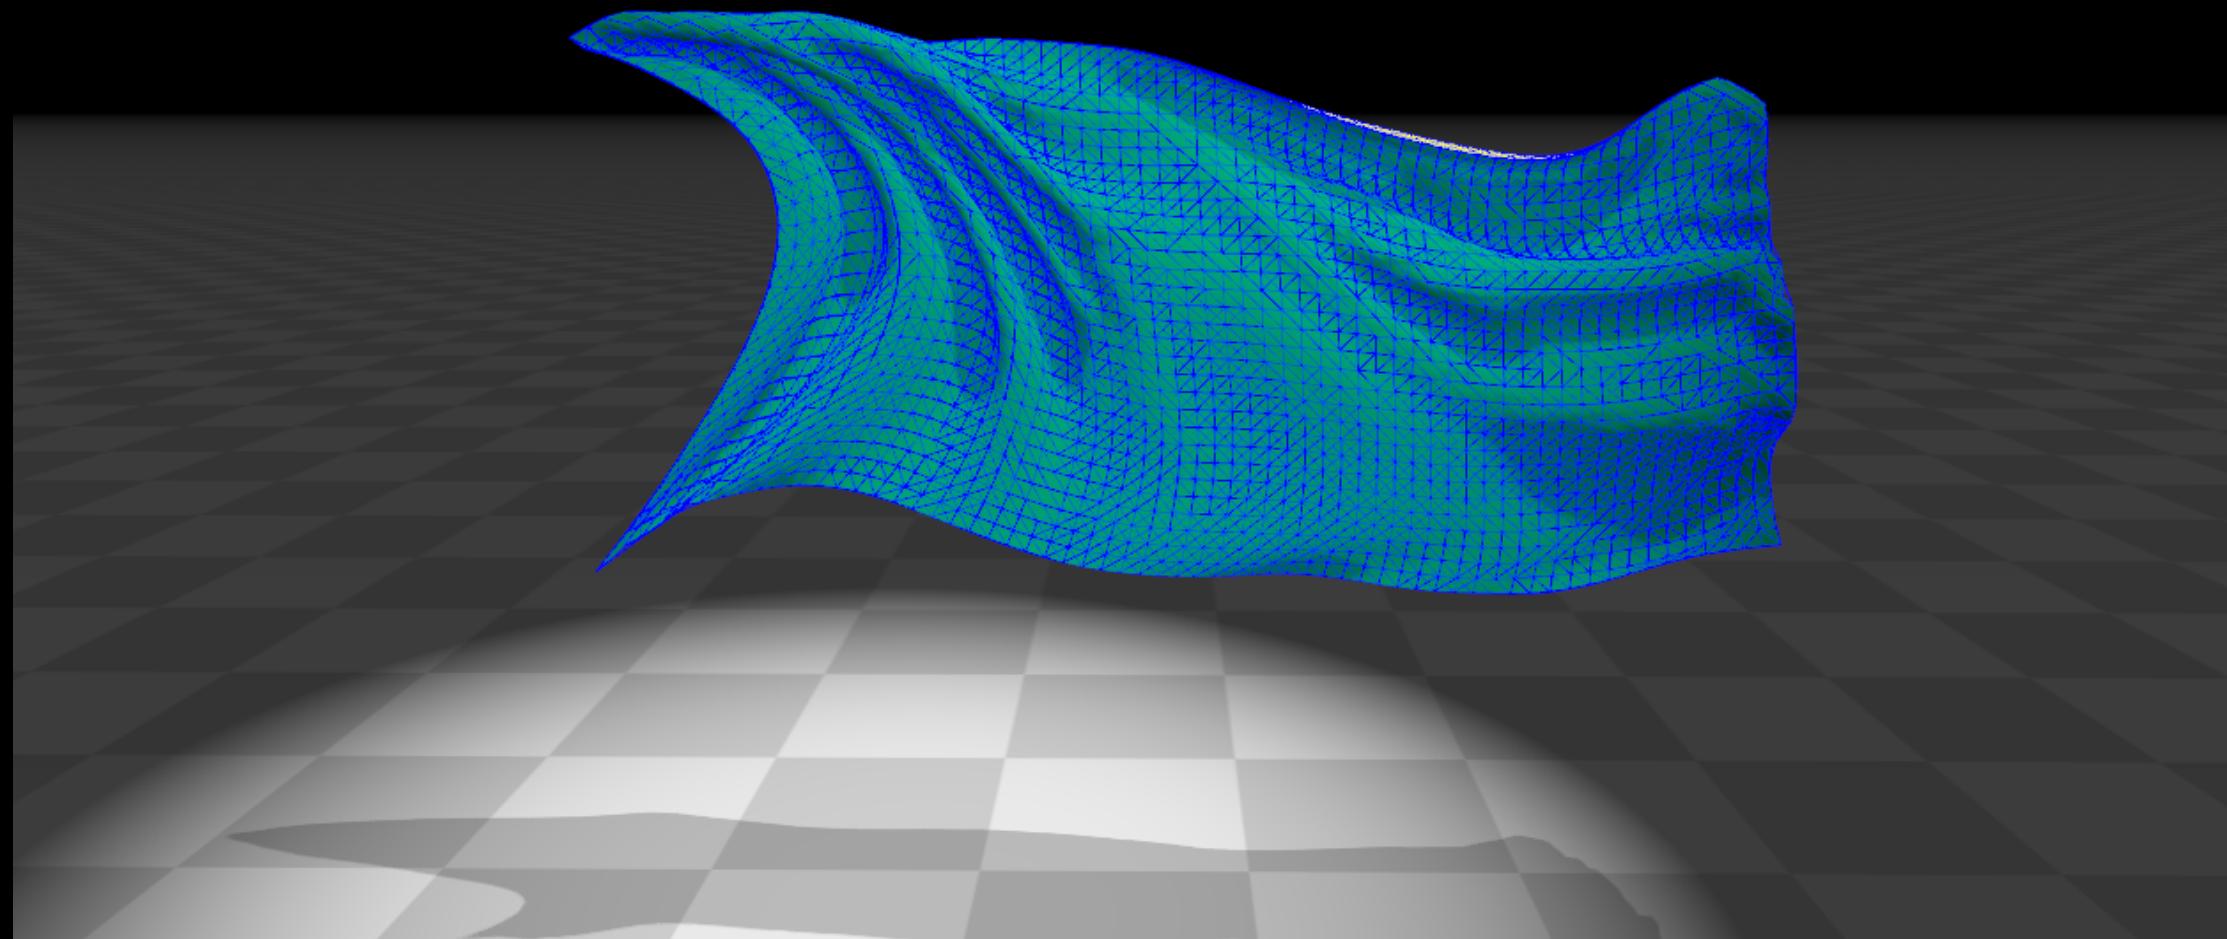
2 Iterations
No Shock Propagation



2 Iterations
With Shock Propagation

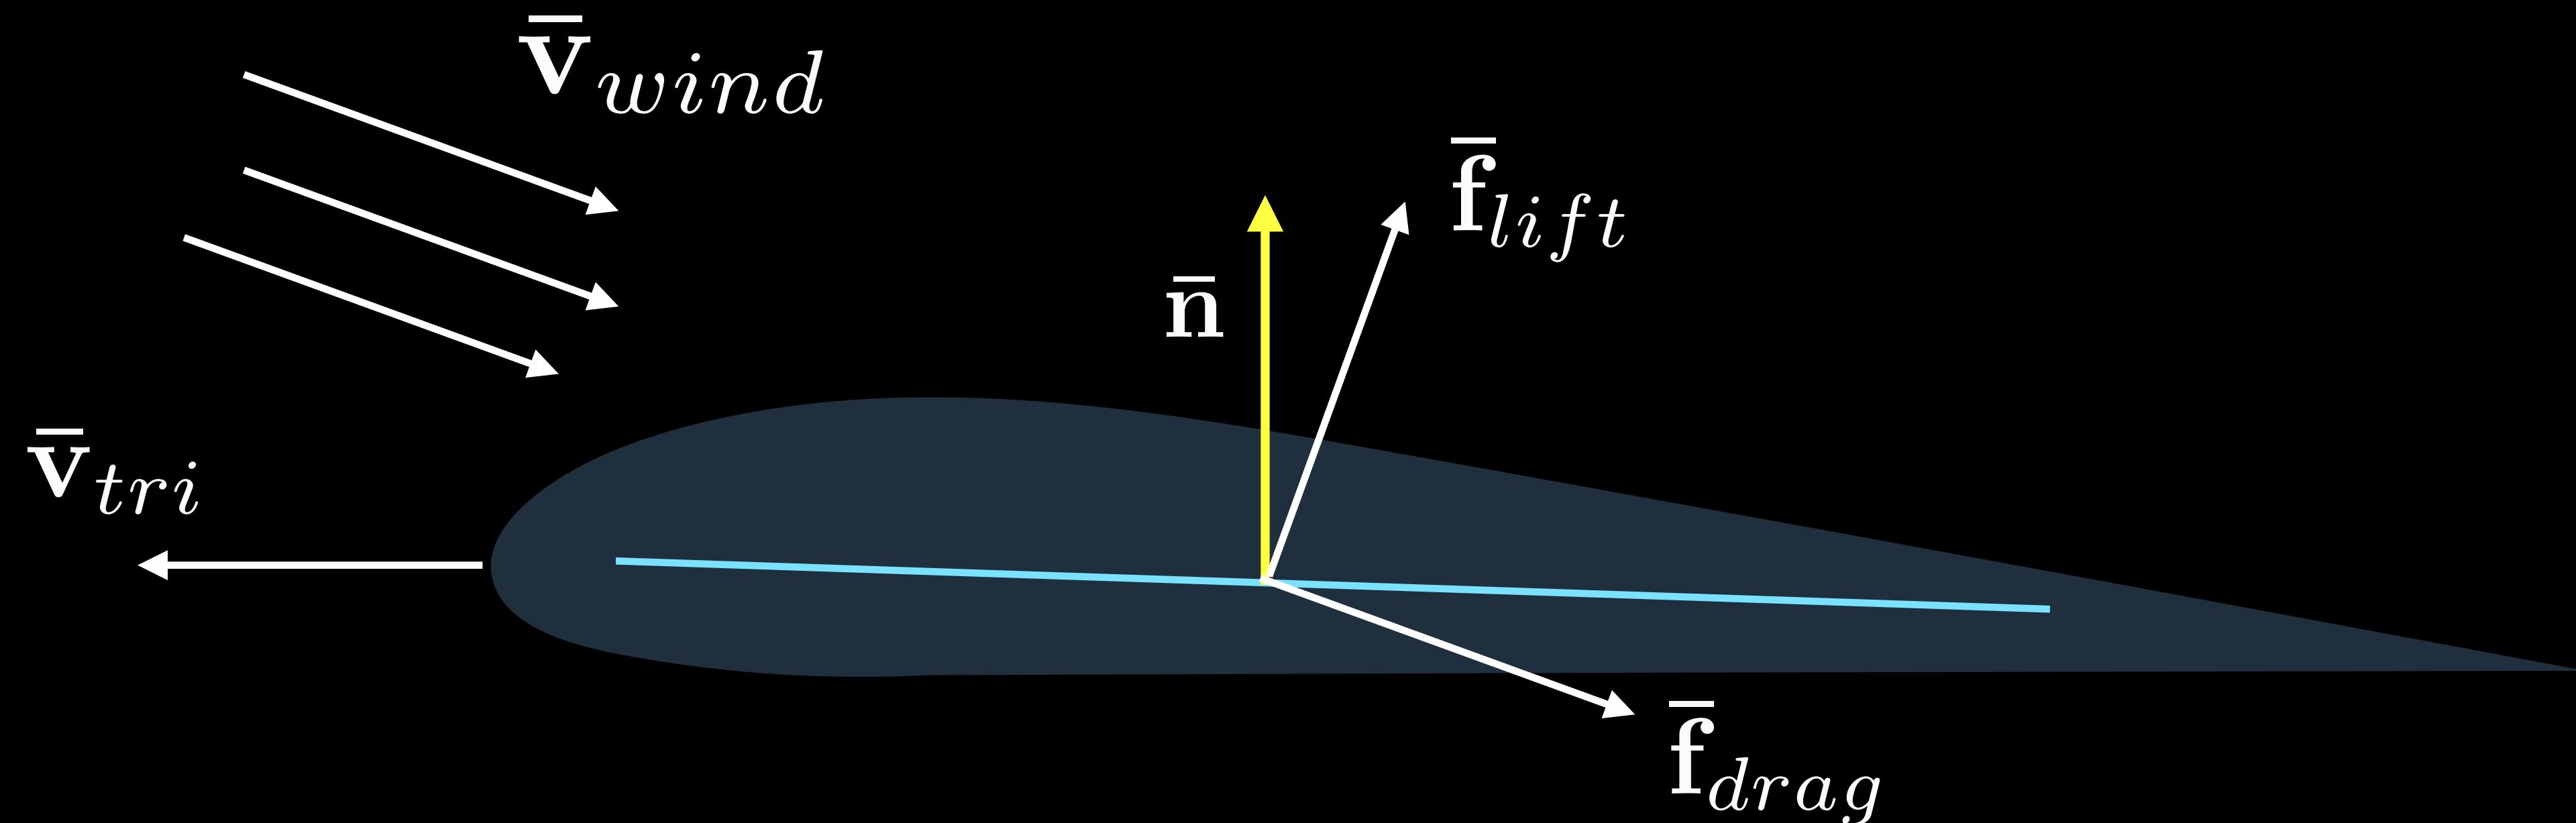
Cloth

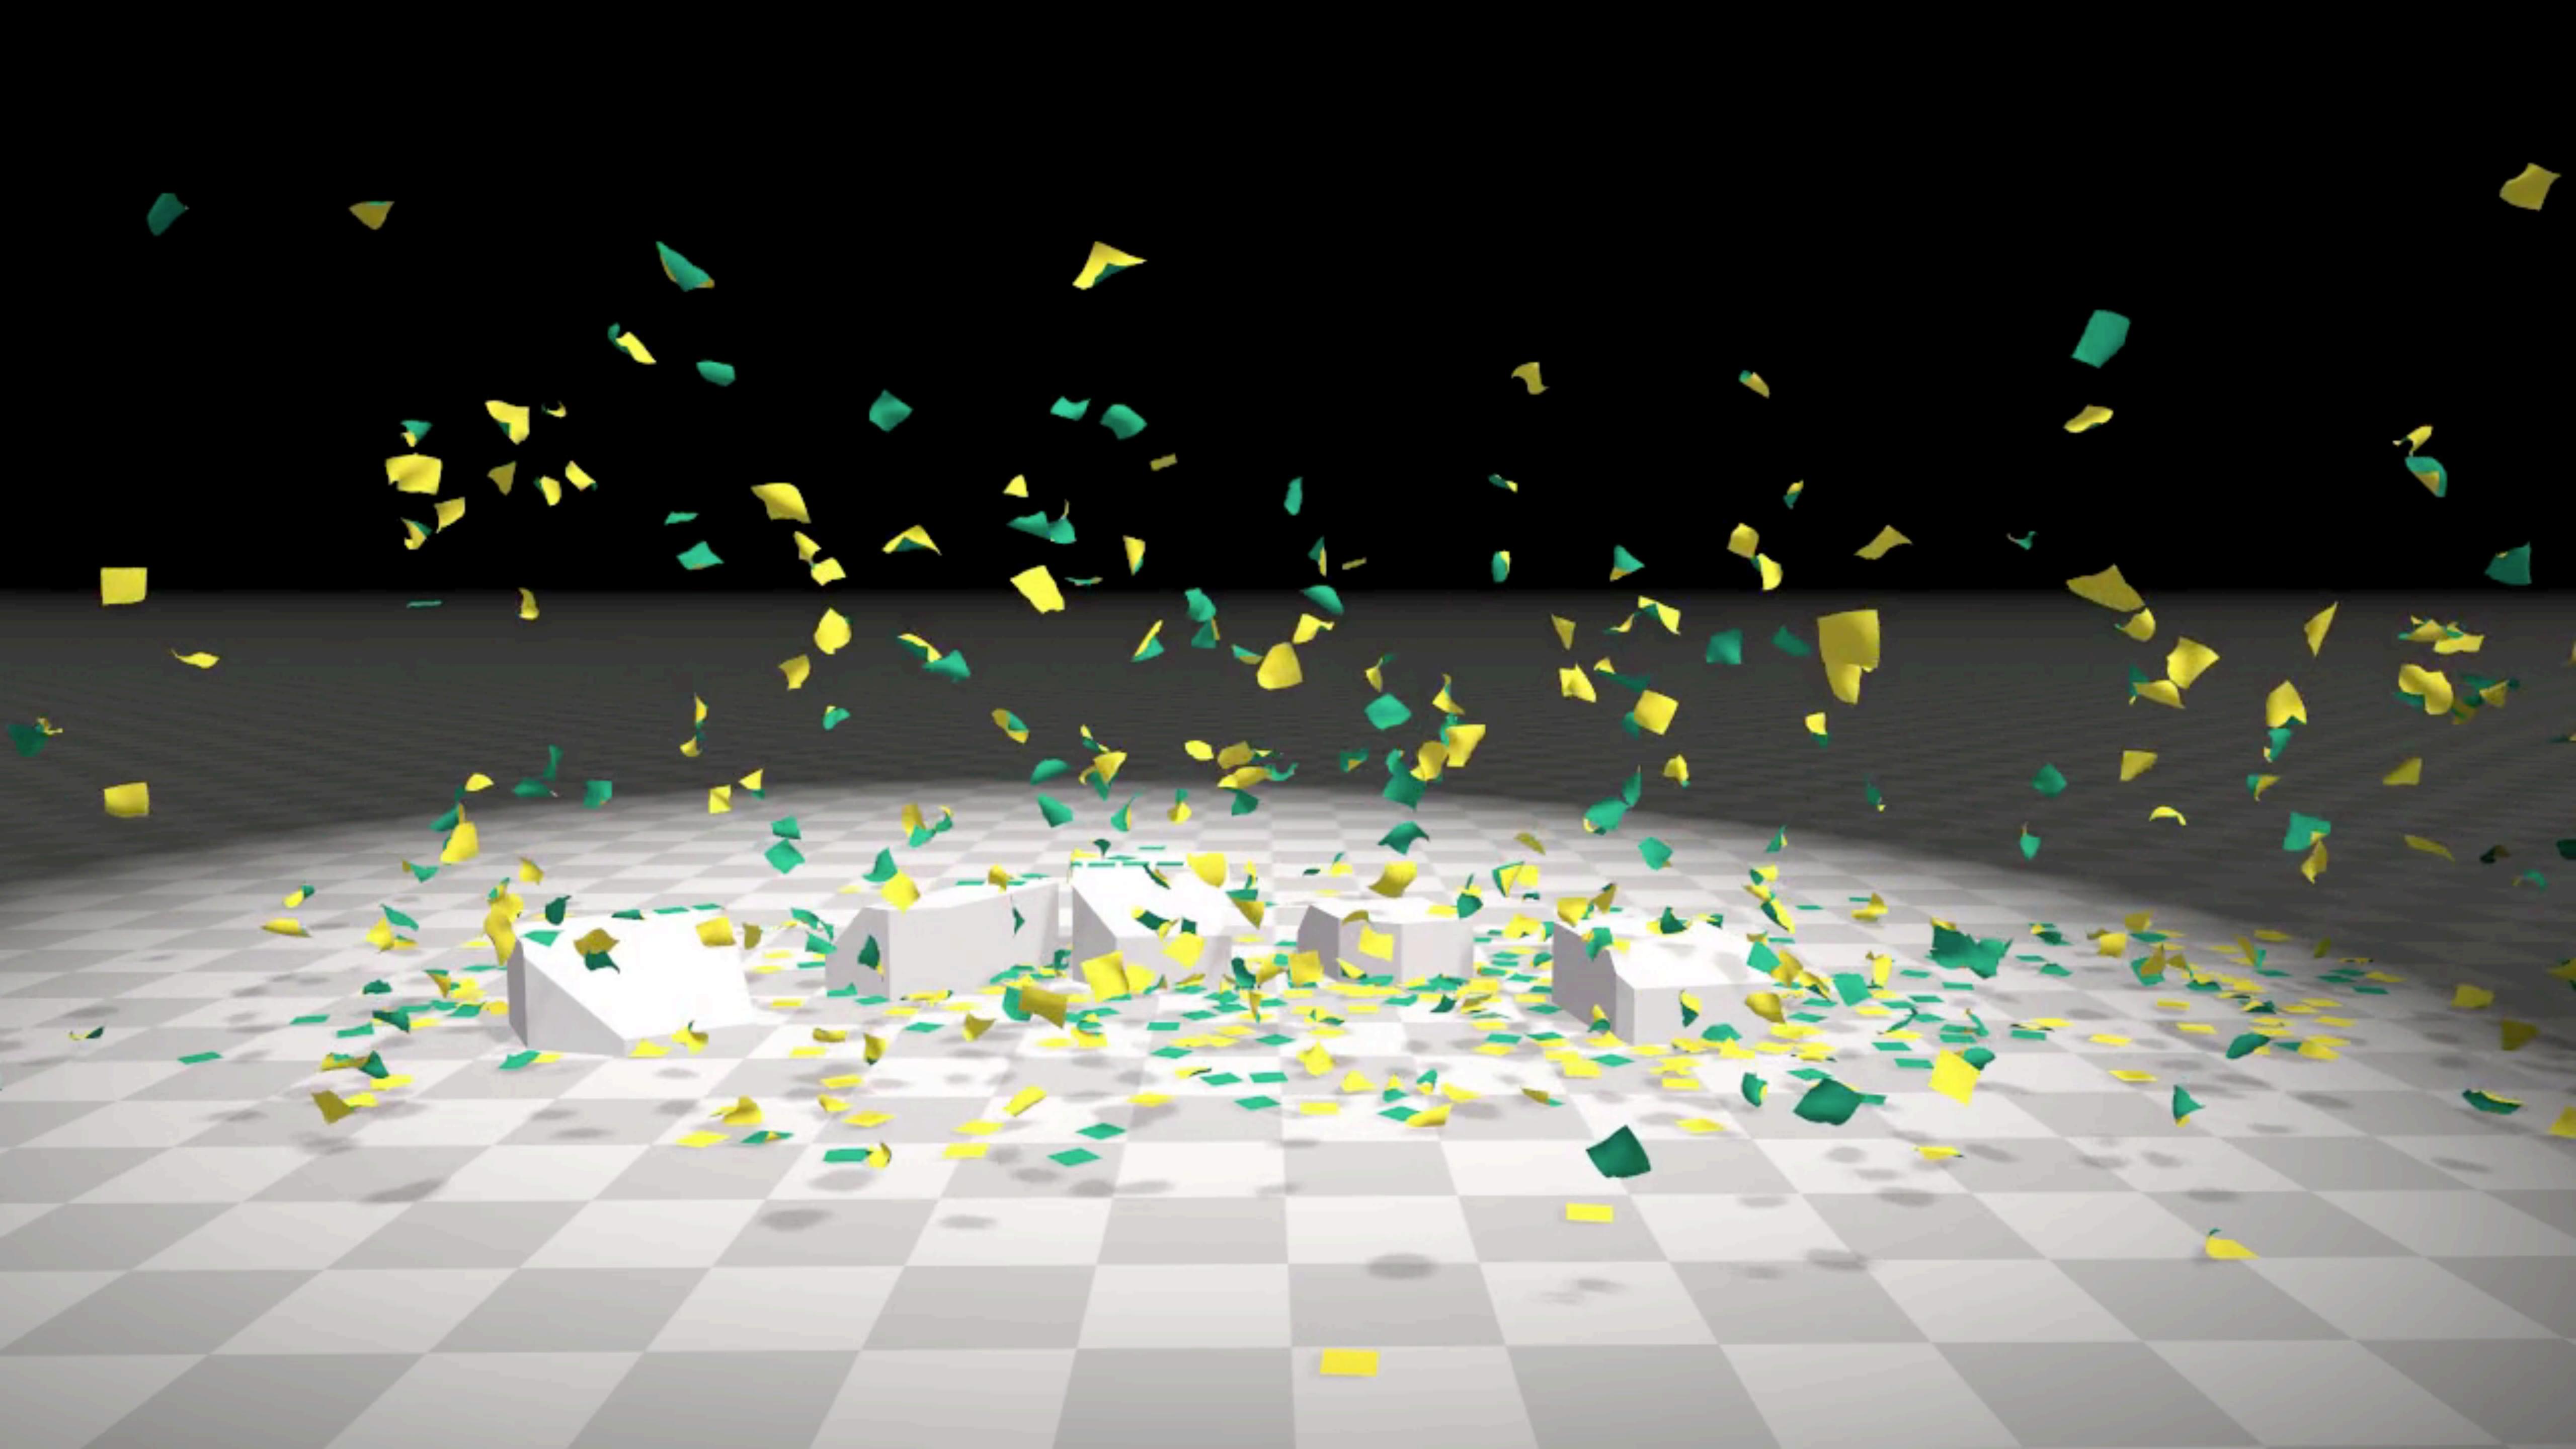
- Graph of distance + tether constraints
- Self-collision / inter-collision automatically handled



Cloth - Forces

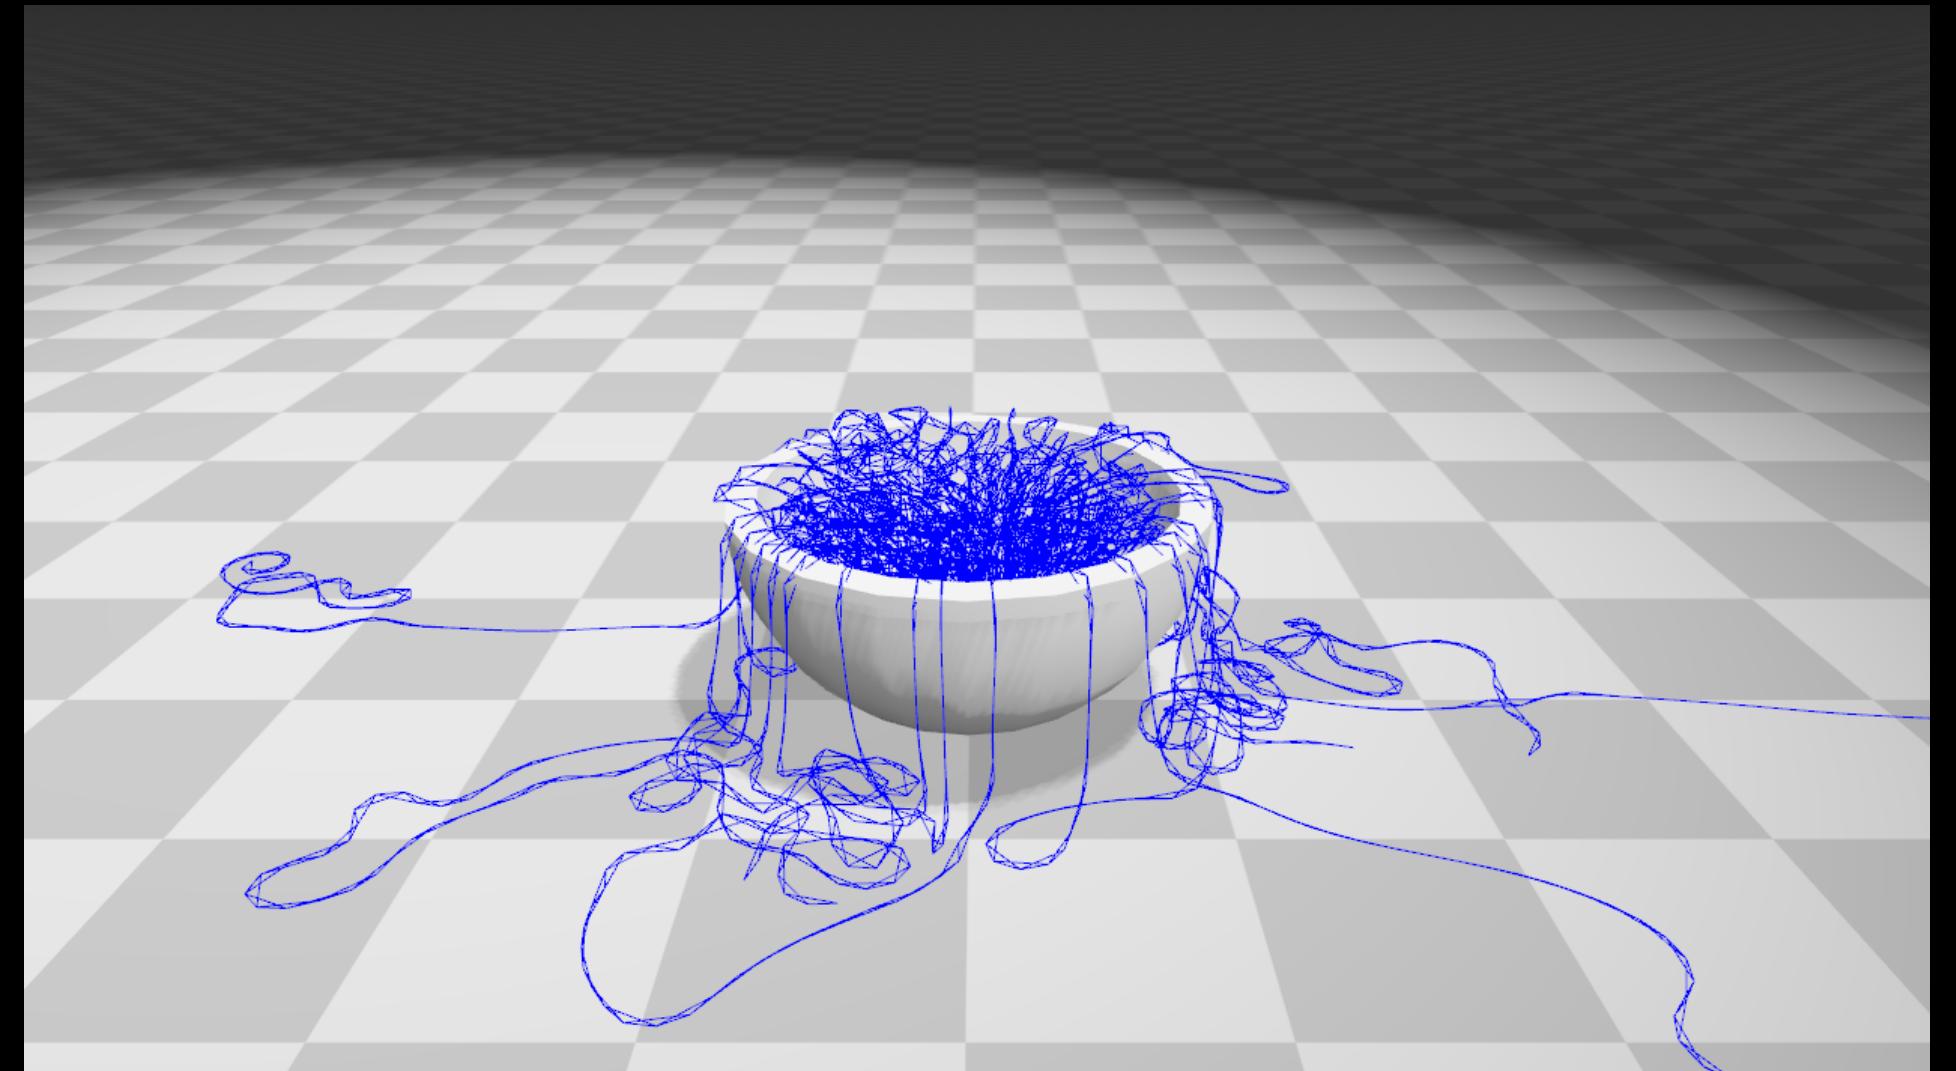
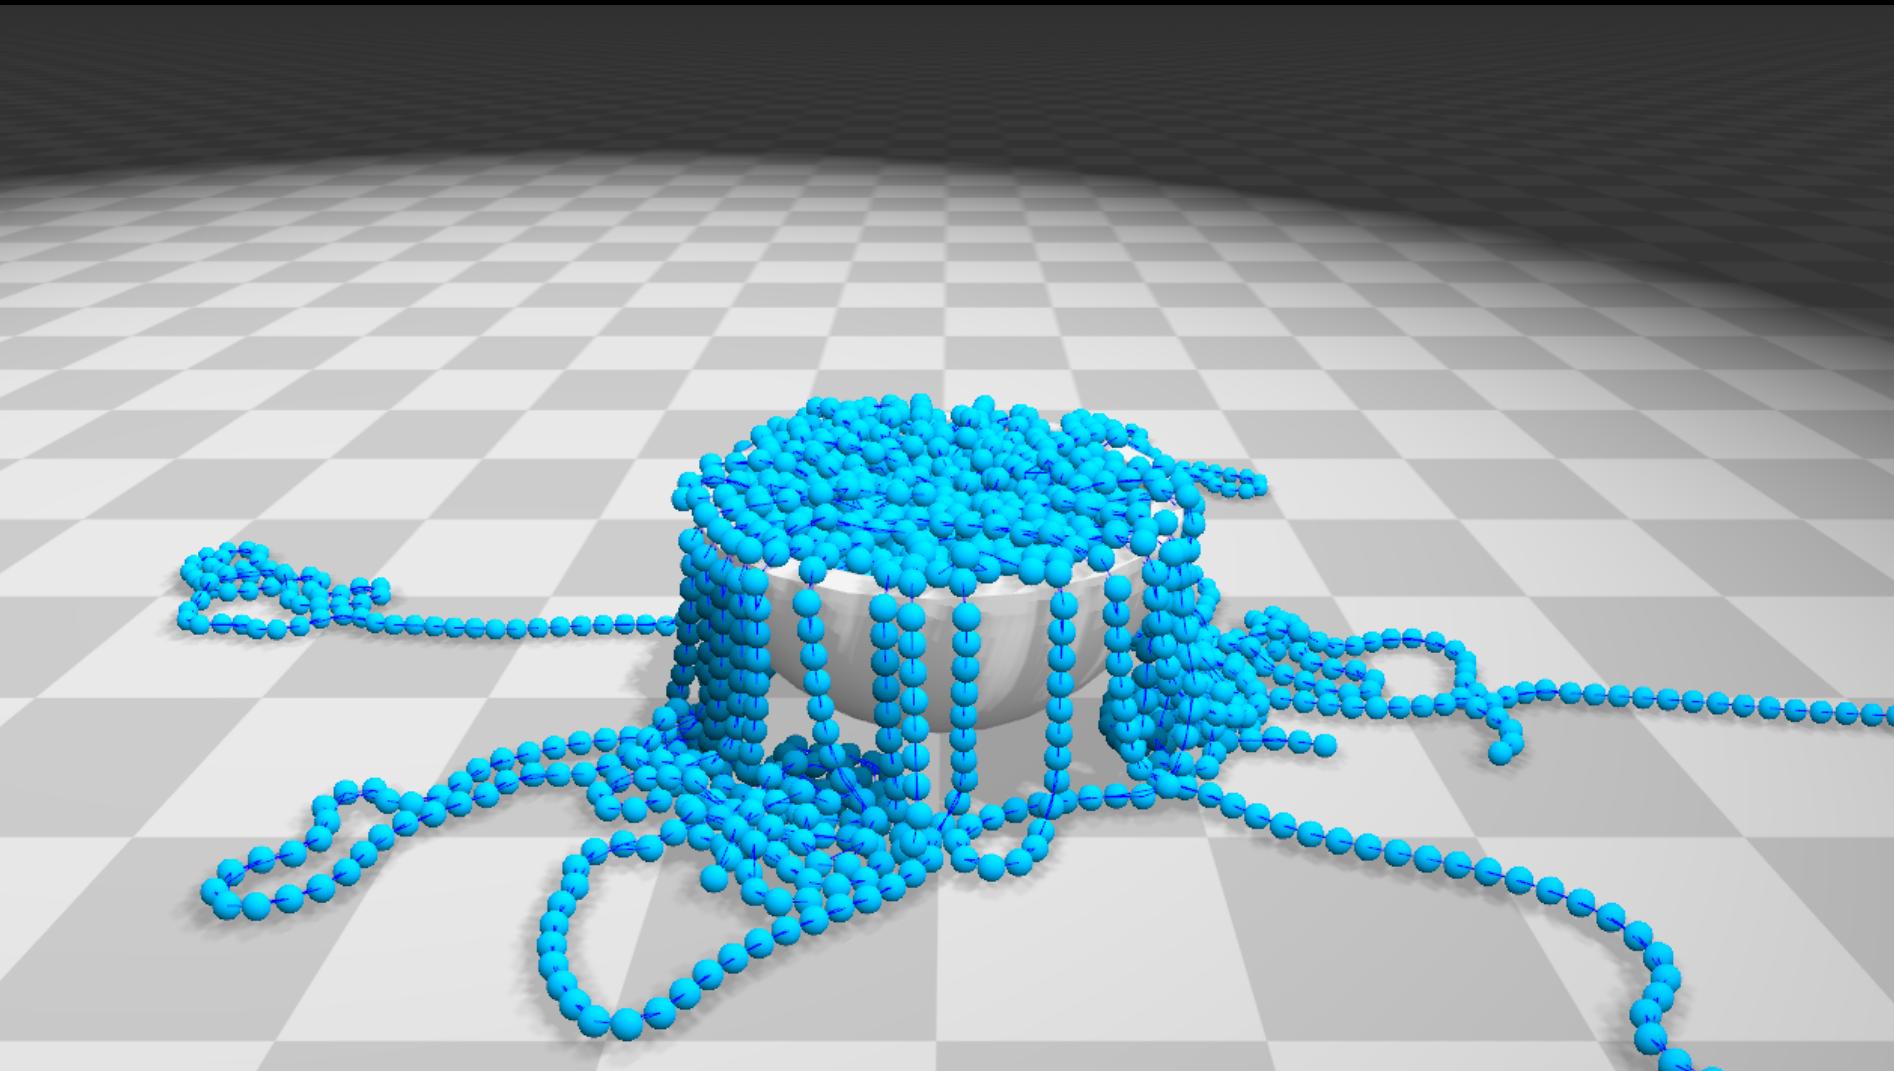
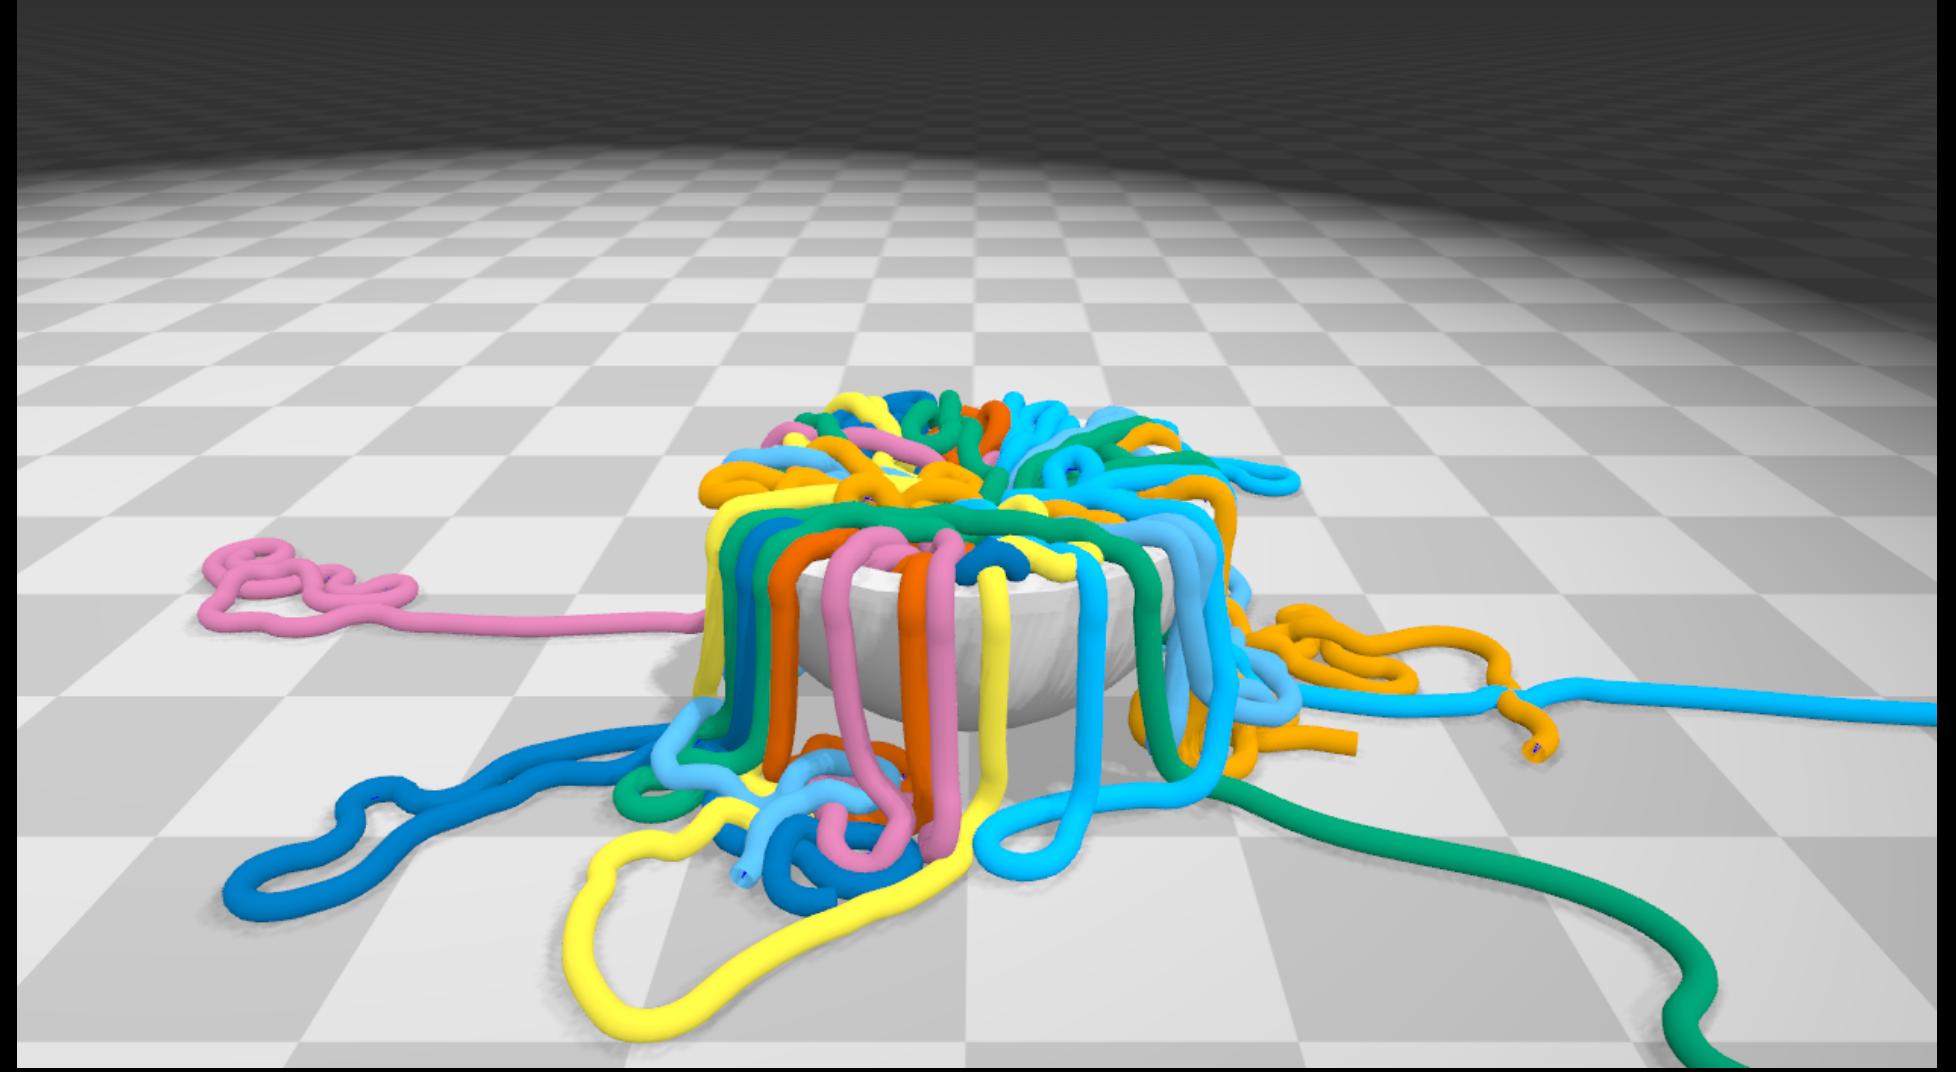
- Basic aerodynamic model
- Treat each triangle as a thin airfoil to generate lift + drag
- Flexible enough to model paper planes





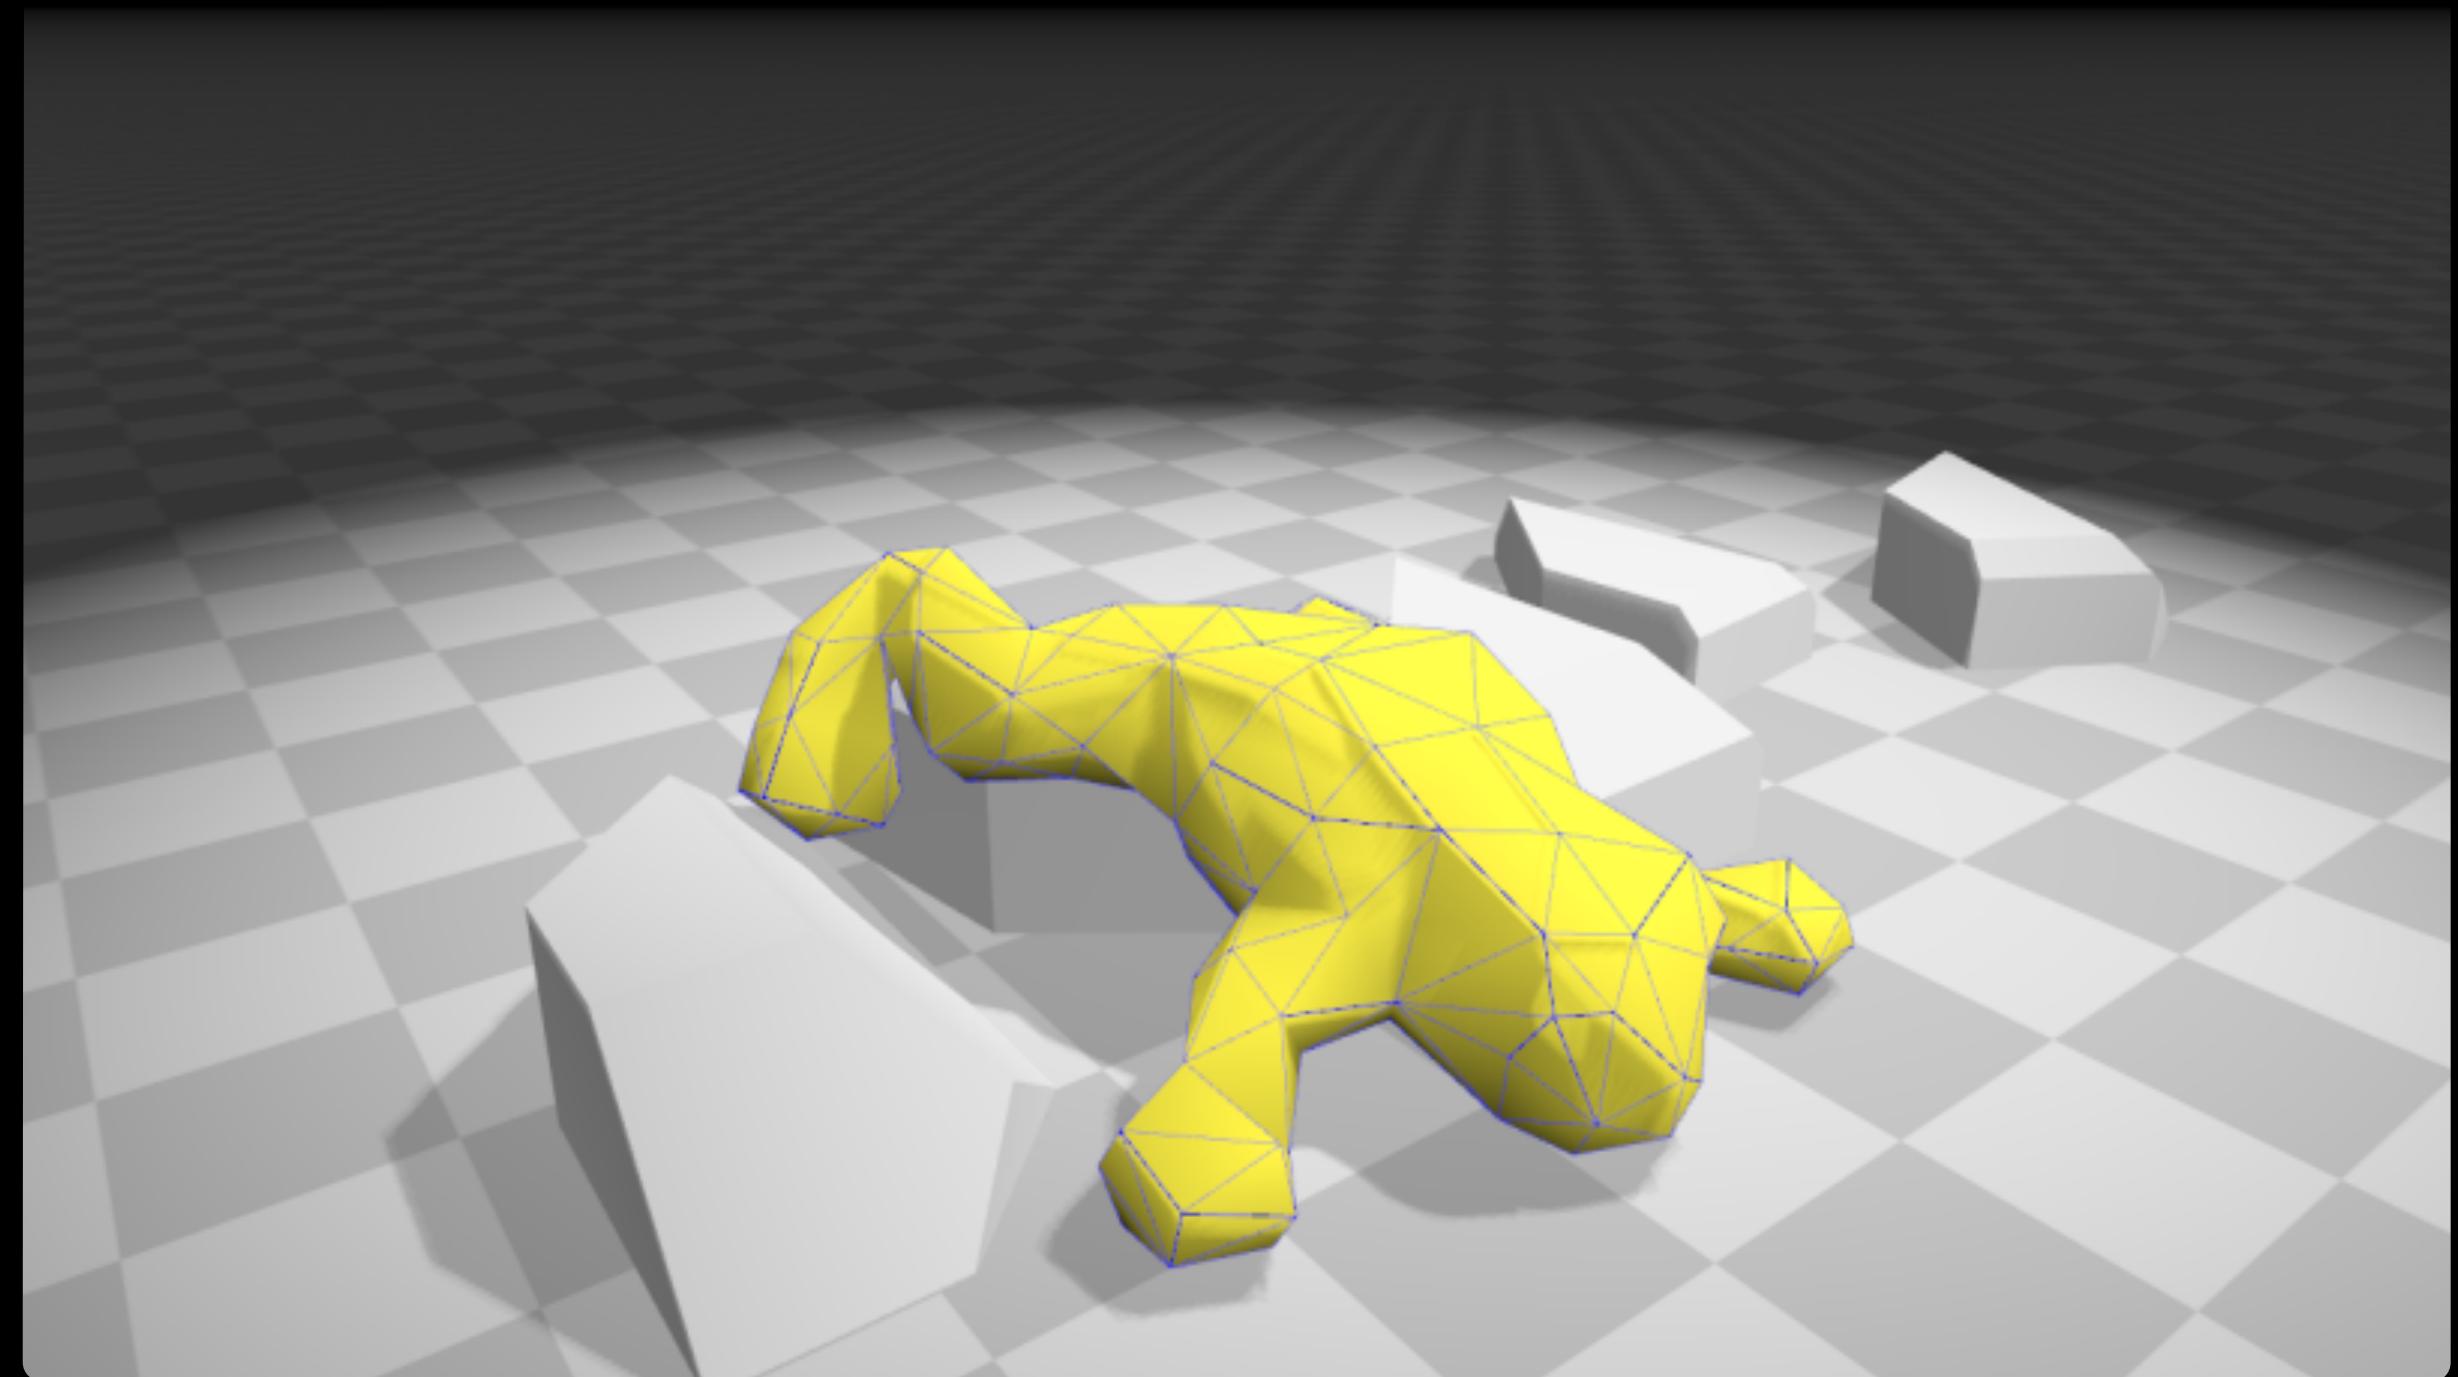
Ropes

- Build ropes from distance + bending constraints
- Fit Catmull-Rom spline to points
- Good candidate for GPU tessellation unit
- No torsion constraint (need orientation)



Deformables

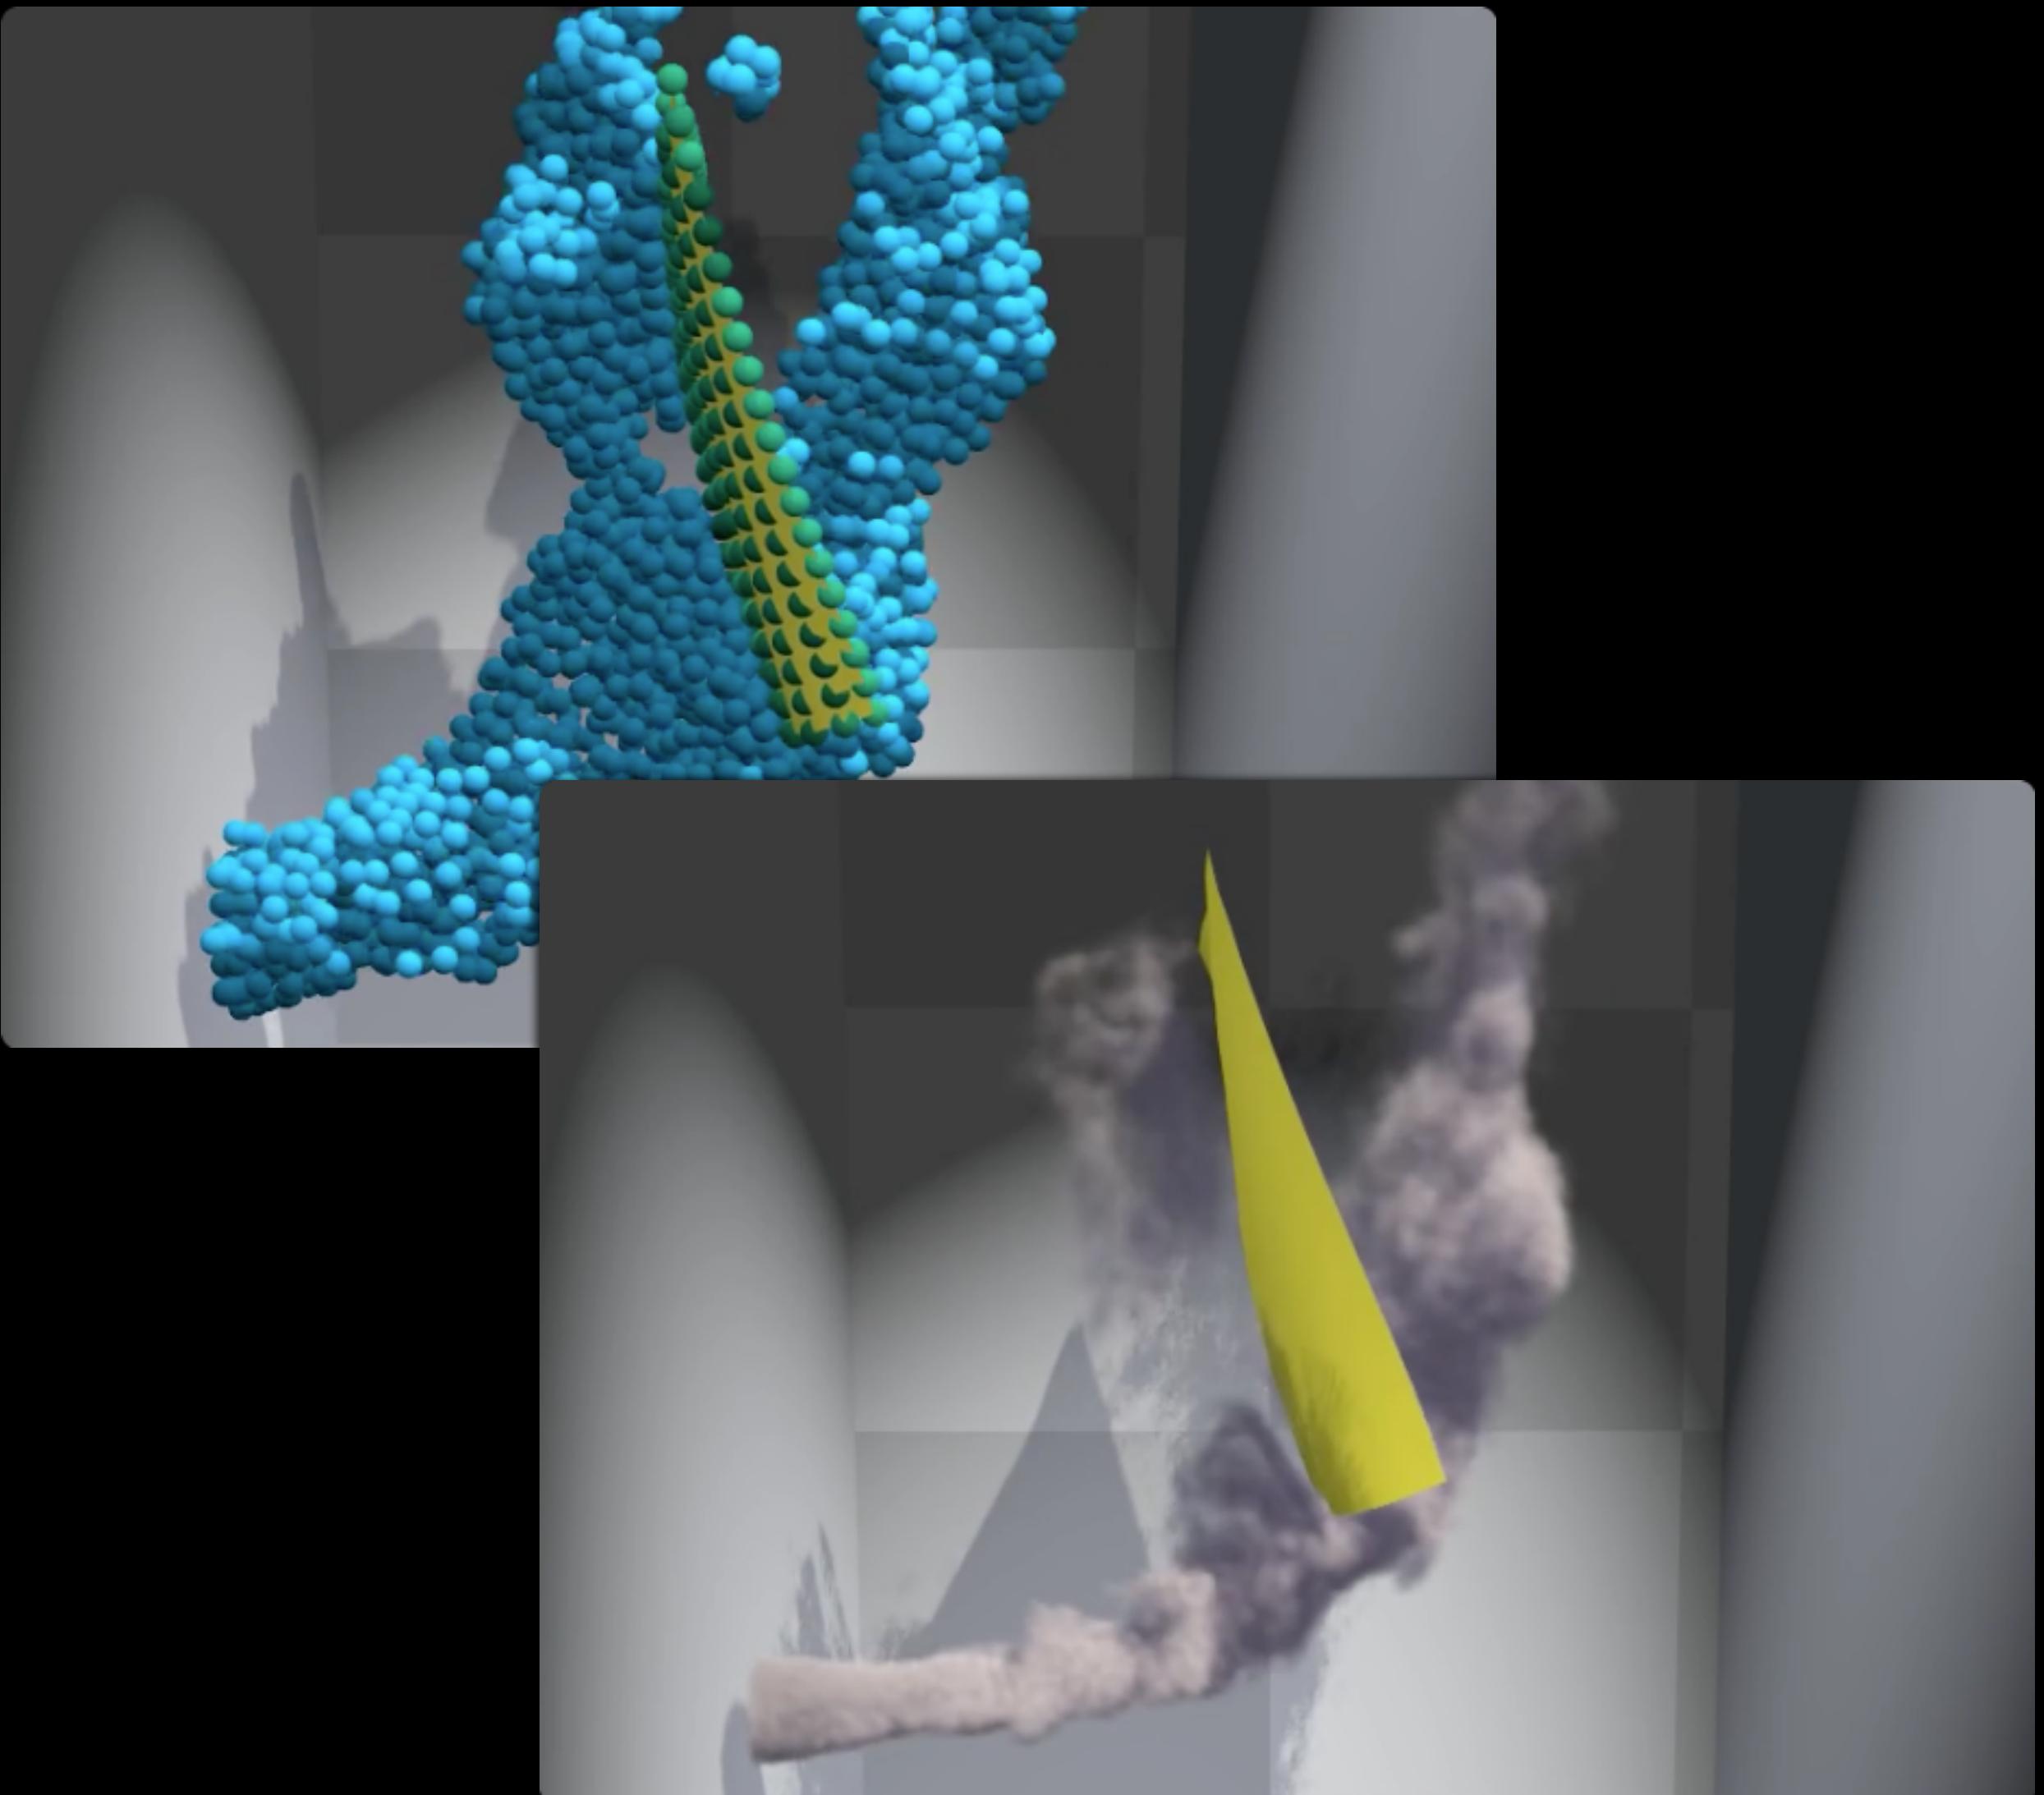
- Tetrahedral meshes -> mass spring system
- Tetrahedral volume constraints
- Soft shape-matching



Gases

Gases

- Treat as an incompressible fluid with density constraint
- Sparse representation
- Passive smoke advection (“diffuse particles”)



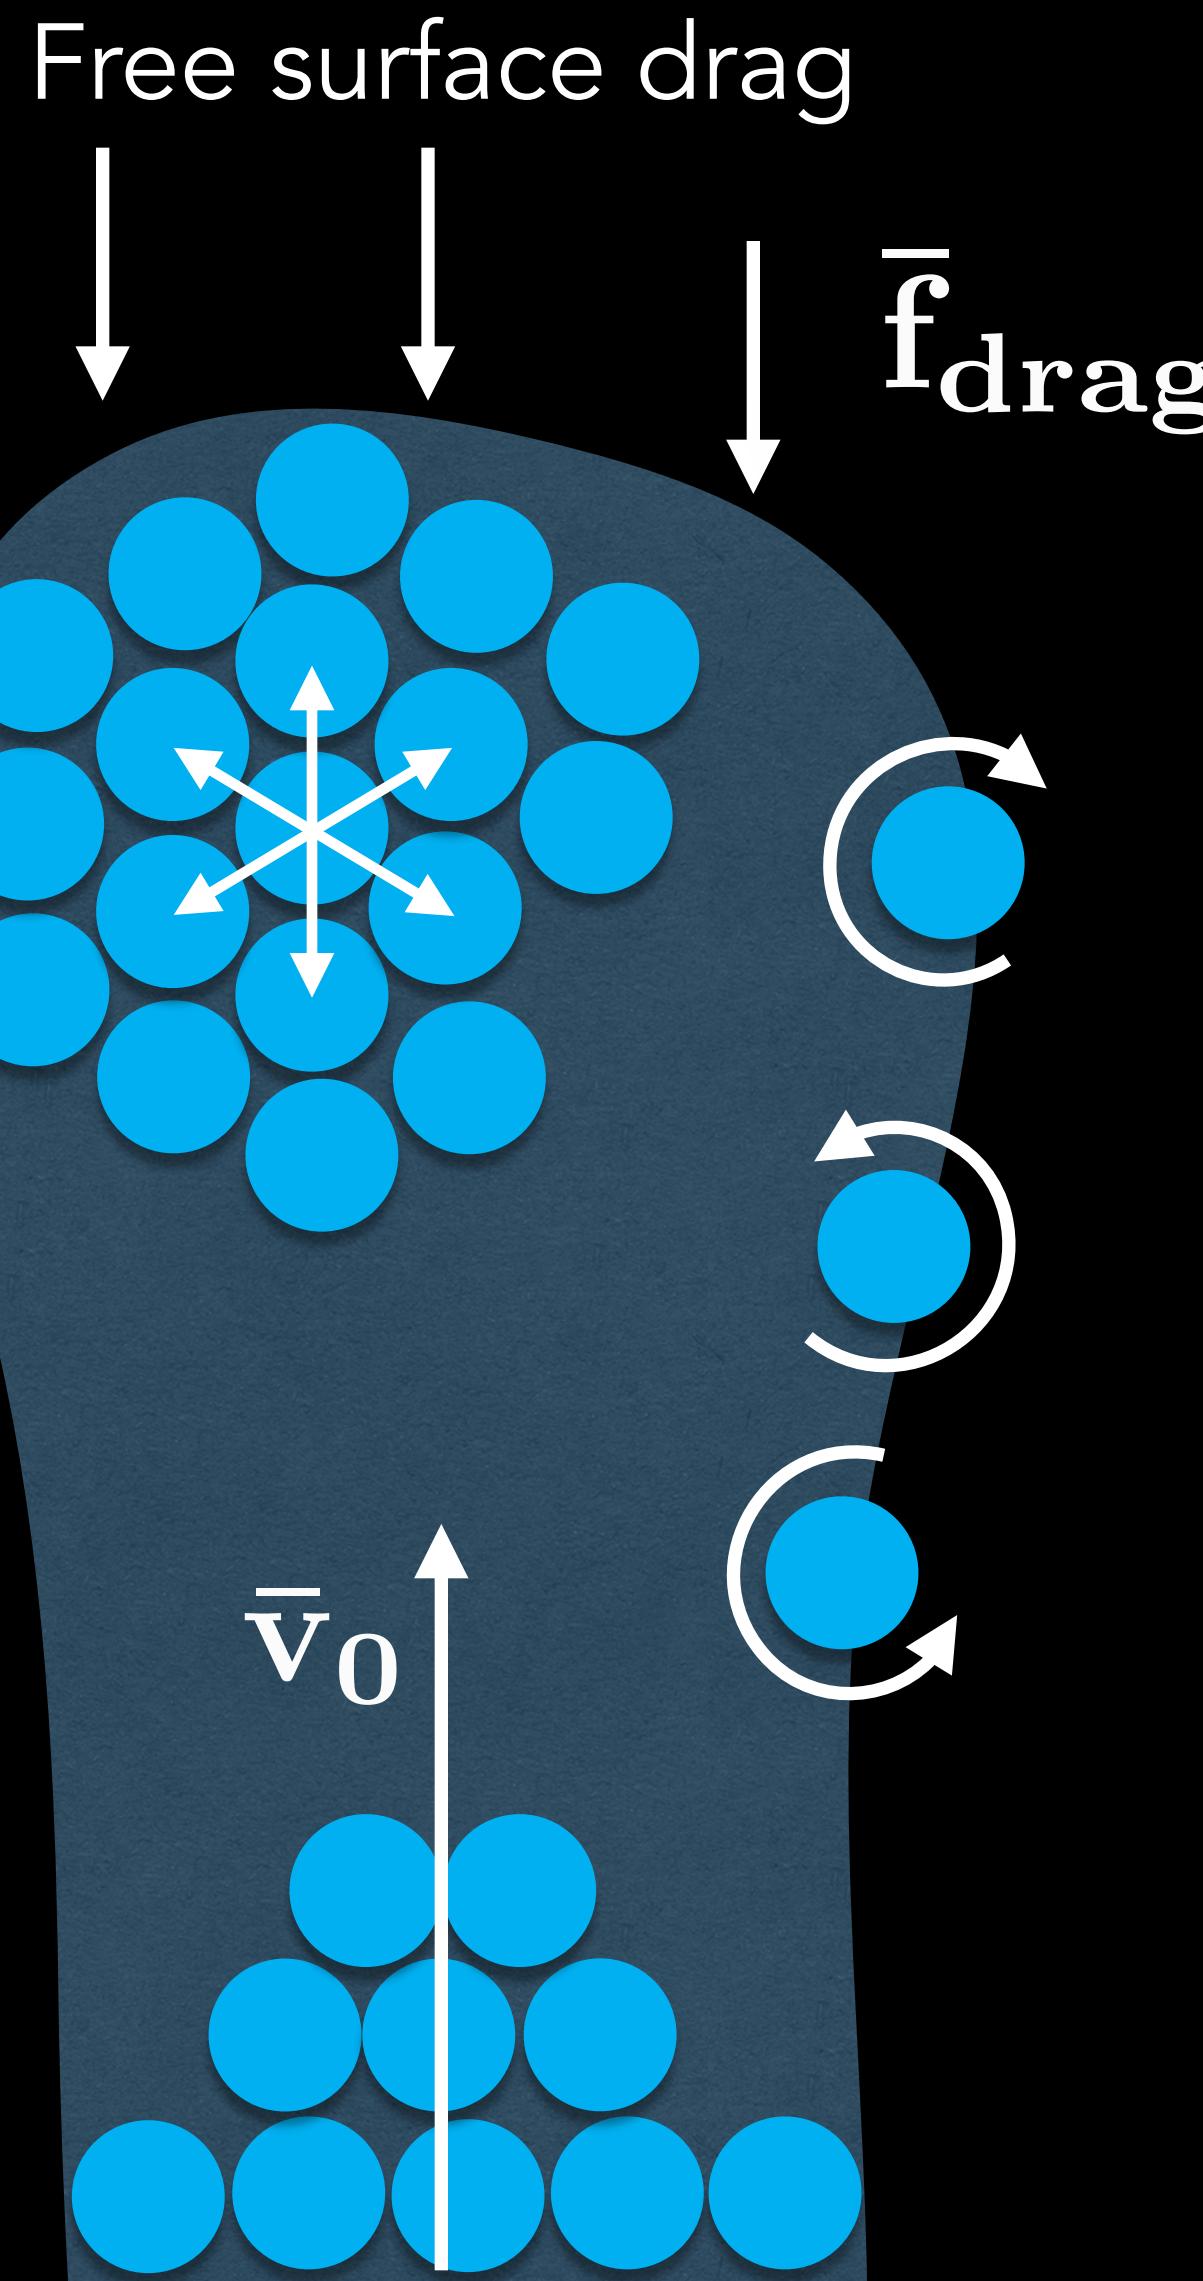
Gas Forces

Density gradient via SPH derivatives:

$$\nabla \rho = \bar{n}$$

Pressure gradient via Boussinesq approximation:

$$\nabla p = \bar{g}$$



Baroclinic vorticity:

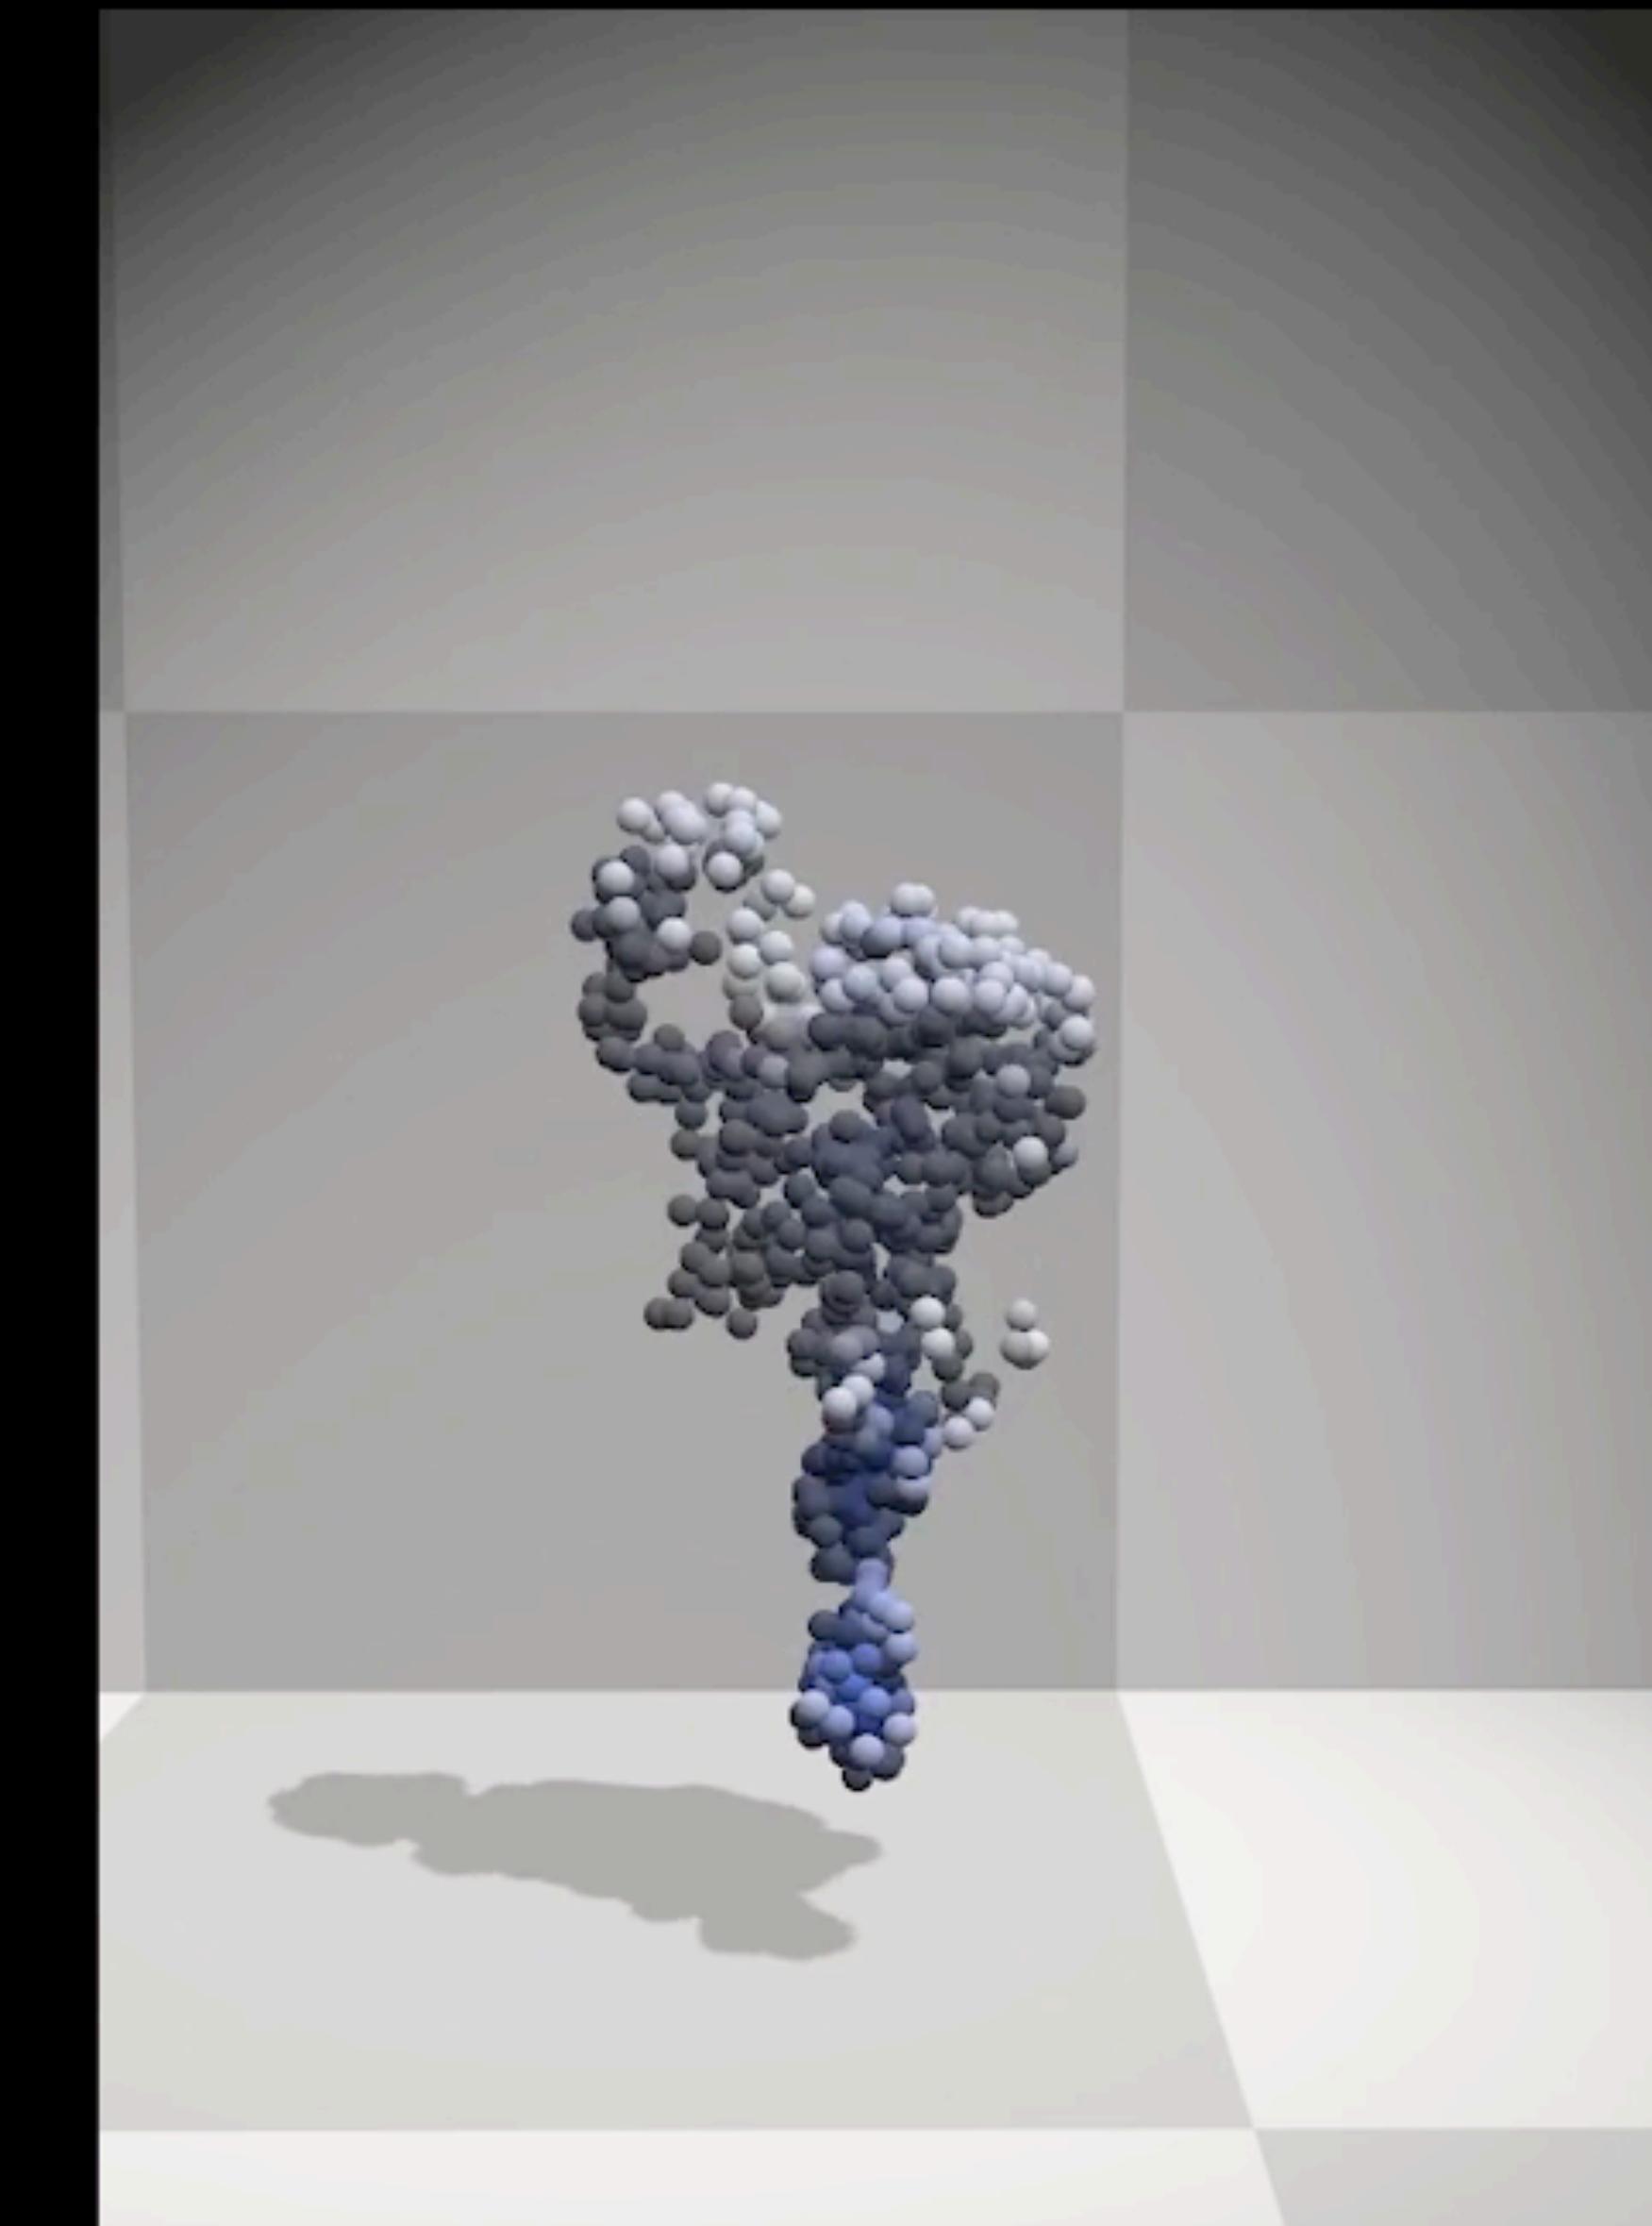
$$\frac{D\bar{\omega}}{dt} = \nabla \rho \times \nabla p$$

Driving vorticity:

$$\bar{f}_{\text{vort}} = \bar{\omega} \times \bar{x}_{ij}$$



Smoke Particles



Fluid Particles

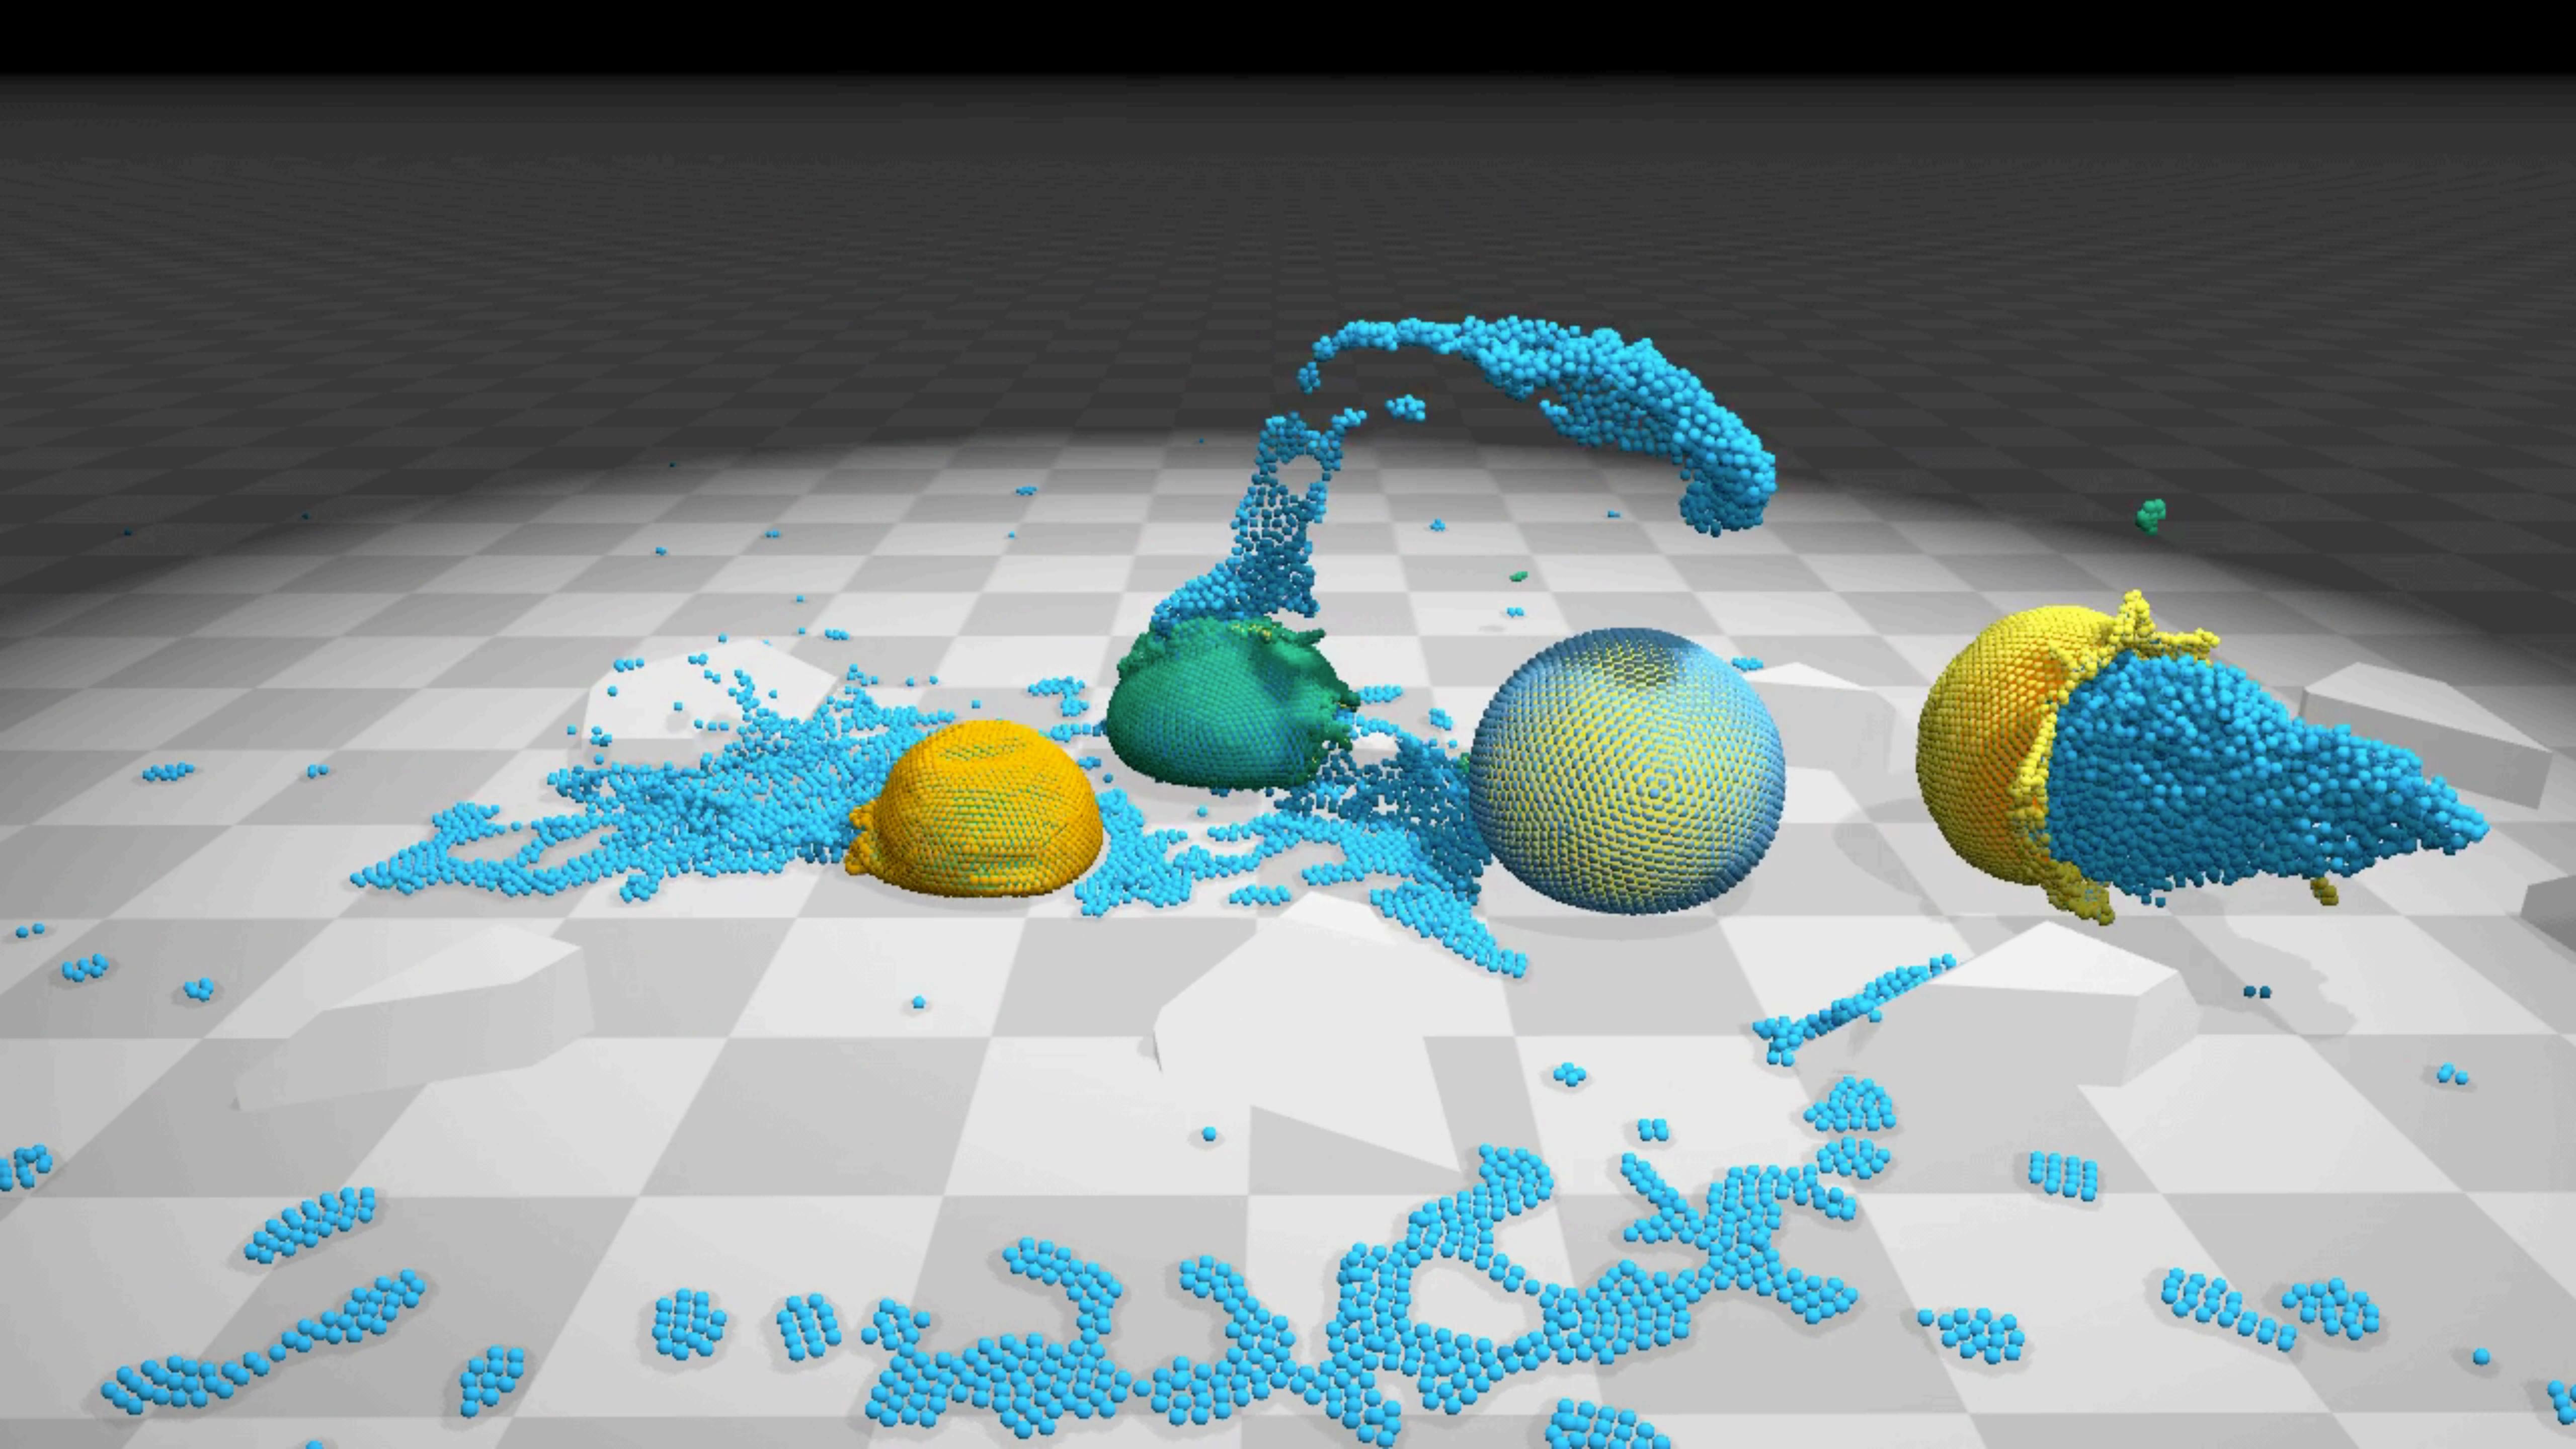
Smoke Rendering

- Back-to-Front sort particles in CUDA
- Point based rendering
- Approximate transmission using shadow-map depth as input to scattering function



Examples





Limitations / Future Work

- Representing smooth surfaces problematic
- Would like parallel and robust collision of simplices
- Dynamic re-seeding for gases
- Iteration independence for non-stiff constraints

Thank you!

Acknowledgements

- Thanks to the PhysX team and the paper reviewers
- Contact details:
- mmacklin@nvidia.com
- [@milesmacklin](https://twitter.com/milesmacklin)

References

- Two-way coupling of fluids to rigid and deformable solids and shells, Avi Robinson-Mosher, Tamar Shinar, Jón Grétarsson, Jonathan Su, and Ronald Fedkiw, SIGGRAPH 2008
- Full two-way coupling of rigid and deformable bodies, T Shinar, C Schroeder, R Fedkiw, SIGGRAPH 2008
- Nucleus: Towards a unified dynamics solver for computer graphics, J Stam - Computer-Aided Design and Computer Graphics, 2009
- Robust treatment of collisions, contact and friction for cloth animation, R Bridson, R Fedkiw, J Anderson, SIGGRAPH 2002
- Example-based elastic materials, S Martin, B Thomaszewski, E Grinspun, SIGGRAPH 2011
- Nonconvex rigid bodies with stacking, E Guendelman, R Bridson, R Fedkiw, SIGGRAPH 2003